

Formation

Introduction Deep Learning

Séquence 2

Convolutional neural network (CNN) - 1/2



<https://fidle.cnrs.fr>



Cette session va être enregistrée.
Retrouvez-nous sur notre chaine YouTube :-)

This session will be recorded.
Find us on our YouTube channel :-)

<https://fidle.cnrs.fr/youtube>

Formation

Introduction Deep Learning

Séquence 2

Convolutional neural network (CNN) - 1/2



<https://fidle.cnrs.fr>

<https://fidle.cnrs.fr>

Powered by CNRS CRIC, and UGA DGDSI
of Grenoble, Thanks !



Course materials (pdf)



Practical work environment*



Corrected notebooks



Videos (YouTube)

You can also subscribe to :



FIDLE

<http://fidle.cnrs.fr/listeinfo>



<https://listes.services.cnrs.fr/www/info/devlog>¹



GROUPE **CALCUL**

<https://listes.math.cnrs.fr/www/info/calcul>²

(1) List of ESR* developers,


(2) List of ESR* « calcul » group

Where ESR is Enseignement Supérieur et Recherche, french universities and public academic research organizations



Few little things and concepts to **keep in mind**

1



**History,
Fundamental
Concepts**



**Basic
Regression
DNN**









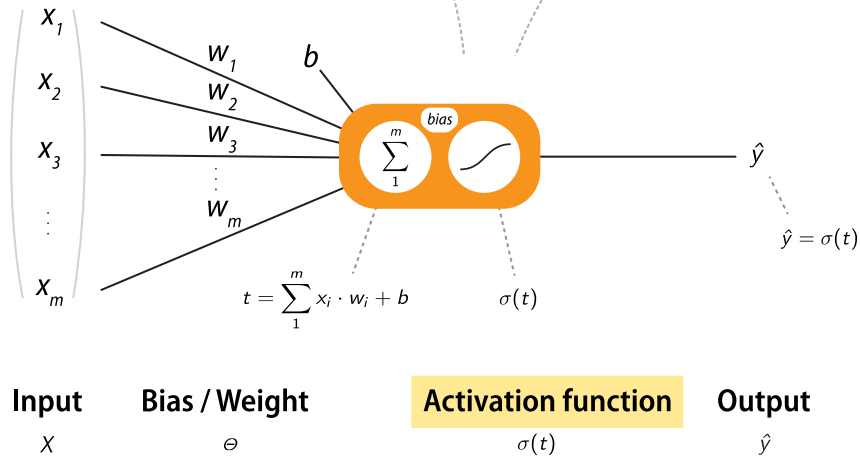
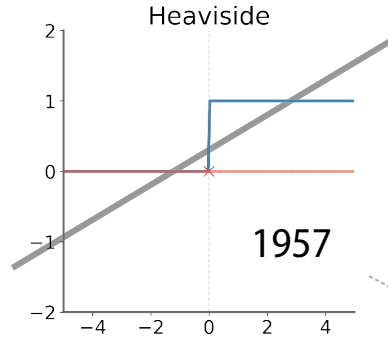
**Basic
Classification
DNN**

- Regression vs. Classification
- Data normalization
- Training and validation
- Epochs and Batches
- Activation functions
- Loss function
- Optimization and gradient descent
- Overfitting
- Metrics
- Softmax and Argmax function
- Numpy shape



Easy :-)

Activation functions



DNN regression and classification

Neurons

$$y = \sigma\left(\sum_1^m y_i \cdot w_i + b\right)$$

Activation : ReLU, etc.

$$\sigma(x) = \max(0, x)$$

Binary classification

Activation : Sigmoid

$$f(x) = \frac{1}{1 + e^{-x}}$$

Loss: Binary cross entropy

$$H(y, \hat{y}) = -\frac{1}{n} \sum_{i=1}^n y_i \cdot \log \hat{y}_i + (1 - y_i) \cdot \log (1 - \hat{y}_i)$$

Muticlass classification

Activation : Softmax

$$\sigma(z)_j = \frac{e^{z_j}}{\sum_{k=1}^K e^{z_k}}$$

Loss: Categorical cross entropy


$$H(y, \hat{y}) = -\sum_{i=1}^n y_i \cdot \log \hat{y}_i$$

Regression



No activation

Loss : MSE, ...

$$\frac{1}{n} \sum_{i=1}^n \left[\hat{y}^{(i)} - y^{(i)} \right]^2$$




About overfitting




$$P_n(x) = a_0 + a_1 \cdot x + a_2 \cdot x^2 + \dots + a_n \cdot x^n = \sum_{i=0}^n a_i \cdot x^i$$

About overfitting






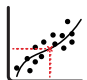
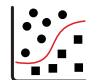







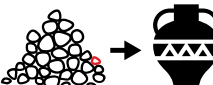



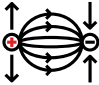




Underfitting



Overfitting



















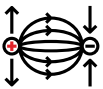


<p>1</p>  <p>History, Fundamental Concepts</p>	<p>2</p>  <p>Hight Dimensionnal Data CNN</p>	<p>4</p>  <p>Demystify mathematics for neural networks.</p>	<p>5</p>  <p>Training strategies Evaluation</p> <p>Sparse data (text) Embedding</p>	<p>6</p>  <p>Sequences data RNN</p>
<p>7</p>  <p>Basic Regression DNN</p>  <p>Basic Classification DNN</p>	<p>7</p>  <p>PyTorch A small detour with PyTorch.</p>	<p>8</p>  <p>«Attention is All You Need» Transformers</p>	<p>9</p>  <p>Graph Neural Network GNN</p> <p>New !</p>	<p>10</p>  <p>Autoencoder networks AE</p>
<p>11</p>  <p>Variational Antoencoder VAE</p>	<p>12</p>  <p>Project session «My project in 180 s»</p>	<p>13</p>  <p>Generative Adversarial Networks GAN</p>	<p>14</p>  <p>Diffusion Model Text to image</p> <p>New !</p>	<p>15</p>  <p>AI, Law, Society and Ethics</p>
<p>16</p>  <p>Model and training optimization Resource efficiency</p> <p>New !</p>	<p>17</p>  <p>Jean-Zay GPU acceleration</p>	<p>18</p>  <p>Physics-Informed Neural Networks PINNS</p> <p>New !</p>	<p>19</p>  <p>Deep Reinforcement Learning RL</p> <p>New !</p>	<p>20</p>  <p>What will be tomorrow's AI Review & perspectives !</p>

20 Séquences
du 17 novembre
au 14 mai 2023



SAISON
22/23



<p>1</p>  <p>History, Fundamental Concepts</p>	<p>2</p>  <p>Hight Dimensionnal Data CNN</p>	<p>4</p>  <p>Demystify mathematics for neural networks.</p>	<p>5</p>  <p>Training strategies Evaluation</p> <p>Sparse data (text) Embedding</p>	<p>6</p>  <p>Sequences data RNN</p>
<p>Basic Regression DNN</p> <p>Basic Classification DNN</p>	<p>7</p>  <p>PyTorch A small detour with PyTorch.</p>	<p>8</p>  <p>«Attention is All You Need» Transformers</p>	<p>9</p> <p>New !</p>  <p>Graph Neural Network GNN</p>	<p>10</p>  <p>Autoencoder networks AE</p>
<p>11</p>  <p>Variational Antoencoder VAE</p>	<p>12</p>  <p>Project session «My project in 180 s»</p>	<p>13</p>  <p>Generative Adversarial Networks GAN</p>	<p>14</p> <p>New !</p>  <p>Diffusion Model Text to image</p>	<p>15</p>  <p>AI, Law, Society and Ethics</p>
<p>16</p> <p>New !</p>  <p>Model and training optimization Resource efficiency</p>	<p>17</p>  <p>Jean-Zay GPU acceleration</p>	<p>18</p> <p>New !</p>  <p>Physics-Informed Neural Networks PINNS</p>	<p>19</p>  <p>Deep Reinforcement Learning RL</p>	<p>20</p>  <p>What will be tomorrow's AI Review & perspectives !</p>

20 Séquences
du 17 novembre
au 14 mai 2023



SAISON
22/23



2.1 What is a **Convolutional Neuron Network** (CNN) ?

- Understanding what a CNN is
- Identify use cases

2.2 **Example 1 : MNIST**

- Implementation of a simple case

3.1 **Example 2 : GTSRB** 🐳

- The devil is also hiding in the data
- How to work with « large » dataset
- Monitoring the training phase and managing our models
- Improve our results with data augmentation
- Datasets and models: how to automate testing
- How to go from notebook to HPC





2.1

What is a **Convolutional Neuron Network (CNN)** ?

- Understanding what a CNN is
- Identify use cases

2.2

Example 1 : MNIST

- Implementation of a simple case

3.1

Example 2 : GTSRB 🐳

- The devil is also hiding in the data
- How to work with « large » dataset
- Monitoring the training phase and managing our models
- Improve our results with data augmentation
- Datasets and models: how to automate testing
- How to go from notebook to HPC



Why convolutional Neural Networks (CNN) ?

For a fully connected layer of (only)
 ± 1000 neurons, we would need to



0.0008 M pixels
28x28, 8 bits



785.000 params




24 M pixels
(r,v,b) 3x8 bits



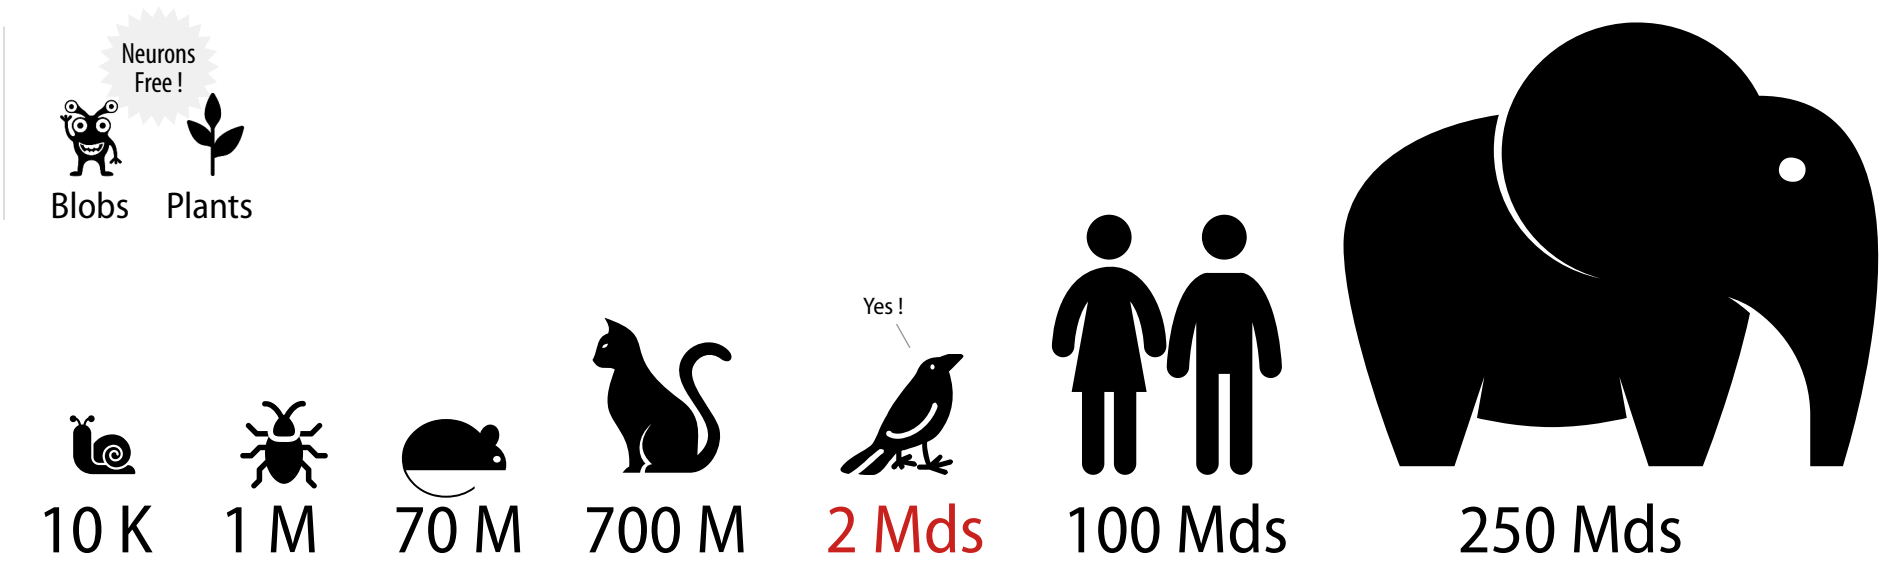
72. 10E9 params...



One neuron is **good**... but more than one is **better** !



One neuron is **good**... but more than one is **better** !



Short
Disgression!


Yes!



Principle of convolutions



By Jan Kroon, from Pexels.com



5	2	1	3	5
4	3	2	3	4
0	2	1	2	2
3	3	2	3	2
2	1	3	2	2

Image piece

5	2	1
4	3	2
0	2	1

x

Kernel 3x3

1	0	1
0	1	0
1	0	1

ω



10

y

$$\begin{aligned} y &= 5 \times 1 + 2 \times 0 + 1 \times 1 \\ &+ 4 \times 0 + 3 \times 1 + 2 \times 0 \\ &+ 0 \times 1 + 2 \times 0 + 1 \times 1 = 10 \end{aligned}$$

$$y = \sum_{i=1}^n \sum_{j=1}^m x_{i,j} \cdot \omega_{i,j} \quad \text{with } \begin{cases} n & \text{kernel width} \\ m & \text{kernel height} \end{cases}$$





⊗ Hadamard product

2D convolution

Principle of convolutions




By Jan Kroon, from Pexels.com




We can perform convolutions in 1, 2, 3 ...or n-dimensional spaces !

Convolutional layers



Number of parameters for a convolutional layer : $kx \cdot ky + 1$



Convolutional layers



If we want to generate m convolutional layers, we will need m convolutional neurons

Convolution parameters



Parameters of a convolutional layer : padding



Keras : padding = 'same' \Rightarrow size is preserved

Convolution parameters



Parameters of a convolutional layer : padding



Keras : padding = 'valid' \Rightarrow Size is not preserved (no padding)


Convolution parameters

Parameters of a convolutional layer : **Strides**




$\text{strides} = (s_x, s_y) \Rightarrow$ A $\text{strides}=(2,2)$ will reduce by 2 the output size

Convolution parameters



padding

‘same’ ‘valid’




kernel size

strides

Number of parameters


For a fully connected layer :



Number of parameters for a DNN layer : $n (m + 1)$

Number of parameters

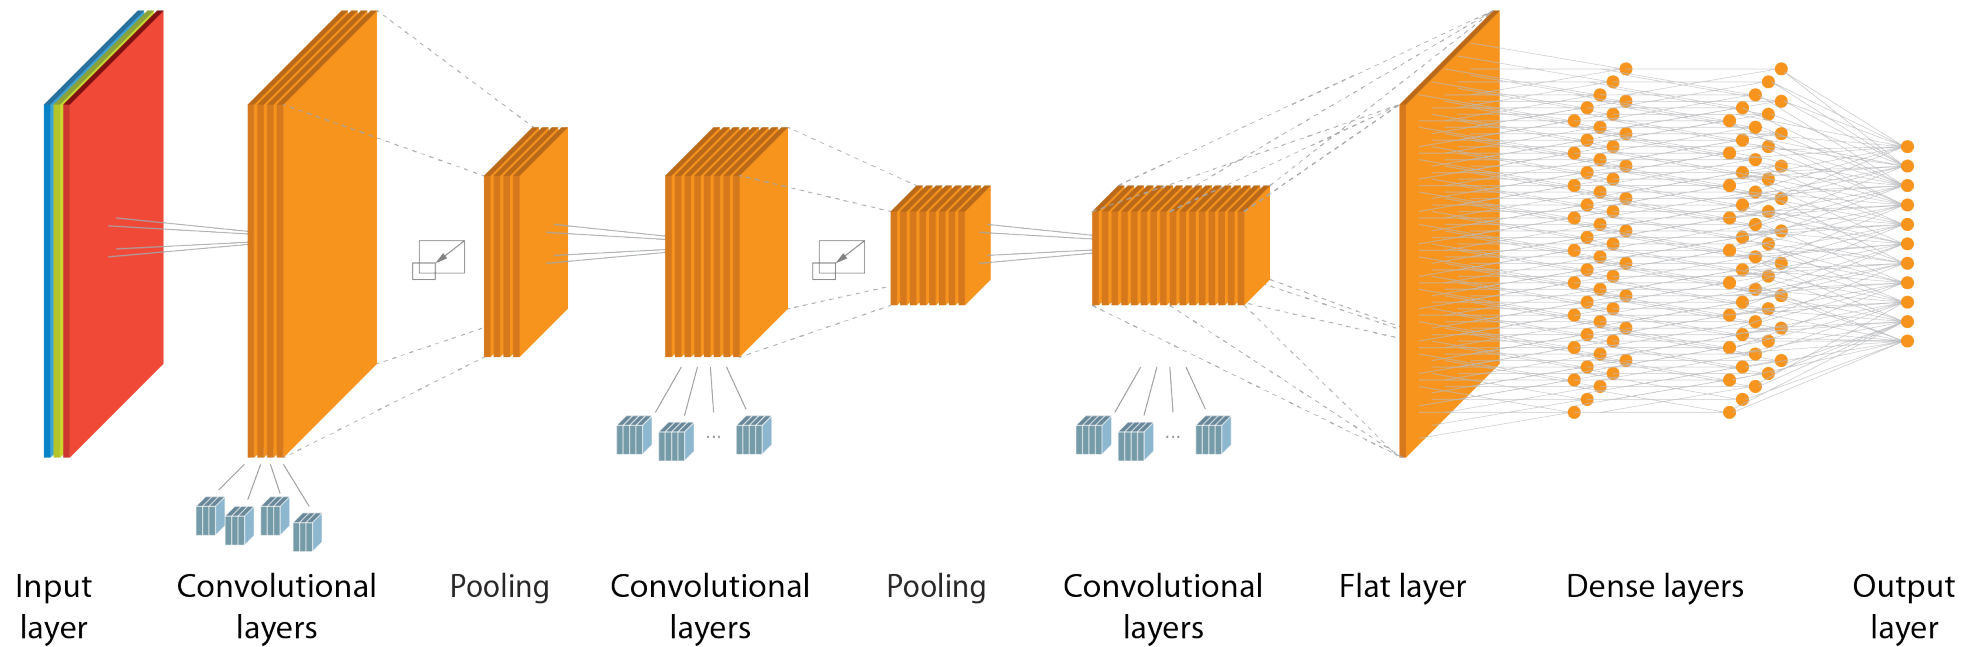
For a convolutional layer :



Number of parameters for a convolutional layer : $n.kx.Ky + 1$





If we want to generate m convolutional layers, we will need m convolutional neurons, so, number of parameters is : $m.(n.kx.ky + 1)$

Convolutional Neural Networks (CNN)



Convolutional Neural Networks (CNN)

Principle of Max Pooling :

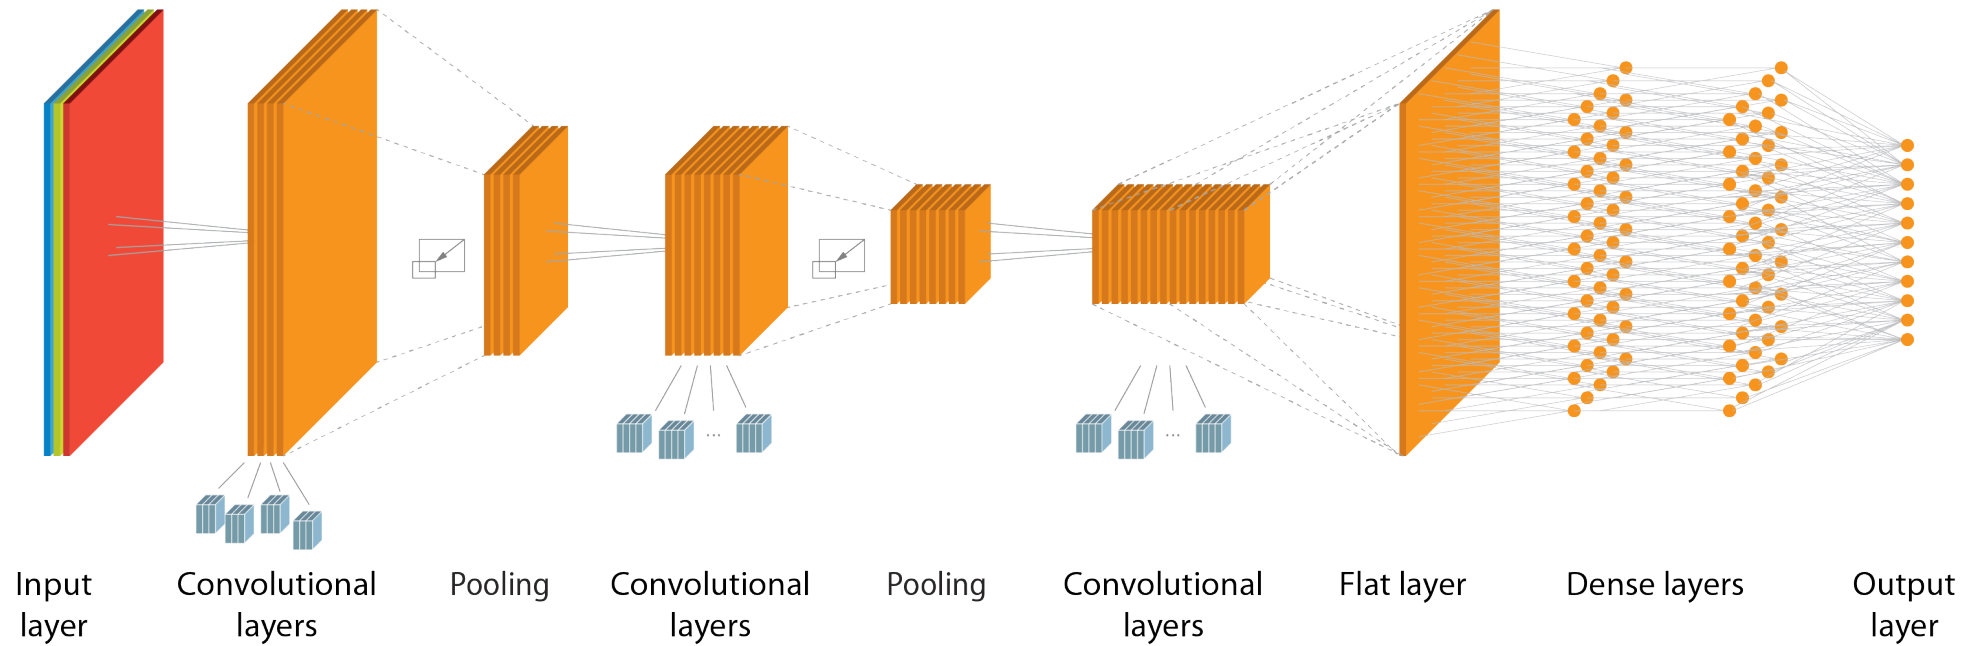


It is possible to set the **window size**, **padding** mode and **strides**.

By default, the strides correspond to the size of the window.

A window (2,2) generates an image twice as small.

Convolutional Neural Networks (CNN)





2.1 What is a **Convolutional Neuron Network (CNN)** ?

- Understanding what a CNN is
- Identify use cases

2.2 **Example 1 : MNIST**

- Implementation of a simple case

3.1 **Example 2 : GTSRB** 🐳

- The devil is also hiding in the data
- How to work with « large » dataset
- Monitoring the training phase and managing our models
- Improve our results with data augmentation
- Datasets and models: how to automate testing
- How to go from notebook to HPC





2.1 What is a **Convolutional Neuron Network** (CNN) ?

- Understanding what a CNN is
- Identify use cases

2.2 **Example 1 : MNIST**

- Implementation of a simple case

3.1 **Example 2 : GTSRB** 🐳

- The devil is also hiding in the data
- How to work with « large » dataset
- Monitoring the training phase and managing our models
- Improve our results with data augmentation
- Datasets and models: how to automate testing
- How to go from notebook to HPC







Image classification with CNN

Notebook : [\[MNIST2\]](#)



Objective :

Recognizing handwritten numbers

Dataset :

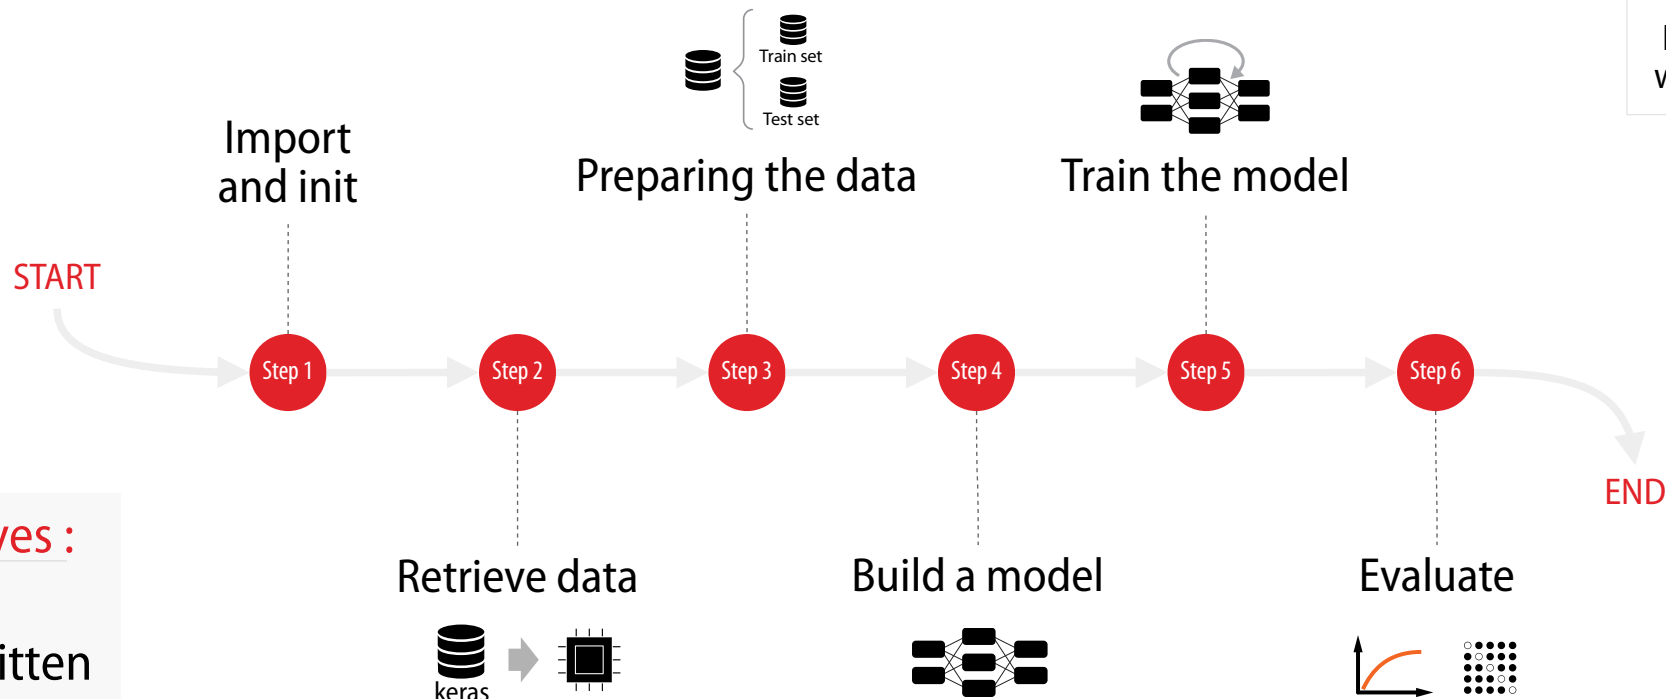
Modified National Institute of Standards and Technology (MNIST)



97.7 %




Previously
with a DNN



Objectives :

Classify
handwritten
numbers
(MNIST
dataset)
via a CNN




Number of parameters for a convolutional layer : $n \cdot k_x \cdot k_y + 1$

If we want to generate m convolutional layers, we will need m convolutional neurons



Understand how it works by understanding where the parameters are...




Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 8)	80
max_pooling2d (MaxPooling2D)	(None, 13, 13, 8)	0
dropout (Dropout)	(None, 13, 13, 8)	0
conv2d_1 (Conv2D)	(None, 11, 11, 16)	1168
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 16)	0
dropout_1 (Dropout)	(None, 5, 5, 16)	0
flatten (Flatten)	(None, 400)	0
dense (Dense)	(None, 100)	40100
dropout_2 (Dropout)	(None, 100)	0
dense_1 (Dense)	(None, 10)	1010

Total params: 42,358
Trainable params: 42,358
Non-trainable params: 0



Previously
with a DNN



Objectives :

Classify
handwritten
numbers
(MNIST
dataset)
via a CNN



Next, on Fidle :



**Jeudi 1 décembre,
14h00**

Séquence 3 :

Réseaux convolutifs, partie 2

Quand les datasets et les calculs grossissent,
problématiques liées à la gestion des données

- Rappel sur les convolutions
- Monitoring et Tensorboard
- Augmentation de données
- Passage à l'échelle (du notebook au batch)
- Points de reprise (checkpoint)

Exemple proposé :

Classification de panneaux routiers

Durée : 2h00

Next on Fidle :




Jeudi 1^{er} décembre, 14h00

Séquence 3 :

Réseaux convolutifs, partie 2

**Quand les datasets et les calculs grossissent,
problématiques liées à la gestion des données**

Merci !



To be continued...