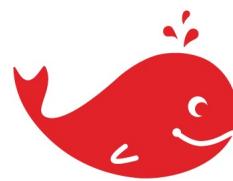




Formation
Introduction au
Deep Learning

Séquence 7

Un petit détour par PyTorch



FIDLE





Cette session va être enregistrée.
Retrouvez-nous sur notre chaîne
YouTube :-)

This session will be recorded.
Find us on our YouTube channel :-)

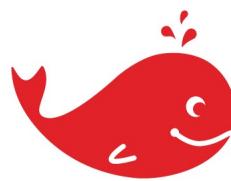
<https://fidle.cnrs.fr/youtube>



Formation
Introduction au
Deep Learning

Séquence 7

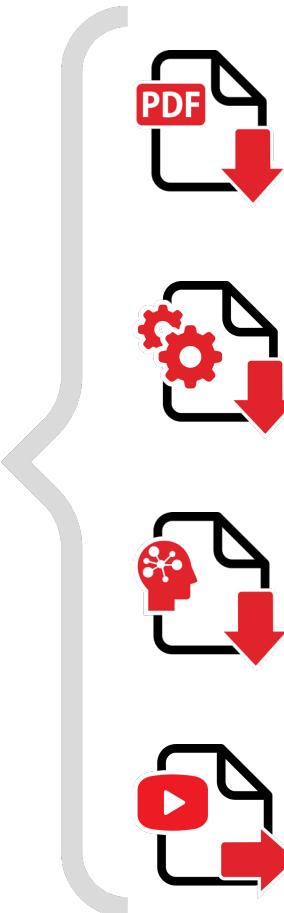
Un petit détour par PyTorch



FIDLE



<https://fiddle.cnrs.fr>



Course materials (pdf)

Practical work environment*

Corrected notebooks

Videos (YouTube)

(*) Procedure via Docket or pip
Remember to get the latest version !

List and networks

You can also subscribe to :



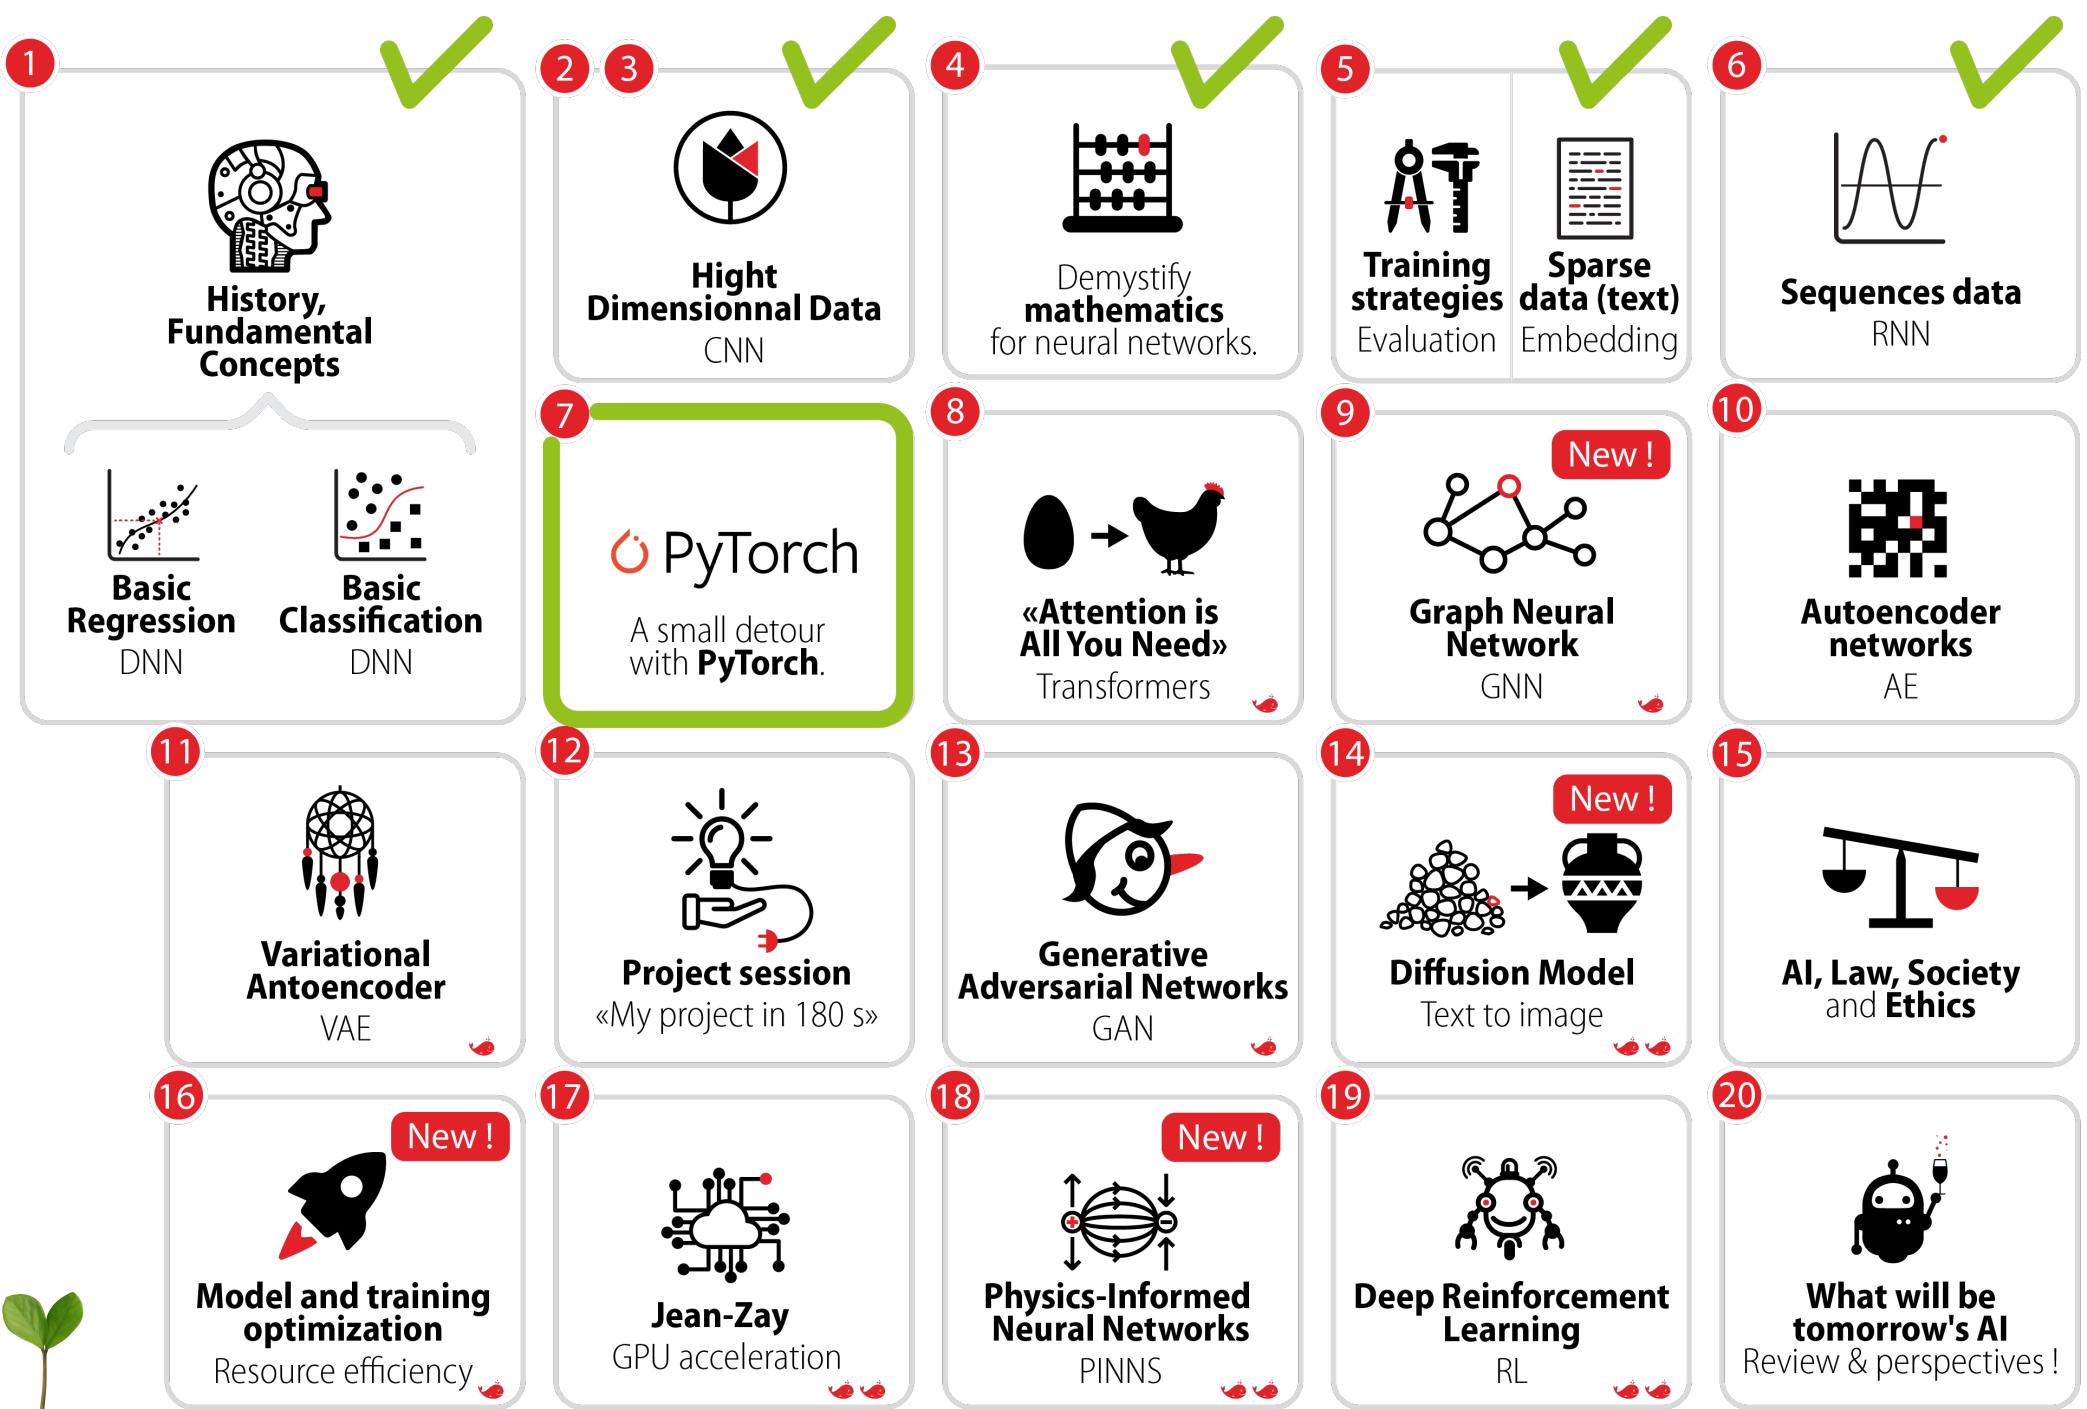
<http://fidle.cnrs.fr/listeinfo>
Fidle information list



<https://listes.services.cnrs.fr/wws/info/devlog>
List of ESR* « calcul » group



<https://listes.math.cnrs.fr/wws/info/calcul>
List of ESR* « calcul » group



SAISON
22/23

7

PyTorch

A small detour
with **PyTorch**.

7.1

Comprendre PyTorch

- Objet pytorch.tensor
- Différentiation automatique
- Calcul GPU

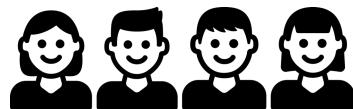
7.2

Etude d'un exemple

7.3

Aller plus loin...

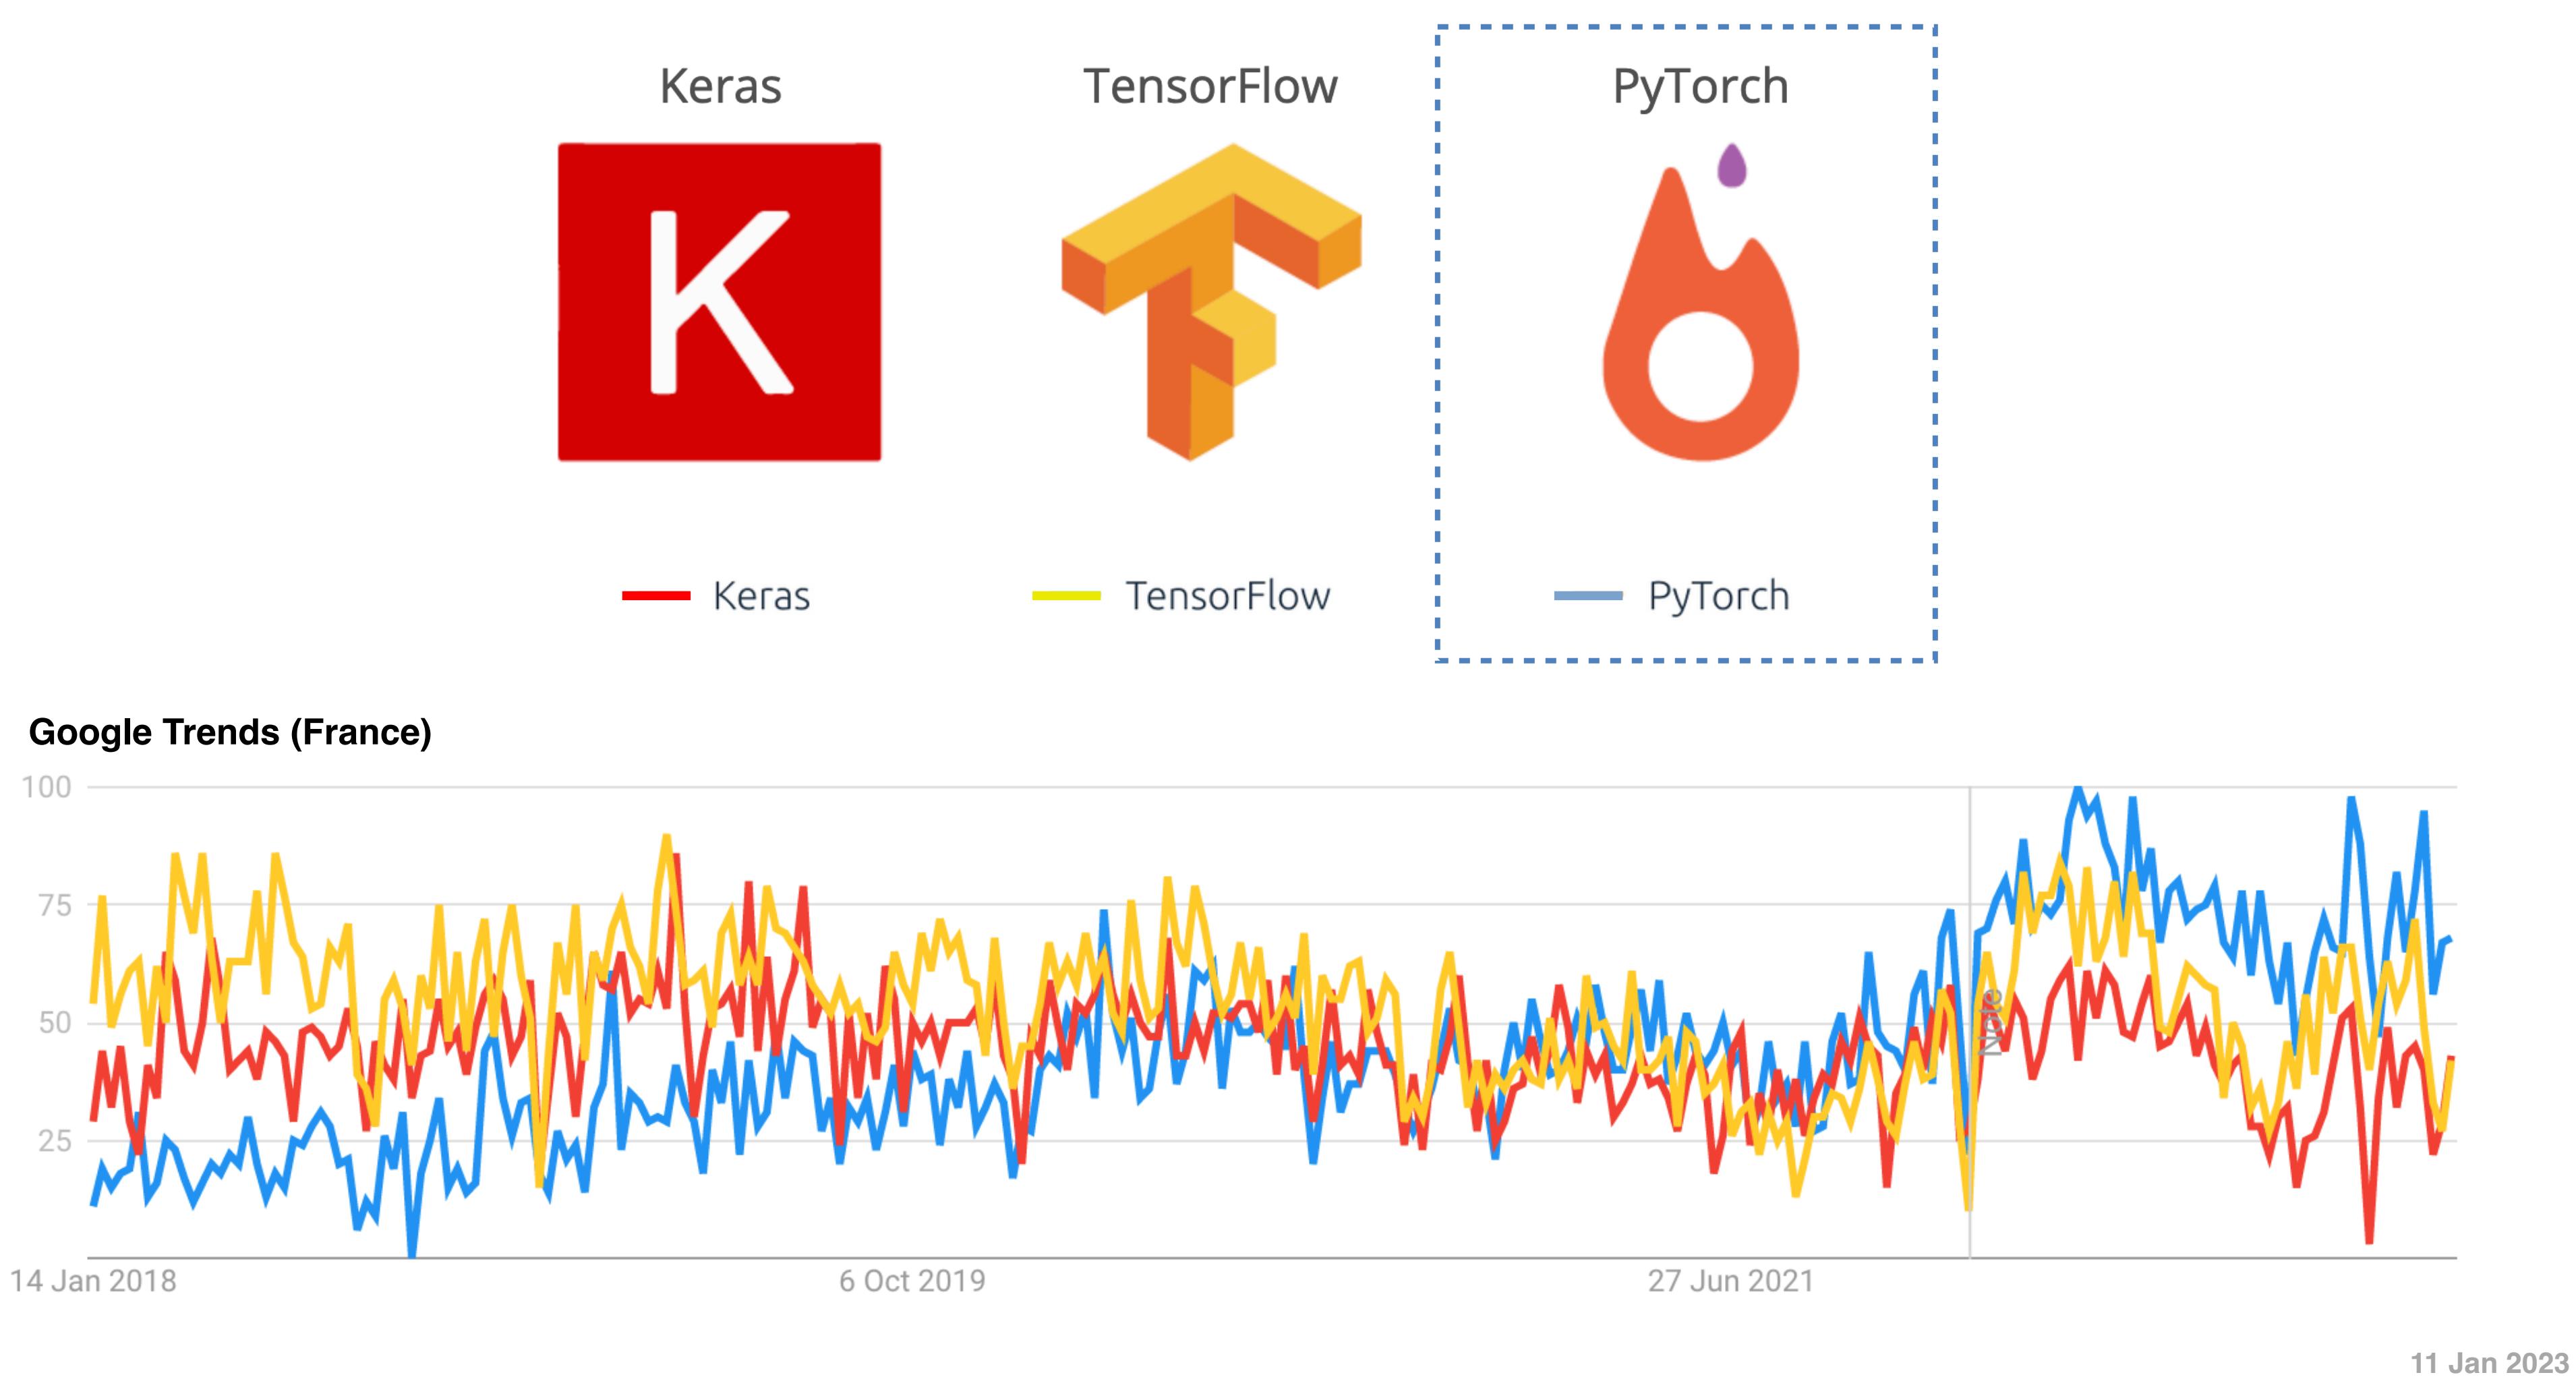
- Charger et adapter une architecture
- Récupérer les paramètres du réseau
- N'apprendre qu'une sous-partie du réseau
- Définir ses propres couches
- Récupérer les données modifiées dans le réseau



Introduction



Introduction



- Complexité intermédiaire entre TensorFlow et Keras
- Rendu open-source par Facebook's AI sur GitHub en 2017
- Relativement simple si les concepts principaux des réseaux de neurones sont assimilés
- Très flexible
- Plus apprécié des *méthodologistes* que des personnes visant une mise en production rapide

Introduction : Installation

→ Trouver la bonne ligne de commande : <https://pytorch.org/get-started/locally/>

PyTorch Build	Stable (1.13.1)	Preview (Nightly)		
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python	C++ / Java		
Compute Platform	CUDA 11.6	CUDA 11.7	ROCM 5.2	CPU
Run this Command:	conda install pytorch torchvision torchaudio pytorch-cuda=11.6 -c pytorch -c nvidia			

Introduction : Comparaison rapide avec Keras

K

fidle/MNIST/01-DNN-MNIST.ipynb

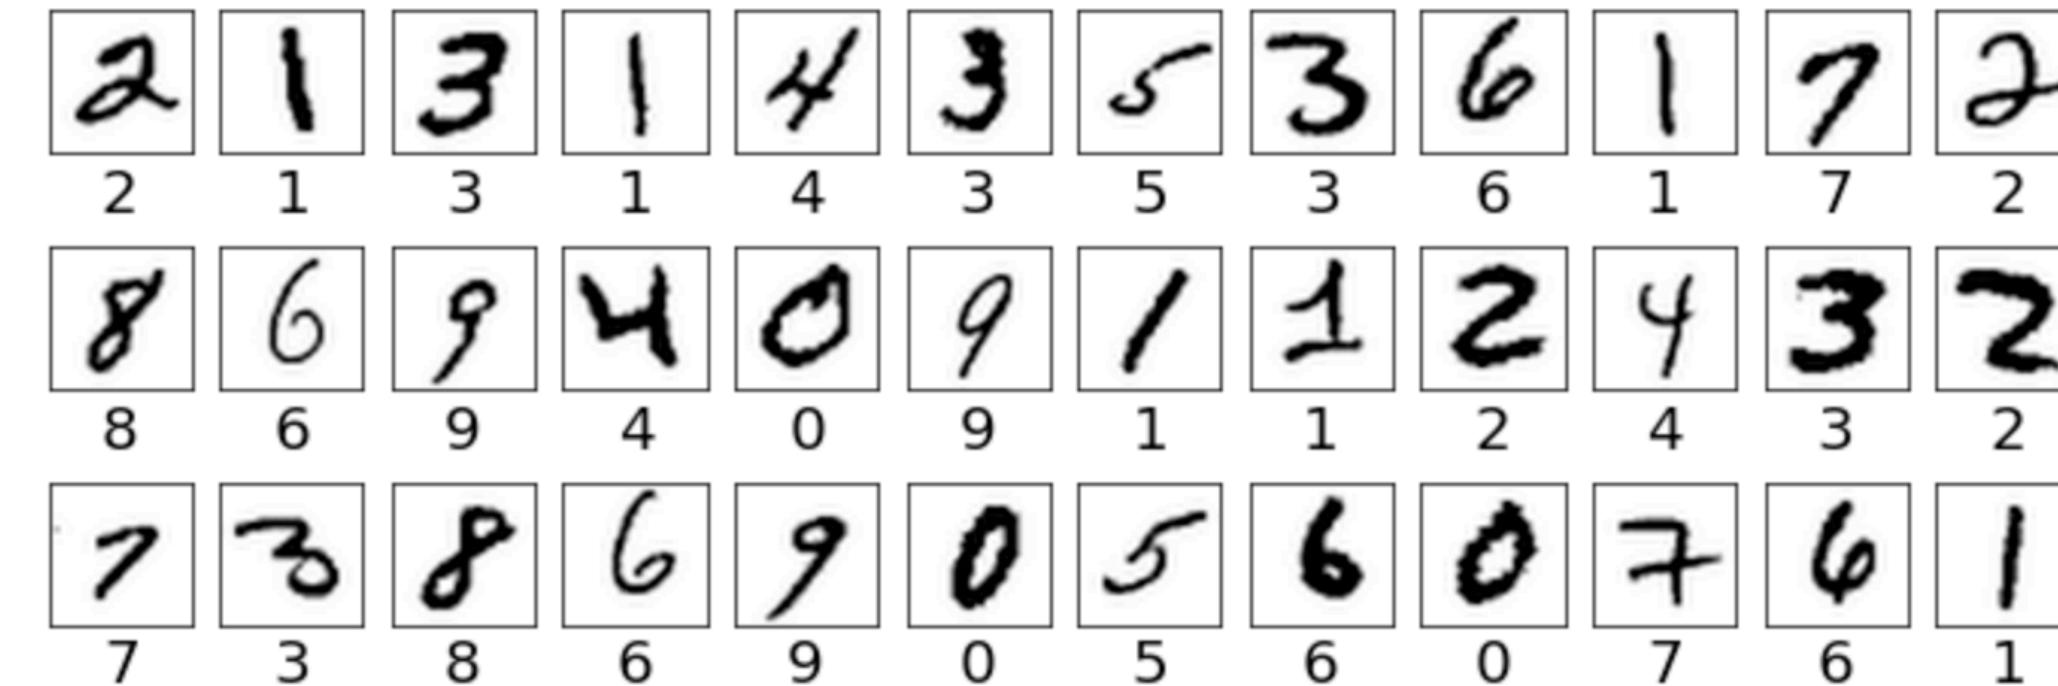
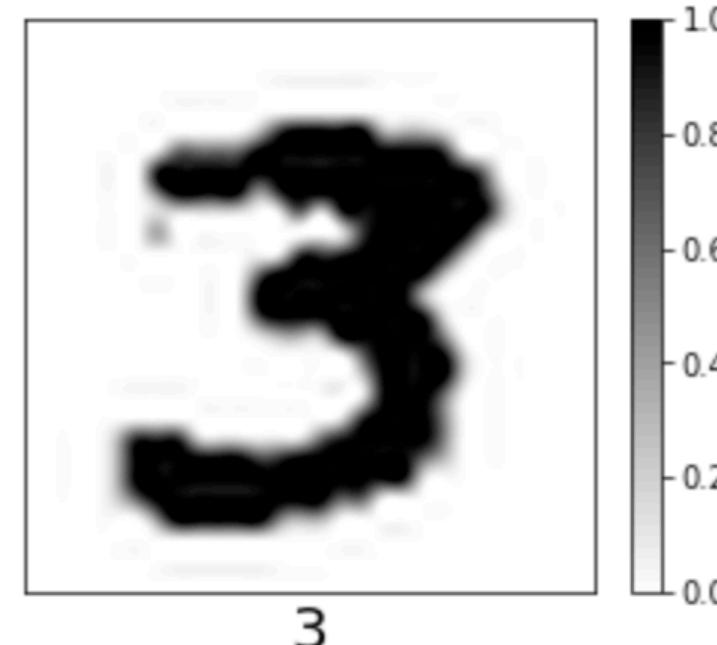
```
(x_train, y_train), (x_test, y_test) = keras.datasets.mnist.load_data()  
  
xmax=x_train.max()  
x_train = x_train / xmax  
x_test = x_test / xmax  
  
pwk.plot_images(x_train, y_train, [27], x_size=5,y_size=5, colorbar=True, save_as='01-one-digit')  
pwk.plot_images(x_train, y_train, range(5,41), columns=12, save_as='02-many-digits')
```

🔥

fidle/MNIST_PyTorch/01-DNN-MNIST_PyTorch.ipynb

```
mnist_trainset = torchvision.datasets.MNIST(root='./data', train=True, download=True, transform=None)  
x_train=mnist_trainset.data.type(torch.DoubleTensor)  
y_train=mnist_trainset.targets  
  
xmax=x_train.max()  
x_train = x_train / xmax  
x_test = x_test / xmax  
np_x_train=x_train.numpy().astype(np.float64)  
np_y_train=y_train.numpy().astype(np.uint8)  
ooo.plot_images(np_x_train,np_y_train , [27], x_size=5,y_size=5, colorbar=True)  
ooo.plot_images(np_x_train,np_y_train, range(5,41), columns=12)
```

Assez similaire... mais utilise autre chose que des numpy.arrays



Introduction : Comparaison rapide avec Keras

K

fidle/MNIST/01-DNN-MNIST.ipynb

```
hidden1      = 100
hidden2      = 100

model = keras.Sequential([
    keras.layers.Input((28,28)),
    keras.layers.Flatten(),
    keras.layers.Dense(hidden1, activation='relu'),
    keras.layers.Dense(hidden2, activation='relu'),
    keras.layers.Dense(10,      activation='softmax')
])

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

F

fidle/MNIST_PyTorch/01-DNN-MNIST_PyTorch.ipynb

```
class MyModel(nn.Module):
    """
    Basic fully connected neural-network
    """

    def __init__(self):
        hidden1      = 100
        hidden2      = 100
        super(MyModel, self).__init__()
        self.hidden1 = nn.Linear(784, hidden1)
        self.hidden2 = nn.Linear(hidden1, hidden2)
        self.hidden3 = nn.Linear(hidden2, 10)

    def forward(self, x):
        x = x.view(-1,784)
        x = self.hidden1(x)
        x = F.relu(x)
        x = self.hidden2(x)
        x = F.relu(x)
        x = self.hidden3(x)
        x = F.softmax(x, dim=0)
        return x

model = MyModel()
loss = nn.CrossEntropyLoss()
optimizer = torch.optim.Adam(model.parameters(), lr=1e-3)
```

Plus verbeux mais potentiellement plus flexible

Introduction : Comparaison rapide avec Keras

K

fidle/MNIST/01-DNN-MNIST.ipynb

```
batch_size = 512
epochs = 16

history = model.fit(x_train, y_train,
                     batch_size = batch_size,
                     epochs = epochs,
                     verbose = 1,
                     validation_data = (x_test, y_test))

score = model.evaluate(x_test, y_test, verbose=0)
```

T

fidle/MNIST_PyTorch/01-DNN-MNIST_PyTorch.ipynb

```
def fit(model,X_train,Y_train,X_test,Y_test, EPOCHS = 5, BATCH_SIZE = 32):

    loss = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-3) #lr is the learning rate
    model.train()

    history=convergence_history_CrossEntropyLoss()

    history.update(model,X_train,Y_train,X_test,Y_test)

    n=X_train.shape[0] #number of observations in the training data

    #stochastic gradient descent
    for epoch in range(EPOCHS):

        batch_start=0
        epoch_shuffler=np.arange(n)
        np.random.shuffle(epoch_shuffler) #remark that 'utilsData.DataLoader' could be used instead

        while batch_start+BATCH_SIZE < n:
            #get mini-batch observation
            mini_batch_observations = epoch_shuffler[batch_start:batch_start+BATCH_SIZE]
            var_X_batch = Variable(X_train[mini_batch_observations,:,:]).float() #the input image is flattened
            var_Y_batch = Variable(Y_train[mini_batch_observations])

            #gradient descent step
            optimizer.zero_grad() #set the parameters gradients to 0
            Y_pred_batch = model(var_X_batch) #predict y with the current NN parameters

            curr_loss = loss(Y_pred_batch, var_Y_batch) #compute the current loss
            curr_loss.backward() #compute the loss gradient w.r.t. all NN parameters
            optimizer.step() #update the NN parameters

            #prepare the next mini-batch of the epoch
            batch_start+=BATCH_SIZE

        history.update(model,X_train,Y_train,X_test,Y_test)

    return history

batch_size = 512
epochs = 128

history=fit(model,x_train,y_train,x_test,y_test,EPOCHS=epochs,BATCH_SIZE = batch_size)

var_x_test = Variable(x_test[:, :, :]).float()
var_y_test = Variable(y_test[:])
y_pred = model(var_x_test)

loss = nn.CrossEntropyLoss()
curr_loss = loss(y_pred, var_y_test)

val_loss = curr_loss.item()
val_accuracy = float( (torch.argmax(y_pred, dim= 1) == var_y_test).float().mean() )
```

Définitivement plus verbeux... mais :

- Permet de facilement rentrer dans des détails algorithmique
- Donne un bon accès aux données intermédiaires



1 : Comprendre PyTorch



PyTorch tensors

- Visiblement très proche d'un numpy array, mais dans PyTorch
- Peut être utilisé pour :
 - Rétro-propager le calcul d'un gradient (voir partie 1.2)
 - Transférer facilement des données sur un GPU (voir partie 1.3)

Instantiation :

```
import torch  
torch.tensor([[1,-2],[3,4]])  
  
tensor([[ 1, -2],  
       [ 3,  4]])
```

```
import numpy as np  
torch.tensor(np.array([[1,-2],[3,4]]))  
  
tensor([[ 1, -2],  
       [ 3,  4]])
```

```
torch.zeros([2,4],dtype=torch.float32)  
  
tensor([[0., 0., 0., 0.],  
       [0., 0., 0., 0.]])
```

1.1 Comprendre PyTorch — Objet pytorch.tensor

Accès à une valeur :

```
A=torch.tensor([[1,-2],[3,4]])  
print(A[0,1])
```

tensor(-2)

```
print(A[0,1].item())
```

-2

Conversion en numpy :

```
print(A.numpy())
```

```
[[ 1 -2]  
 [ 3  4]]
```

1.1 Comprendre PyTorch — Objet pytorch.tensor

Echanger deux axes (transposée) :

```
A=torch.randn([10,20,30])  
print(A.shape)
```

```
torch.Size([10, 20, 30])
```

```
A_t=torch.transpose(A,0,2)  
print(A.shape)  
print(A_t.shape)
```

```
torch.Size([10, 20, 30])  
torch.Size([30, 20, 10])
```

Equivalent du « Reshape » :

```
A=torch.randint(0,10,[4,4])  
print(A)
```

```
tensor([[7, 6, 2, 9],  
       [4, 1, 0, 9],  
       [9, 3, 9, 6],  
       [3, 2, 2, 5]])
```

```
print(A.view([2,8]))
```

```
tensor([[7, 6, 2, 9, 4, 1, 0, 9],  
       [9, 3, 9, 6, 3, 2, 2, 5]])
```

```
print(A.view([-1,2,2]))
```

```
tensor([[[7, 6],  
        [2, 9]],  
      [[4, 1],  
       [0, 9]],  
      [[9, 3],  
       [9, 6]],  
      [[3, 2],  
       [2, 5]]])
```

```
print(A.flatten())
```

```
tensor([7, 6, 2, 9, 4, 1, 0, 9, 9, 3, 9, 6, 3,  
       2, 2, 5])
```

1.1 Comprendre PyTorch — Objet pytorch.tensor

Changer le type des données (cast) :

```
A=torch.randn([2,2],dtype=torch.float32)
```

```
print(A.type())
```

```
torch.FloatTensor
```

```
B=A.type(torch.DoubleTensor)
print(A.type())
print(B.type())
```

```
torch.FloatTensor
```

```
torch.DoubleTensor
```

```
C=A.int()
print(A.type())
print(C.type())
```

```
torch.FloatTensor
```

```
torch.IntTensor
```

Pour aller plus loin : <https://pytorch.org/docs/stable/tensors.html>

1.2 Comprendre PyTorch — Différentiation automatique

```

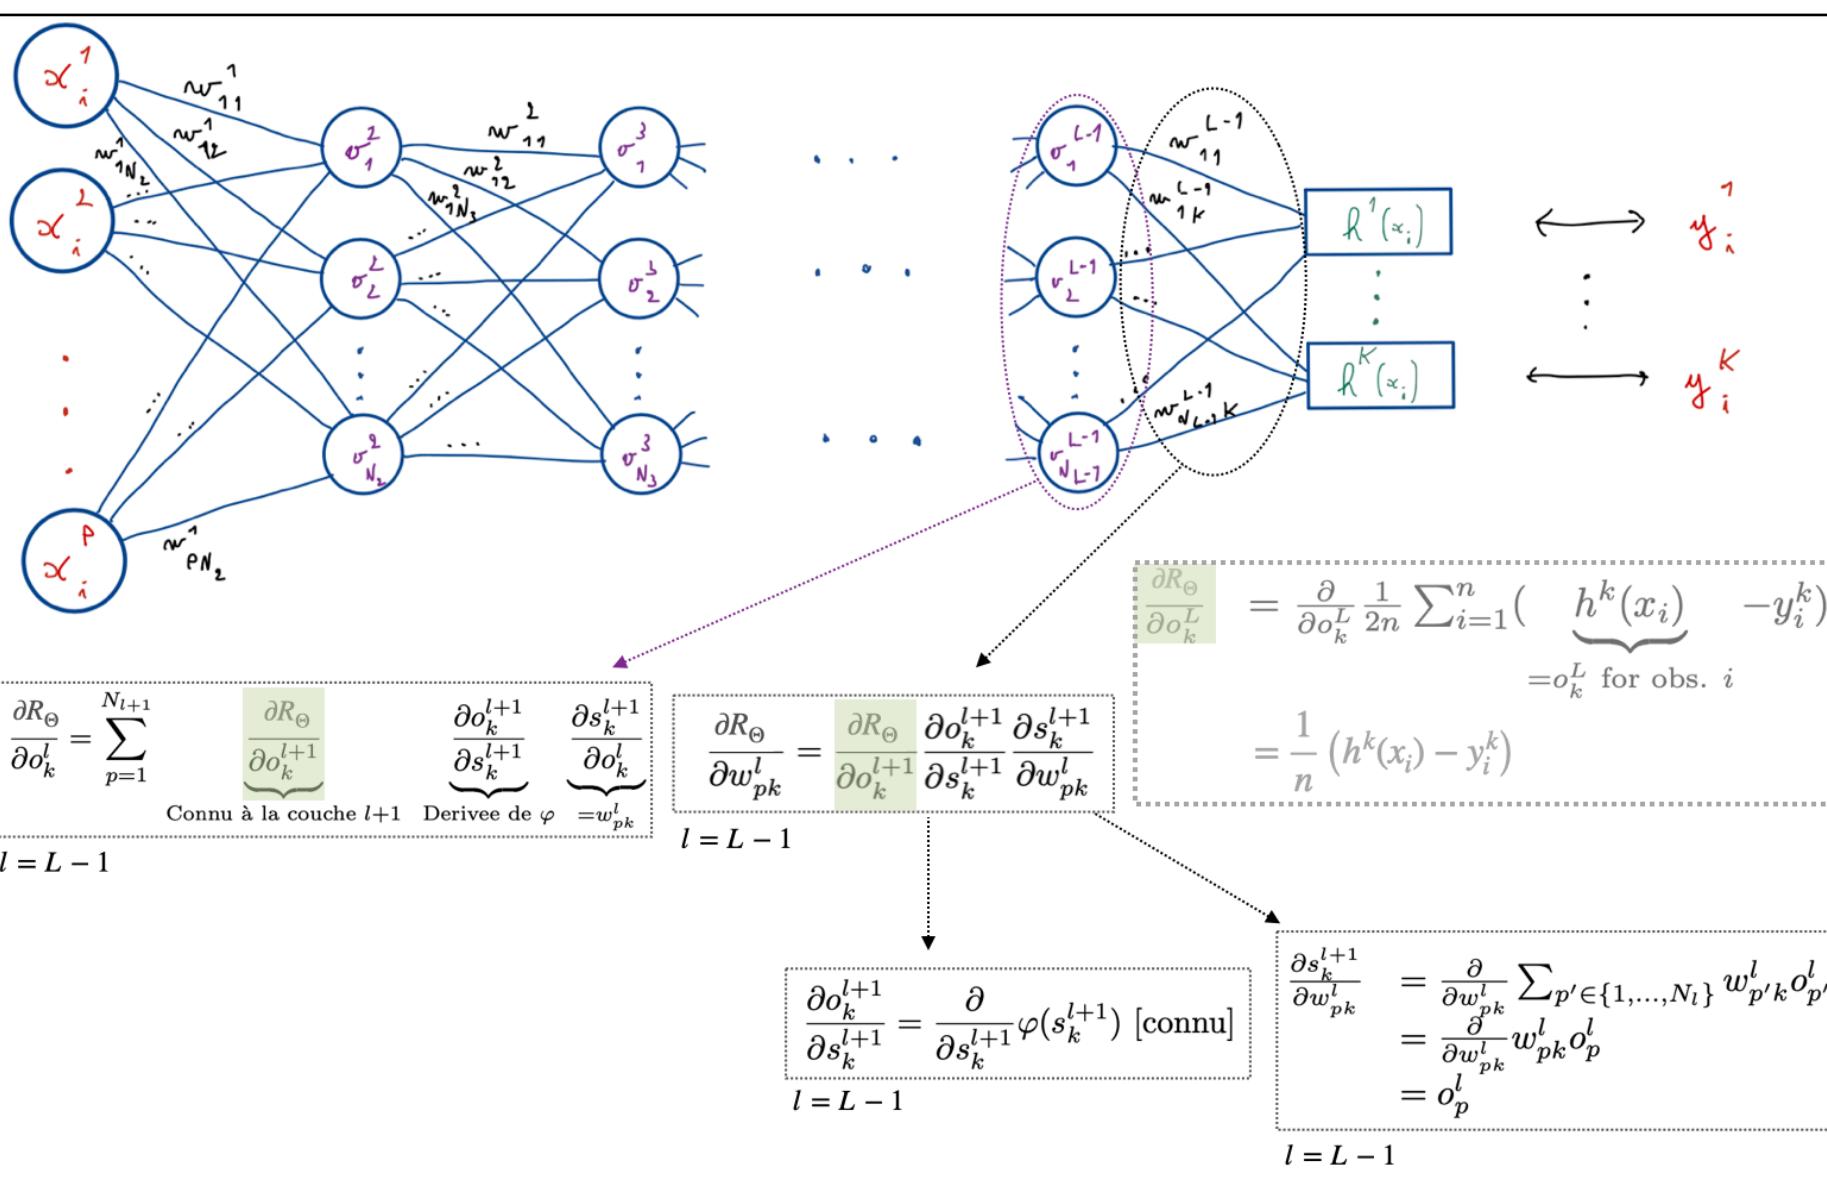
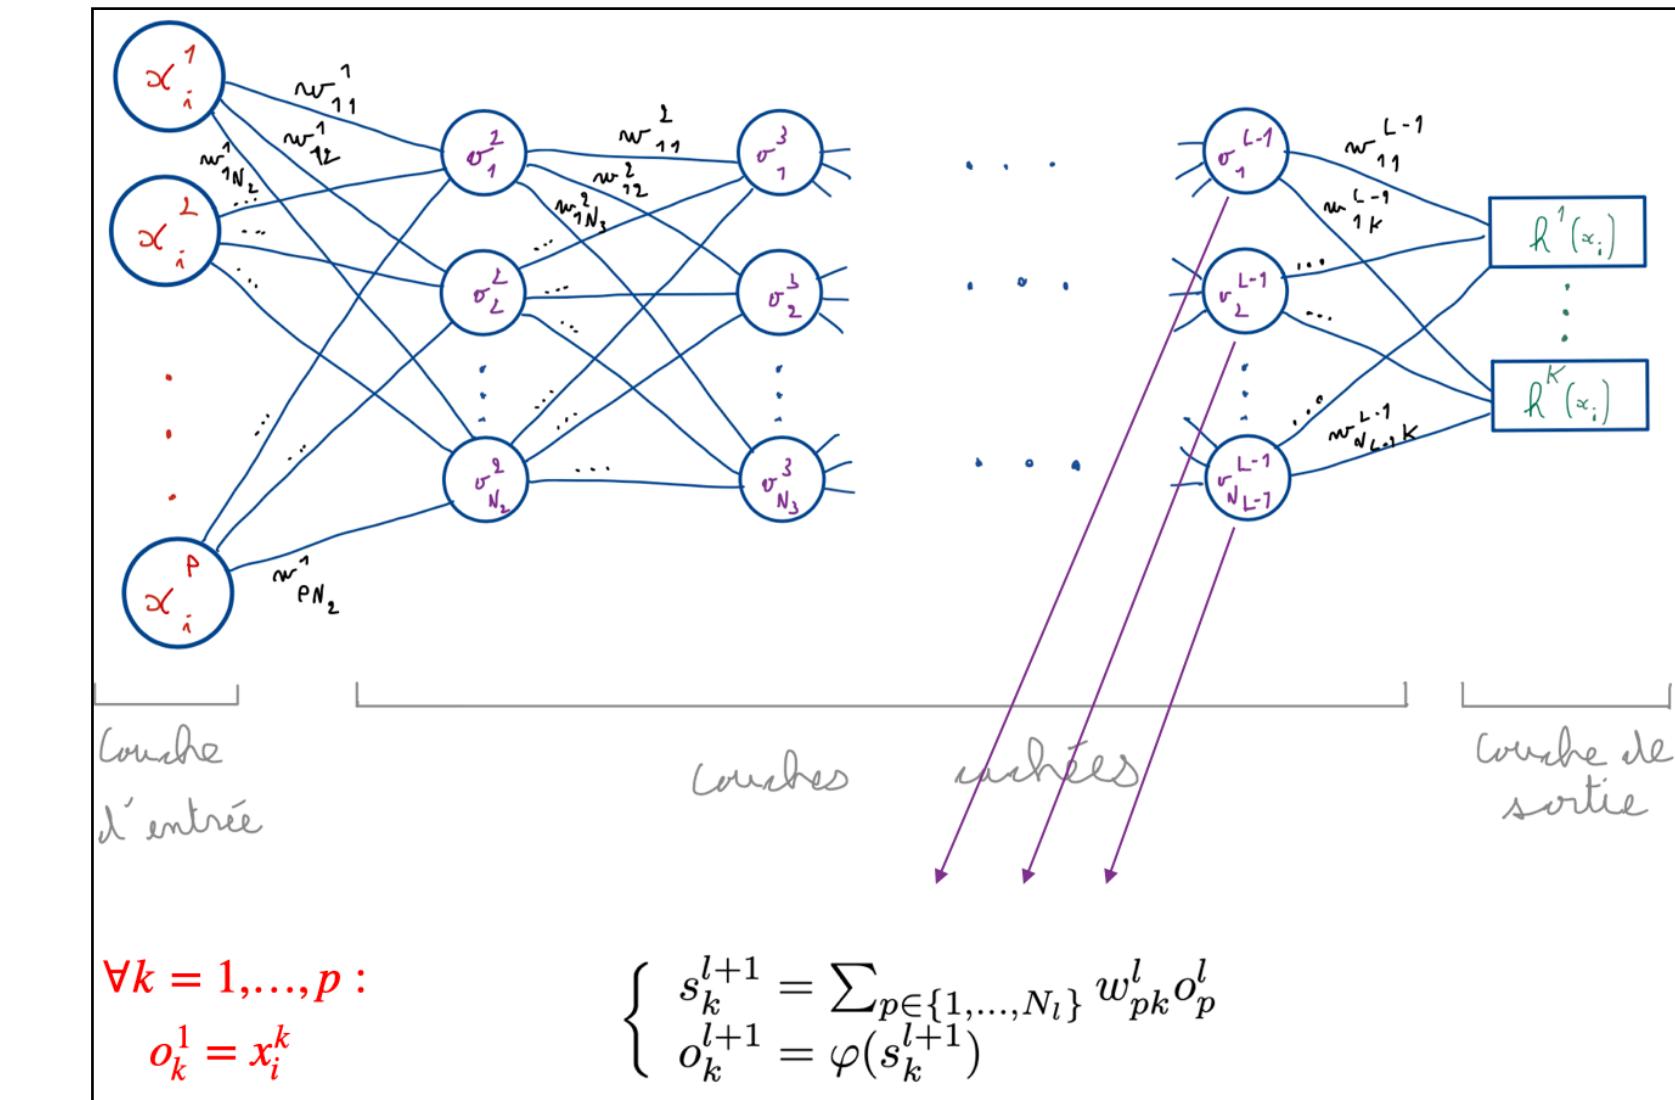
class MyModel(nn.Module):
    """
    Basic fully connected neural-network
    """

    def __init__(self):
        hidden1 = 100
        hidden2 = 100
        super(MyModel, self).__init__()
        self.hidden1 = nn.Linear(784, hidden1)
        self.hidden2 = nn.Linear(hidden1, hidden2)
        self.hidden3 = nn.Linear(hidden2, 10)

    def forward(self, x):
        x = x.view(-1, 784)
        x = self.hidden1(x)
        x = F.relu(x)
        x = self.hidden2(x)
        x = F.relu(x)
        x = self.hidden3(x)
        x = F.softmax(x, dim=0)
        return x

```

Modèle
« forward »
bien défini



Calcul du gradient du loss
par rapport aux paramètres
du modèle grâce à la
différentiation automatique !

Exemple élémentaire (régression linéaire)

- Estimation \hat{y} de la note à un futur concours en fonction de notes déjà obtenues en Français (X^1) et en Mathématiques (X^2).
- On suppose que \hat{y} est la moyenne pondérée de X^1 et $X^2 \rightarrow \hat{y} = f_{\Theta}(X) = w_1 X^1 + w_2 X^2$.
- Les poids w_1 et w_2 sont appris sur un jeu de données $\{X_i, y_i\}_{i=1,\dots,n}$ contenant les notes obtenues l'an dernier.

Apprentissage des paramètres w_1 et w_2

- On minimise par descente de gradient $R(w_1, w_2) = \sum_{i=1}^n (\hat{y}_i - y_i)^2 = \sum_{i=1}^n ((w_1 X_i^1 + w_2 X_i^2) - y)^2$
- Calcul du gradient $\nabla R(w_1, w_2) = \left(\frac{\partial R}{\partial w_1}, \frac{\partial R}{\partial w_2} \right)^T$ par différentiation automatique

1.2 Comprendre PyTorch — Différentiation automatique

```
import torch

X = torch.randn([5,2], requires_grad=False)
w = torch.randn([2], requires_grad=True)

y_pred=w[0]*X[:,0]+w[1]*X[:,1]
#y_pred,retain_grad()

print('X :',X)
print('w :',w)
print('y_pred :',y_pred)

print('X.grad :',X.grad)
print('w.grad :',w.grad)
print('y_pred.grad :',y_pred.grad)
```

X : tensor([[0.2590, 1.1822],
 [1.8838, 0.7159],
 [0.3180, -0.9601],
 [0.4899, 2.0838],
 [1.2439, 0.1527]])
w : tensor([0.3187, 1.3714], requires_grad=True)
y_pred : tensor([1.7039, 1.5821, -1.2153, 3.0139, 0.6058], grad_fn=<AddBackward0>)
X.grad : None
w.grad : None
y_pred.grad : None

```
y_true=torch.zeros(5, requires_grad=False)

print('y_true :',y_true)
y_true : tensor([0., 0., 0., 0., 0.])

Risk=(torch.pow(y_pred-y_true,2.)).mean()

Risk.backward()

print('X.grad :',X.grad)
X.grad : None

print('y_pred.grad :',y_pred.grad)

y_pred.grad : None

print('w.grad :',w.grad)

w.grad : tensor([2.1062, 4.2746])
```

1.2 Comprendre PyTorch — Différentiation automatique

```
import torch

X = torch.randn([5,2], requires_grad=False)
w = torch.randn([2], requires_grad=True)

y_pred=w[0]*X[:,0]+w[1]*X[:,1]
y_pred,retain_grad()

print('X :',X)
print('w :',w)
print('y_pred :',y_pred)

print('X.grad :',X.grad)
print('w.grad :',w.grad)
print('y_pred.grad :',y_pred.grad)
```

```
y_true=torch.zeros([5], requires_grad=False)

print('y_true :',y_true)
y_true : tensor([0., 0., 0., 0., 0.])

Risk=(torch.pow(y_pred-y_true,2.)).mean()

Risk.backward()

print('X.grad :',X.grad)
X.grad : None

print('y_pred.grad :',y_pred.grad)
y_pred.grad : tensor([ 0.6815,  0.6328, -0.4861,
1.2056,  0.2423])

print('w.grad :',w.grad)
w.grad : tensor([2.1062, 4.2746])
```

```
X : tensor([[ 0.2590,  1.1822],
           [ 1.8838,  0.7159],
           [ 0.3180, -0.9601],
           [ 0.4899,  2.0838],
           [ 1.2439,  0.1527]])
w : tensor([0.3187, 1.3714], requires_grad=True)
y_pred : tensor([ 1.7039,  1.5821, -1.2153,  3.0139,  0.6058], grad_fn=<AddBackward0>)
X.grad : None
w.grad : None
y_pred.grad : None
```

1.2 Comprendre PyTorch — Différentiation automatique

Utilisation de detach() : pour qu'un tenseur ne soit plus qu'une constante (on ne calcule plus son gradient)

```
x=torch.ones(10, requires_grad=True)

y=x**2
z=x**3

r=(y+z).sum()
r.backward()
print(x.grad)
```

```
tensor([5., 5., 5., 5., 5., 5., 5., 5., 5., 5.])
```

```
x=torch.ones(10, requires_grad=True)

y=x**2
z=x.detach()**3

r=(y+z).sum()
r.backward()

print(x.grad)
```

```
tensor([2., 2., 2., 2., 2., 2., 2., 2., 2., 2.])
```

1.2 Comprendre PyTorch — Différentiation automatique

Utilisation de torch.no_grad() : pour ne calculer aucun gradient et limiter la mémoire nécessaire

```
x = torch.ones(3, requires_grad=True)
y = x**2
z = x**3
r = (y+z).sum()
r.backward()
print('r :',r)
print('x.grad :',x.grad)

r : tensor(6., grad_fn=<SumBackward0>)
x.grad : tensor([5., 5., 5.])
```

```
x = torch.ones(3, requires_grad=True)
with torch.no_grad():
    y = x**2
    z = x**3
    r = (y+z).sum()
print('r :',r)
print('x.grad :',x.grad)
```

```
r : tensor(6.)
x.grad : None
```

```
r.backward()
```

RuntimeError Traceback
(most recent call last)
<ipython-input-232-d0c0300fbddc> in <module>
----> 1 r.backward()

1.2 Comprendre PyTorch — Différentiation automatique

Mieux comprendre la différentiation automatique

5 gradient/derivative related PyTorch functions

 Attyuttam Saha Jun 16, 2020 · 6 min read



In this article, I will be talking about the 5 PyTorch functions that I have studied through. Examples will be provided along with scenarios when the functions might break. This article is a great head start to explore PyTorch and the various plethora of functionalities it provides.

medium.com/@attyuttam/5-gradient-derivative-related-pytorch-functions-8fd0e02f13c6

PyTorch

Table of Contents

AUTOGRADE MECHANICS

This note will present an overview of how autograd works and records the operations. It's not strictly necessary to understand all this, but we recommend getting familiar with it, as it will help you write more efficient, cleaner programs, and can aid you in debugging.

pytorch.org/docs/stable/notes/autograd.html

Automatic differentiation for applied mathematicians

Is PyTorch the right tool for you?

Jean Feydy

February 2018

Écoles Normales Supérieures de Paris et Paris-Saclay

www.jeanfeydy.com/Talks/autodiff_appliedmaths/AutoDiff_AppliedMaths.pdf

PyTorch

Table of Contents

AUTOMATIC DIFFERENTIATION PACKAGE - TORCH.AUTOGRAD

pytorch.org/docs/stable/autograd.html

1.3 Comprendre PyTorch — Calcul sur GPU

Voir si une carte GPU (Nvidia) est disponible :

```
import torch

print('Cuda disponible : ',torch.cuda.is_available())
print('Nombre de devices : ',torch.cuda.device_count())
print('Device utilisé : ',torch.cuda.current_device())
print('Nom du device utilisé : ',torch.cuda.get_device_name(0))
```

```
Cuda disponible :  True
Nombre de devices :  1
Device utilisé :  0
Nom du device utilisé :  Tesla K80
```

1.3 Comprendre PyTorch — Calcul sur GPU

Mécanisme de transferts similaire à ce qui se fait usuellement avec CUDA :

```
import torch
import torchvision.models as models

import torch.nn as nn

h_resnet18 = models.resnet18()
h_resnet18.fc = nn.Linear(512,5)

d_resnet18 = h_resnet18.cuda()

h_simulated_data = torch.randn([1000,3,64,64])

d_mini_batch = h_simulated_data[0:10,:,:,:].cuda()

d_outputs = d_resnet18(d_mini_batch)

h_outputs = d_outputs.cpu()

print(h_outputs)

tensor([[ 0.2605, -0.4549, -0.5667,  0.8204, -0.4016],
        [ 0.5252, -0.7144, -0.4012,  0.7757, -0.8269],
        [ 0.6739, -0.2131, -1.0644,  1.0779, -0.7419],
        ...
        [ 0.7548,  0.0197, -0.2869,  1.0721, -0.6495]],
grad_fn=<ToCopyBackward0>)
```

1.3 Comprendre PyTorch — Calcul sur GPU

Peut être rendu générique (CPU ou GPU) :

```
import torch
import torchvision.models as models

DEVICE = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
print(DEVICE)

cuda:0

import torch.nn as nn

h_resnet18 = models.resnet18()
h_resnet18.fc = nn.Linear(512,5)

d_resnet18 = h_resnet18.to(DEVICE)

h_simulated_data = torch.randn([1000,3,64,64])

d_mini_batch = h_simulated_data[0:10,:,:,:].to(DEVICE)

d_outputs = d_resnet18(d_mini_batch)

h_outputs = d_outputs.to('cpu')

print(h_outputs)

tensor([[ 0.2605, -0.4549, -0.5667,  0.8204, -0.4016],
        [ 0.5252, -0.7144, -0.4012,  0.7757, -0.8269],
        [ 0.6739, -0.2131, -1.0644,  1.0779, -0.7419],
        ...
        [ 0.7548,  0.0197, -0.2869,  1.0721, -0.6495]],

grad_fn=<ToCopyBackward0>)
```

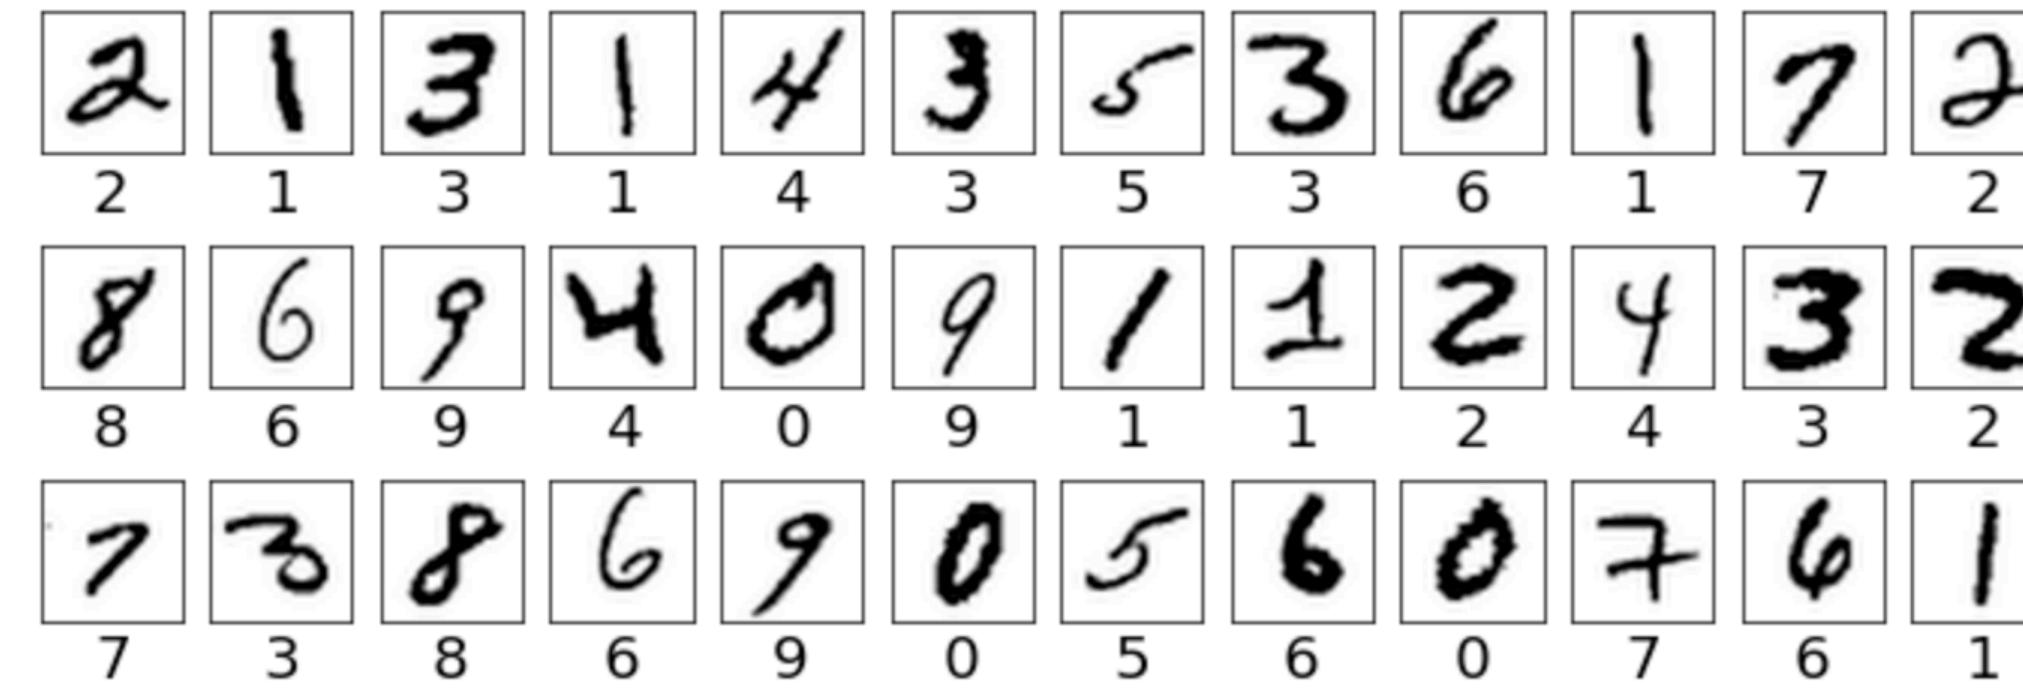
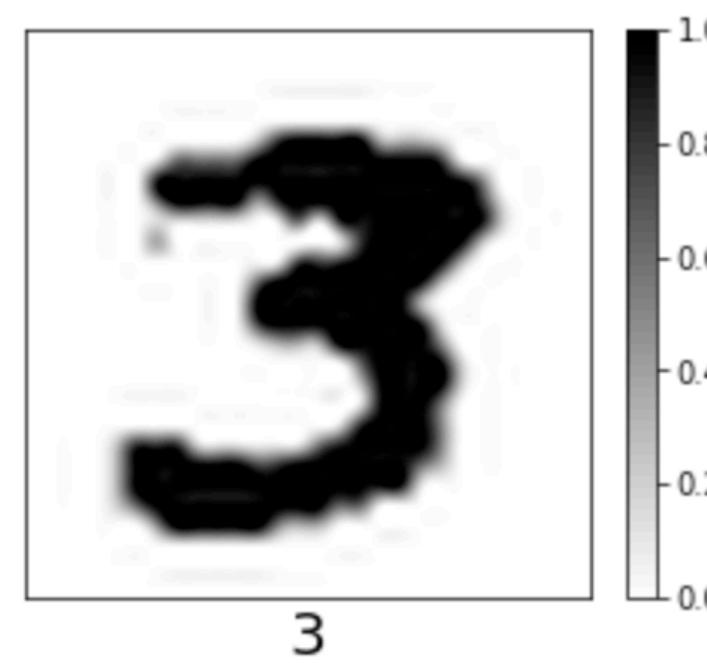
2 : Etude d'un exemple



2 Etude d'un exemple — Détail du code d'apprentissage

```
#get and format the training set
mnist_trainset = torchvision.datasets.MNIST(root='./data', train=True, download=True, transform=None)
x_train=mnist_trainset.data.type(torch.DoubleTensor)
y_train=mnist_trainset.targets

#get and format the test set
mnist_testset = torchvision.datasets.MNIST(root='./data', train=False, download=True, transform=None)
x_test=mnist_testset.data.type(torch.DoubleTensor)
y_test=mnist_testset.targets
```



2 Etude d'un exemple — Détail du code d'apprentissage

```
class MyModel(nn.Module):
    """
    Basic fully connected neural-network
    """

    def __init__(self):
        hidden1      = 100
        hidden2      = 100
        super(MyModel, self).__init__()
        self.hidden1 = nn.Linear(784, hidden1)
        self.hidden2 = nn.Linear(hidden1, hidden2)
        self.hidden3 = nn.Linear(hidden2, 10)

    def forward(self, x):
        x = x.view(-1,784)  #flatten the images before using fully-connected layers
        x = self.hidden1(x)
        x = F.relu(x)
        x = self.hidden2(x)
        x = F.relu(x)
        x = self.hidden3(x)
        x = F.softmax(x, dim=0)
        return x

model = MyModel().to(DEVICE)
```

2 Etude d'un exemple — Détail du code d'apprentissage

```
def fit(model,X_train,Y_train,X_test,Y_test, EPOCHS = 5, BATCH_SIZE = 32):
    loss = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(),lr=1e-3) #lr is the learning rate
    model.train()

    n=X_train.shape[0] #number of observations in the training data

    #stochastic gradient descent
    for epoch in range(EPOCHS):
        batch_start=0
        epoch_shuffler=np.arange(n)
        np.random.shuffle(epoch_shuffler) #remark that 'utilsData.DataLoader' could be used instead

        while batch_start+BATCH_SIZE < n:
            #get mini-batch observation
            mini_batch_observations = epoch_shuffler[batch_start:batch_start+BATCH_SIZE]
            var_X_batch = X_train[mini_batch_observations,:,:].float().to(DEVICE)
            var_Y_batch = Y_train[mini_batch_observations]

            #gradient descent step
            optimizer.zero_grad()          #set the parameters gradients to 0
            Y_pred_batch = model(var_X_batch) #predict y with the current NN parameters
            curr_loss = loss(Y_pred_batch.to('cpu'), var_Y_batch) #compute the current loss
            curr_loss.backward()           #compute the loss gradient
            optimizer.step()              #update the NN parameters

            #prepare the next mini-batch of the epoch
            batch_start+=BATCH_SIZE
```

2 Etude d'un exemple — Détail du code d'apprentissage

```
def fit(model,X_train,Y_train,X_test,Y_test, EPOCHS = 5, BATCH_SIZE = 32):
    loss = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-3) #lr is the learning rate
    model.train()

    n=X_train.shape[0] #number of observations in the training data

    #stochastic gradient descent
    for epoch in range(EPOCHS):
        batch_start=0
        epoch_shuffler=np.arange(n)
        np.random.shuffle(epoch_shuffler) #remark that 'utilsData.DataLoader' could be used instead

        while batch_start+BATCH_SIZE < n:
            #get mini-batch observation
            mini_batch_observations = epoch_shuffler[batch_start:batch_start+BATCH_SIZE]
            var_X_batch = X_train[mini_batch_observations,:,:].float().to(DEVICE)
            var_Y_batch = Y_train[mini_batch_observations]

            #gradient descent step
            optimizer.zero_grad()                      #set the parameters gradients to 0
            Y_pred_batch = model(var_X_batch)           #predict y with the current NN parameters
            curr_loss = loss(Y_pred_batch.to('cpu'), var_Y_batch) #compute the current loss
            curr_loss.backward()                        #compute the loss gradient
            optimizer.step()                           #update the NN parameters

            #prepare the next mini-batch of the epoch
            batch_start+=BATCH_SIZE
```

- Définition de la fonction *loss*
- Définition de la *stratégie d'optimisation* (remarque : accède aux paramètres du modèle)

2 Etude d'un exemple — Détail du code d'apprentissage

```
def fit(model,X_train,Y_train,X_test,Y_test, EPOCHS = 5, BATCH_SIZE = 32):
    loss = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-3) #lr is the learning rate
    model.train()

    n=X_train.shape[0] #number of observations in the training data

    #stochastic gradient descent
    for epoch in range(EPOCHS):
        batch_start=0
        epoch_shuffler=np.arange(n)
        np.random.shuffle(epoch_shuffler) #remark that 'utilsData.DataLoader' could be used instead

        while batch_start+BATCH_SIZE < n:
            #get mini-batch observation
            mini_batch_observations = epoch_shuffler[batch_start:batch_start+BATCH_SIZE]
            var_X_batch = X_train[mini_batch_observations,:,:].float().to(DEVICE)
            var_Y_batch = Y_train[mini_batch_observations]

            #gradient descent step
            optimizer.zero_grad()                      #set the parameters gradients to 0
            Y_pred_batch = model(var_X_batch)           #predict y with the current NN parameters
            curr_loss = loss(Y_pred_batch.to('cpu'), var_Y_batch) #compute the current loss
            curr_loss.backward()                        #compute the loss gradient
            optimizer.step()                           #update the NN parameters

            #prepare the next mini-batch of the epoch
            batch_start+=BATCH_SIZE
```

Précise que l'on va entraîner le modèle

- Il existe le mode **train()** et le mode **eval()**
- **self.training** est **True** par défaut

2 Etude d'un exemple — Détail du code d'apprentissage

```
def fit(model,X_train,Y_train,X_test,Y_test, EPOCHS = 5, BATCH_SIZE = 32):
    loss = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-3) #lr is the learning rate
    model.train()

    n=X_train.shape[0] #number of observations in the training data

    #stochastic gradient descent
    for epoch in range(EPOCHS):
        batch_start=0
        epoch_shuffler=np.arange(n)
        np.random.shuffle(epoch_shuffler) #remark that 'utilsData.DataLoader' could be used instead

        while batch_start+BATCH_SIZE < n:
            #get mini-batch observation
            mini_batch_observations = epoch_shuffler[batch_start:batch_start+BATCH_SIZE]
            var_X_batch = X_train[mini_batch_observations,:,:].float().to(DEVICE)
            var_Y_batch = Y_train[mini_batch_observations]

            #gradient descent step
            optimizer.zero_grad()                      #set the parameters gradients to 0
            Y_pred_batch = model(var_X_batch)           #predict y with the current NN parameters
            curr_loss = loss(Y_pred_batch.to('cpu'), var_Y_batch) #compute the current loss
            curr_loss.backward()                        #compute the loss gradient
            optimizer.step()                           #update the NN parameters

            #prepare the next mini-batch of the epoch
            batch_start+=BATCH_SIZE
```

Gestion des *epochs* et du choix des observations dans chaque *mini-batch*

2 Etude d'un exemple — Détail du code d'apprentissage

```
def fit(model,X_train,Y_train,X_test,Y_test, EPOCHS = 5, BATCH_SIZE = 32):
    loss = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(),lr=1e-3) #lr is the learning rate
    model.train()

    n=X_train.shape[0] #number of observations in the training data

    #stochastic gradient descent
    for epoch in range(EPOCHS):
        batch_start=0
        epoch_shuffler=np.arange(n)
        np.random.shuffle(epoch_shuffler) #remark that 'utilsData.DataLoader' could be used instead

        while batch_start+BATCH_SIZE < n:
            #get mini-batch observation
            mini_batch_observations = epoch_shuffler[batch_start:batch_start+BATCH_SIZE]
            var_X_batch = X_train[mini_batch_observations,:,:].float().to(DEVICE)
            var_Y_batch = Y_train[mini_batch_observations]

            #gradient descent step
            optimizer.zero_grad()                      #set the parameters gradients to 0
            Y_pred_batch = model(var_X_batch)           #predict y with the current NN parameters
            curr_loss = loss(Y_pred_batch.to('cpu'), var_Y_batch) #compute the current loss
            curr_loss.backward()                        #compute the loss gradient
            optimizer.step()                           #update the NN parameters

            #prepare the next mini-batch of the epoch
            batch_start+=BATCH_SIZE
```

Récupération des observations nécessaires pour le **mini-batch** de l'itération courante

2 Etude d'un exemple — Détail du code d'apprentissage

```
def fit(model,X_train,Y_train,X_test,Y_test, EPOCHS = 5, BATCH_SIZE = 32):
    loss = nn.CrossEntropyLoss()
    optimizer = torch.optim.Adam(model.parameters(), lr=1e-3) #lr is the learning rate
    model.train()

    n=X_train.shape[0] #number of observations in the training data

    #stochastic gradient descent
    for epoch in range(EPOCHS):
        batch_start=0
        epoch_shuffler=np.arange(n)
        np.random.shuffle(epoch_shuffler) #remark that 'utilsData.DataLoader' could be used instead

        while batch_start+BATCH_SIZE < n:
            #get mini-batch observation
            mini_batch_observations = epoch_shuffler[batch_start:batch_start+BATCH_SIZE]
            var_X_batch = X_train[mini_batch_observations,:,:].float().to(DEVICE)
            var_Y_batch = Y_train[mini_batch_observations]

            #gradient descent step
            optimizer.zero_grad()                      #set the parameters gradients to 0
            Y_pred_batch = model(var_X_batch)           #predict y with the current NN parameters
            curr_loss = loss(Y_pred_batch.to('cpu'), var_Y_batch) #compute the current loss
            curr_loss.backward()                        #compute the loss gradient
            optimizer.step()                           #update the NN parameters

            #prepare the next mini-batch of the epoch
            batch_start+=BATCH_SIZE
```

- Mise à zéro des gradients de paramètres du modèle
- Prédiction des sorties avec les paramètres du modèle courant (\rightarrow calcul de $\{h_{\Theta}(x_i)\}_{i \in B}$)
- Calcul du loss pour toutes les observations du mini-batch
- Calcul des gradients (backpropagation) puis mise à jour des paramètres en fonction des gradients



Simple DNN classification using PyTorch

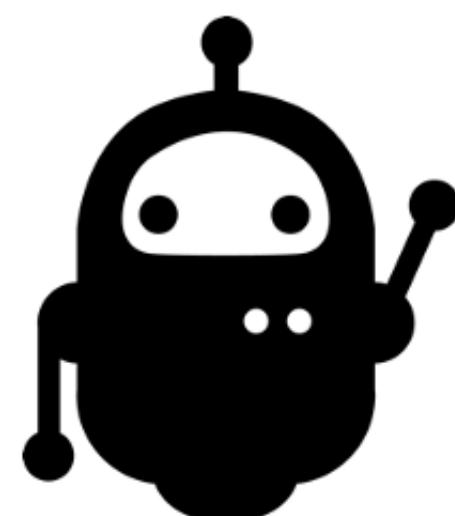
Notebook : [01-DNN-MNIST_PyTorch-2022.ipynb](#)

Objective :

Donner un exemple d'utilisation élémentaire
de PyTorch

Dataset :

MNIST que l'on ne présente plus :-)



3 : Aller plus loin ...



3.1 Aller plus loin – Charger et adapter une architecture

→ Poids tirés au hasard

```
import torchvision.models as models
resnet18 = models.resnet18()
alexnet = models.alexnet()
vgg16 = models.vgg16()
squeezenet = models.squeezenet1_0()
densenet = models.densenet161()
inception = models.inception_v3()
googlenet = models.googlenet()
shufflenet = models.shufflenet_v2_x1_0()
mobilenet = models.mobilenet_v2()
resnext50_32x4d = models.resnext50_32x4d()
wide_resnet50_2 = models.wide_resnet50_2()
mnasnet = models.mnasnet1_0()
```

3.1 Aller plus loin – Charger et adapter une architecture

→ Poids pré-entraînés

```
import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)
alexnet = models.alexnet(pretrained=True)
squeezenet = models.squeeze1_0(pretrained=True)
vgg16 = models.vgg16(pretrained=True)
densenet = models.densenet161(pretrained=True)
inception = models.inception_v3(pretrained=True)
googlenet = models.googlenet(pretrained=True)
shufflenet = models.shufflenet_v2_x1_0(pretrained=True)
mobilenet = models.mobilenet_v2(pretrained=True)
resnext50_32x4d = models.resnext50_32x4d(pretrained=True)
wide_resnet50_2 = models.wide_resnet50_2(pretrained=True)
mnasnet = models.mnasnet1_0(pretrained=True)
```

3.1 Aller plus loin – Charger et adapter une architecture

```
import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)

print(resnet18)

ResNet(
    (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    (layer1): Sequential(
        (0): BasicBlock(
            (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
        :
        :
        :
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    :
    (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
    (fc): Linear(in_features=512, out_features=1000, bias=True)
)
```

← Entrée : Image 2D à 3 canaux → Sortie : Vecteur de taille 1000

Besoin d'adapter la première et dernière couche si par exemple (MNIST)

- Les entrées n'ont qu'un canal
- Les sorties sont en dimension 10

3.1 Aller plus loin – Charger et adapter une architecture

```
import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)

print(resnet18)

ResNet(
    (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    (layer1): Sequential(
        (0): BasicBlock(
            (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
        :
        :
        :
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    :
    (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
    (fc): Linear(in_features=512, out_features=10, bias=True)
)
```

← Entrée : Image 2D à 3 canaux

→ Sortie : Vecteur de taille 10

Modification des sorties (basique)

```
import torch.nn as nn

output_classes=10
resnet18.fc = nn.Linear(512, output_classes)
```

3.1 Aller plus loin – Charger et adapter une architecture

```
import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)

print(resnet18)

ResNet(
    (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    (layer1): Sequential(
        (0): BasicBlock(
            (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
        :
        :
        :
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
    (fc): Sequential(
        (fc1): Linear(in_features=512, out_features=128, bias=True)
        (relu): ReLU()
        (fc2): Linear(in_features=128, out_features=10, bias=True)
        (output): Sigmoid()
    )
)
)
```

← Entrée : Image 2D à 3 canaux

→ Sortie : Vecteur de taille 10

Modification des sorties (avancé) :

```
import torch.nn as nn
from collections import OrderedDict

output_classes=10

resnet18.fc = nn.Sequential(OrderedDict([
    ('fc1', nn.Linear(512, 128)), ('relu', nn.ReLU()),
    ('fc2', nn.Linear(128, output_classes)), ('output', nn.Sigmoid())
]))
```

3.1 Aller plus loin – Charger et adapter une architecture

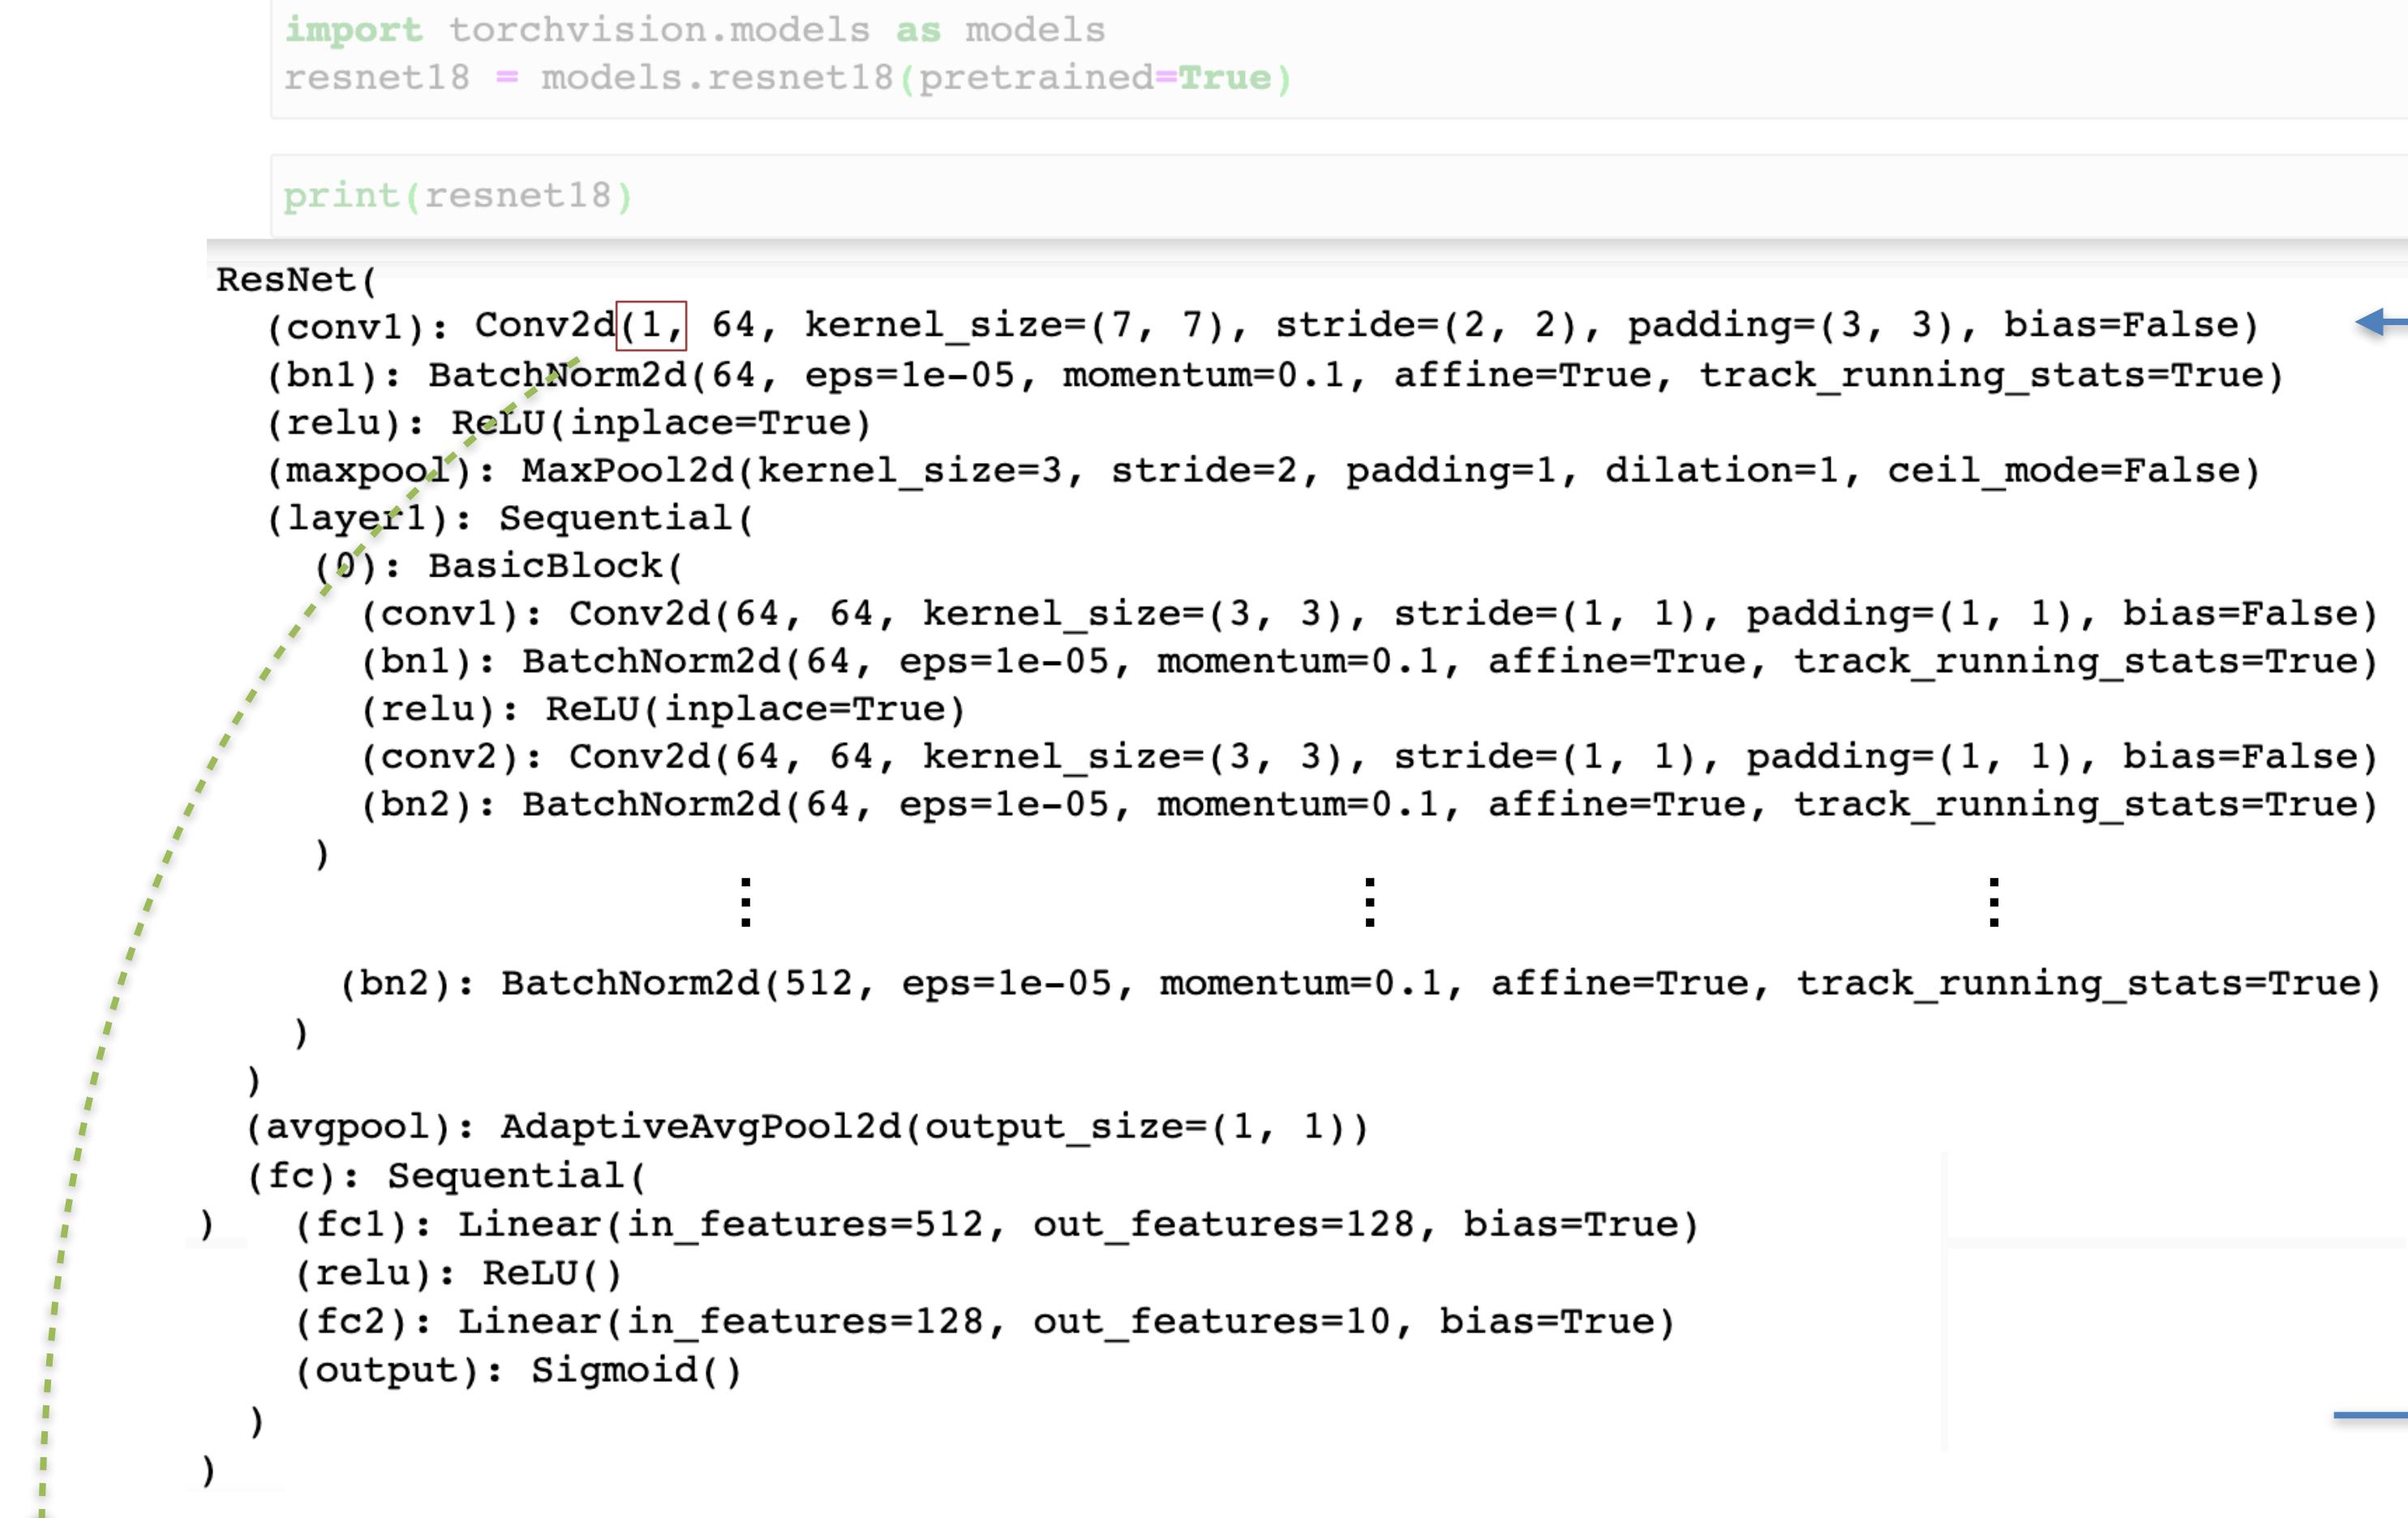
```
import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)

print(resnet18)

ResNet(
    (conv1): Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    (layer1): Sequential(
        (0): BasicBlock(
            (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace=True)
            (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
        :
        :
        :
        (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
    (fc): Sequential(
        (fc1): Linear(in_features=512, out_features=128, bias=True)
        (relu): ReLU()
        (fc2): Linear(in_features=128, out_features=10, bias=True)
        (output): Sigmoid()
    )
)
)
```

← Entrée : Image 2D à 1 canal

→ Sortie : Vecteur de taille 10



Modification des entrées :

```
resnet18.conv1=nn.Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
```

3.1 Aller plus loin – Charger et adapter une architecture

```
import torchvision.models as models
resnet18 = models.resnet18(pretrained=True)

print(resnet18)

ResNet(
    (conv1): Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
    (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (relu): ReLU(inplace=True)
    (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
    (layer1): Identity() ← Entrée : Image 2D à 1 canal
    (layer2): Sequential(
        (0): BasicBlock(
            (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
            (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            (relu): ReLU(inplace=True)
            (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
            (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
            :
            :
            (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        )
    )
    (avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
    (fc): Sequential(
        (fc1): Linear(in_features=512, out_features=128, bias=True)
        (relu): ReLU()
        (fc2): Linear(in_features=128, out_features=10, bias=True)
        (output): Sigmoid()
    )
)
)
```

→ Sortie : Vecteur de taille 10

Simplification du réseau :

```
class Identity(nn.Module):
    def __init__(self):
        super(Identity, self).__init__()
    def forward(self, x):
        return x
resnet18.layer1=Identity()
```

3.2 Aller plus loin — Récupérer les paramètres du réseau

Poids du réseau sous la forme d'un dictionnaire : « resnet18.state_dict() »

```
for param_tensor in resnet18.state_dict():
    print(param_tensor, "\t", resnet18.state_dict()[param_tensor].size())

conv1.weight      torch.Size([64, 1, 7, 7])
bn1.weight       torch.Size([64])
bn1.bias         torch.Size([64])
bn1.running_mean      torch.Size([64])
bn1.running_var       torch.Size([64])
bn1.num_batches_tracked  torch.Size([])
layer2.0.conv1.weight  torch.Size([128, 64, 3, 3])
layer2.0.bn1.weight   torch.Size([128])
layer2.0.bn1.bias     torch.Size([128])

:
:
:

layer4.1.conv2.weight  torch.Size([512, 512, 3, 3])
layer4.1.bn2.weight    torch.Size([512])
layer4.1.bn2.bias      torch.Size([512])
layer4.1.bn2.running_mean  torch.Size([512])
layer4.1.bn2.running_var  torch.Size([512])
layer4.1.bn2.num_batches_tracked  torch.Size([])
fc.fc1.weight      torch.Size([128, 512])
fc.fc1.bias        torch.Size([128])
fc.fc2.weight      torch.Size([10, 128])
fc.fc2.bias        torch.Size([10])
```

3.2 Aller plus loin – Récupérer les paramètres du réseau

Poids du réseau sous la forme d'un dictionnaire : « resnet18.state_dict() »

```
resnet_dict=resnet18.state_dict()  
  
print(resnet_dict['layer2.0.conv2.weight'])
```

```
tensor([[[[-0.0074, -0.0098,  0.0028],  
         [-0.0108,  0.0258,  0.0455],  
         [-0.0272,  0.0053,  0.0132]],  
  
        [[ 0.0354,  0.0251,  0.0078],  
         [ 0.0040,  0.0199,  0.0274],  
         [ 0.0353,  0.0355,  0.0133]],  
  
        [[ 0.0193, -0.0213, -0.0362],  
         [-0.0196, -0.0189, -0.0595],  
         [-0.0218, -0.0077,  0.0039]],  
  
        ...,  
  
        [[-0.0068,  0.0108, -0.0037],  
         [ 0.0135,  0.0114, -0.0013],  
         [ 0.0081,  0.0002,  0.0006]],  
  
        [[ 0.0077,  0.0077,  0.0044],  
         [ 0.0100,  0.0117,  0.0061]]])
```

```
print(resnet_dict['fc.fc2.bias'])
```

```
tensor([ 0.0310, -0.0075, -0.0521,  0.0285,  0.0192, -0.08  
11,  0.0776,  0.0501,  
       -0.0611,  0.0259])
```

3.3 Aller plus loin — N'apprendre que les paramètres d'une sous partie du réseau

Fixer (freezer) l'apprentissage de paramètres spécifiques

```
for param in resnet18.conv1.parameters():
    param.requires_grad = False

for name, param in resnet18.named_parameters():
    print(name, param.requires_grad)

conv1.weight False
bn1.weight True
bn1.bias True
layer2.0.conv1.weight True
layer2.0.bn1.weight True
layer2.0.bn1.bias True
layer2.0.conv2.weight True
layer2.0.bn2.weight True
layer2.0.bn2.bias True
layer2.0.downsample.0.weight True

    :       :       :

layer4.1.bn2.weight True
layer4.1.bn2.bias True
fc.fc1.weight True
fc.fc1.bias True
fc.fc2.weight True
fc.fc2.bias True
```

3.3 Aller plus loin — N'apprendre que les paramètres d'une sous partie du réseau

Fixer (freezer) l'apprentissage de paramètres spécifiques

```
for name, param in resnet18.named_parameters():
    if name[:3] != 'fc.':
        param.requires_grad = False

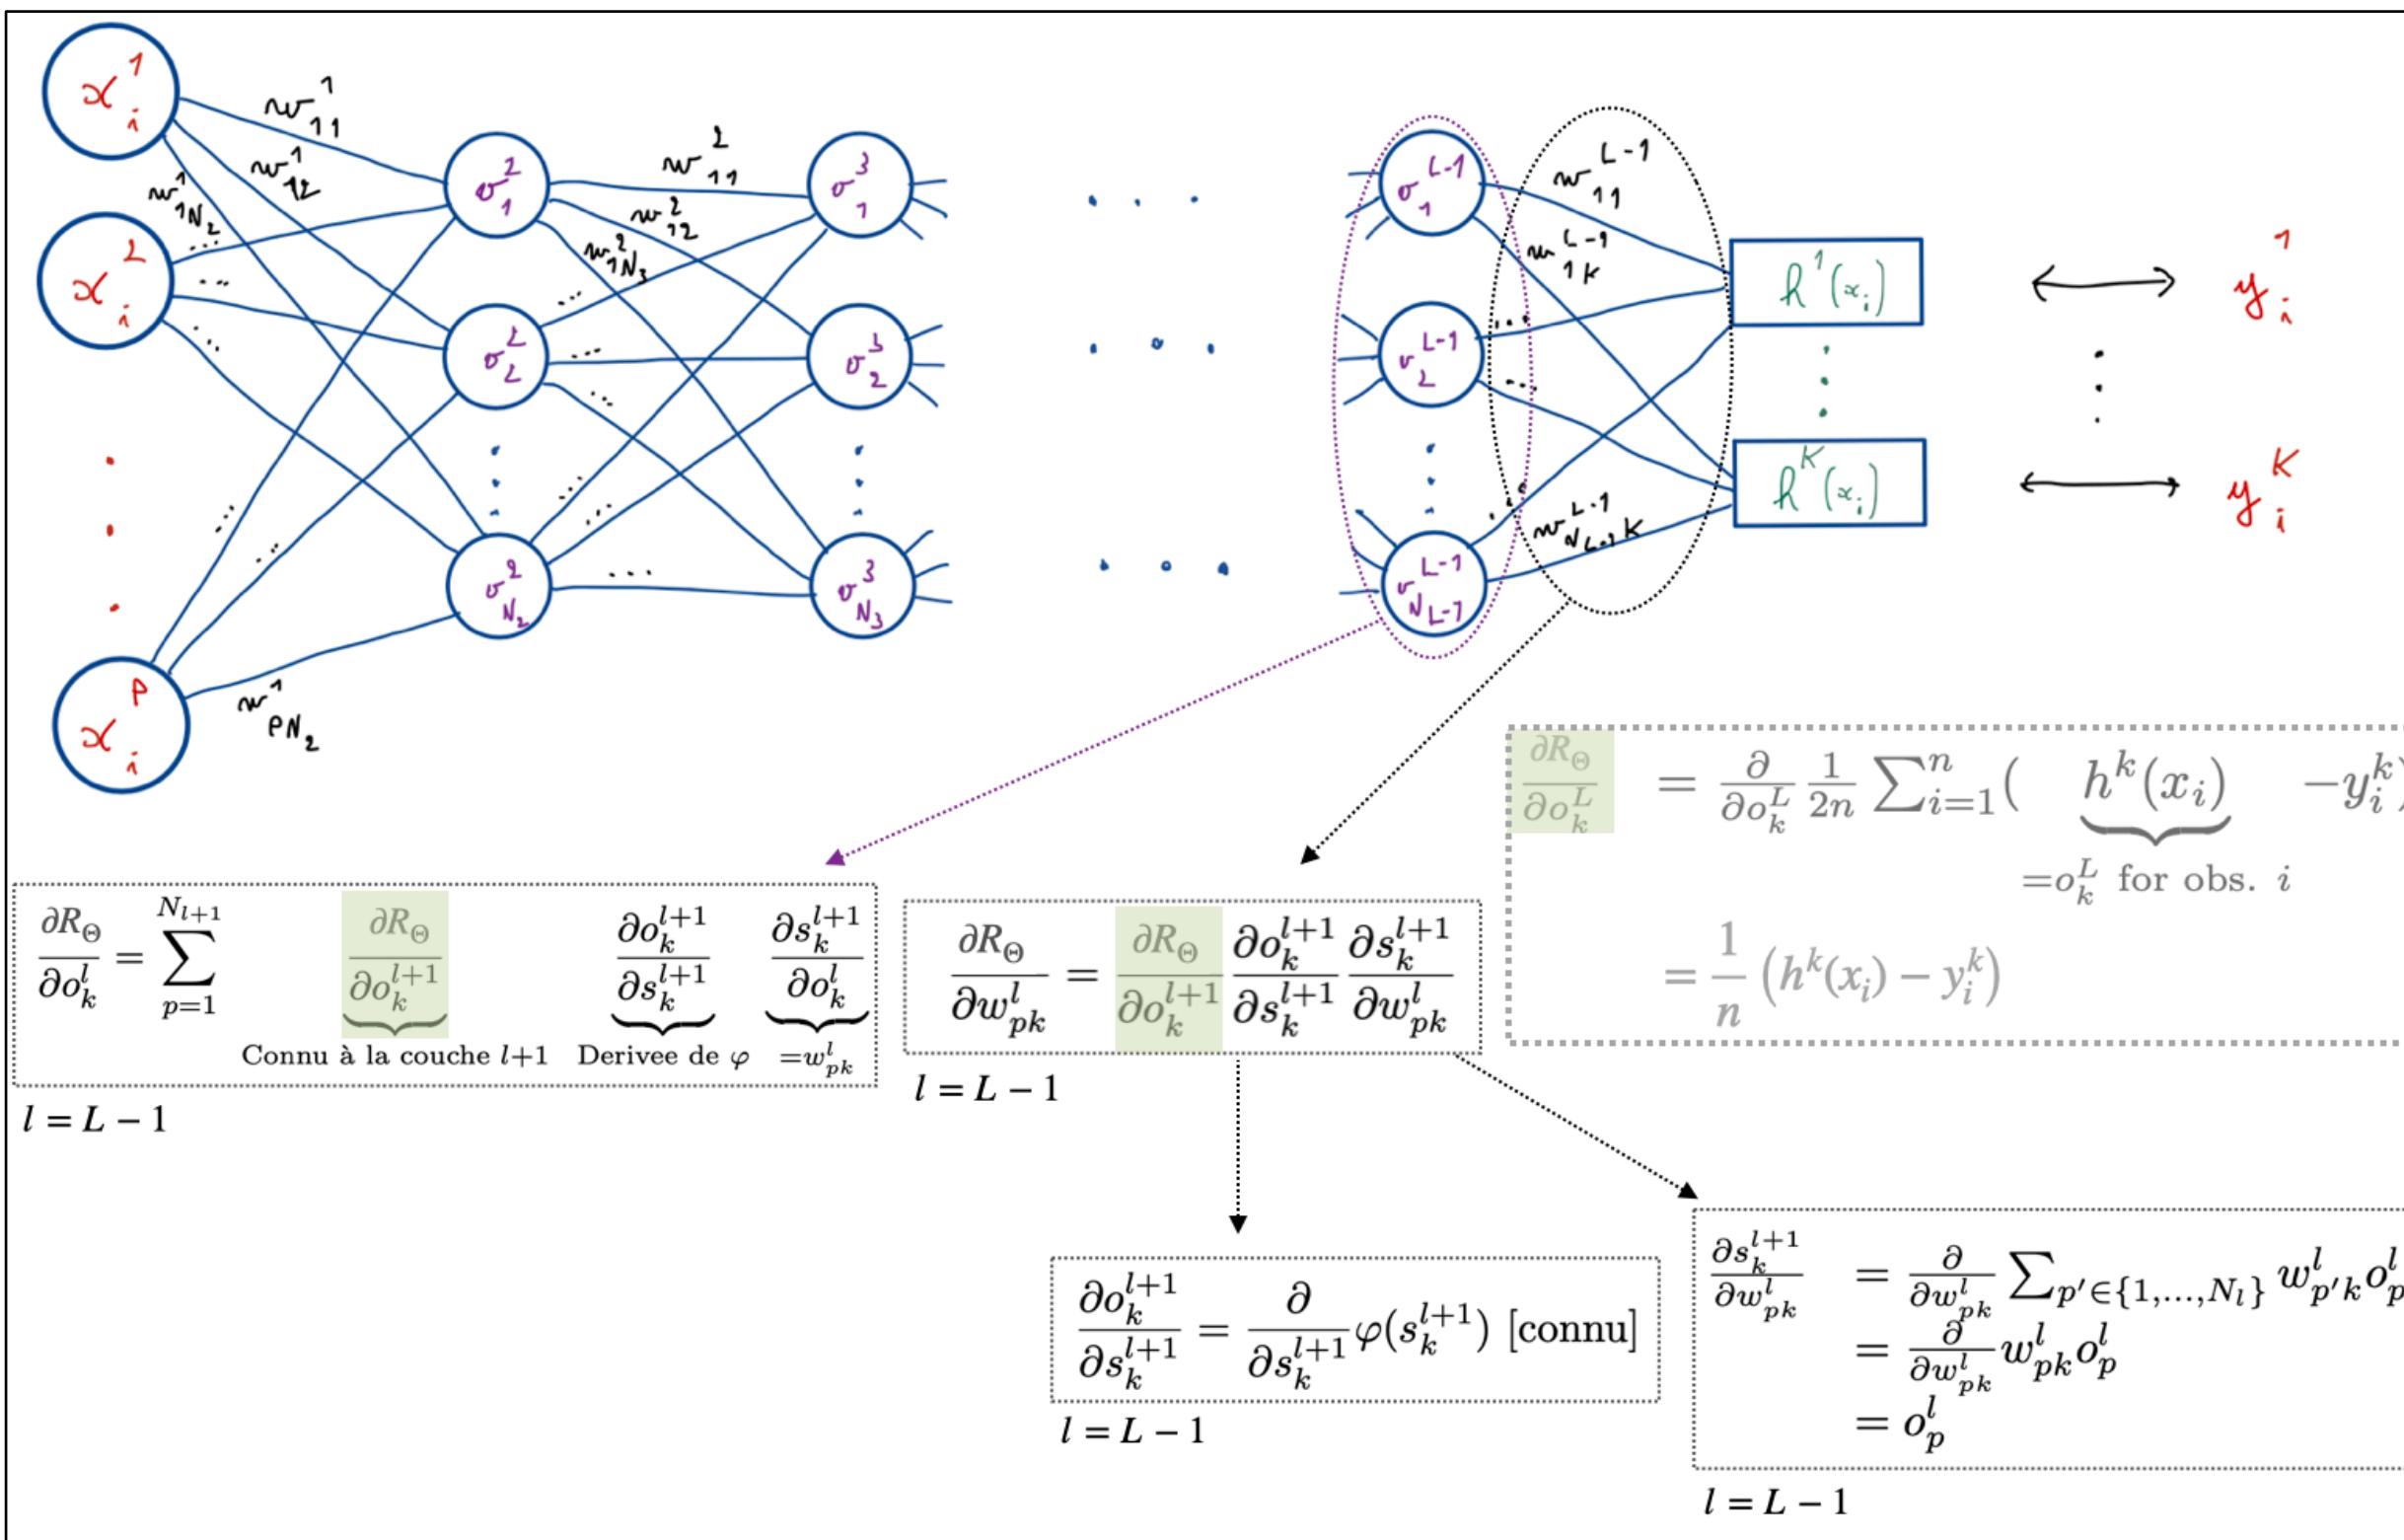
for name, param in resnet18.named_parameters():
    print(name, param.requires_grad)
```

```
conv1.weight False
bn1.weight False
bn1.bias False
layer2.0.conv1.weight False
layer2.0.bn1.weight False
layer2.0.bn1.bias False
layer2.0.conv2.weight False
layer2.0.bn2.weight False
layer2.0.bn2.bias False
```

```
:      :      :
```

```
layer4.1.bn2.weight False
layer4.1.bn2.bias False
fc.fc1.weight True
fc.fc1.bias True
fc.fc2.weight True
fc.fc2.bias True
```

3.4 Aller plus loin – Définir ses propres couches



On défini : $\text{out} = f(\text{in}, \text{param}_1, \dots, \text{param}_p)$

On connaît : $\frac{\partial R}{\partial \text{out}}$

→

$$\frac{\partial R}{\partial \text{in}} = \frac{\partial R}{\partial \text{out}} \frac{\partial \text{out}}{\partial \text{in}} = \frac{\partial R}{\partial \text{out}} \frac{\partial f(\text{in}, \text{param}_1, \dots, \text{param}_p)}{\partial \text{in}}$$

$$\dots$$

$$\frac{\partial R}{\partial \text{param}_p} = \frac{\partial R}{\partial \text{out}} \frac{\partial \text{out}}{\partial \text{param}_p} = \frac{\partial R}{\partial \text{out}} \frac{\partial f(\text{in}, \text{param}_1, \dots, \text{param}_p)}{\partial \text{param}_p}$$

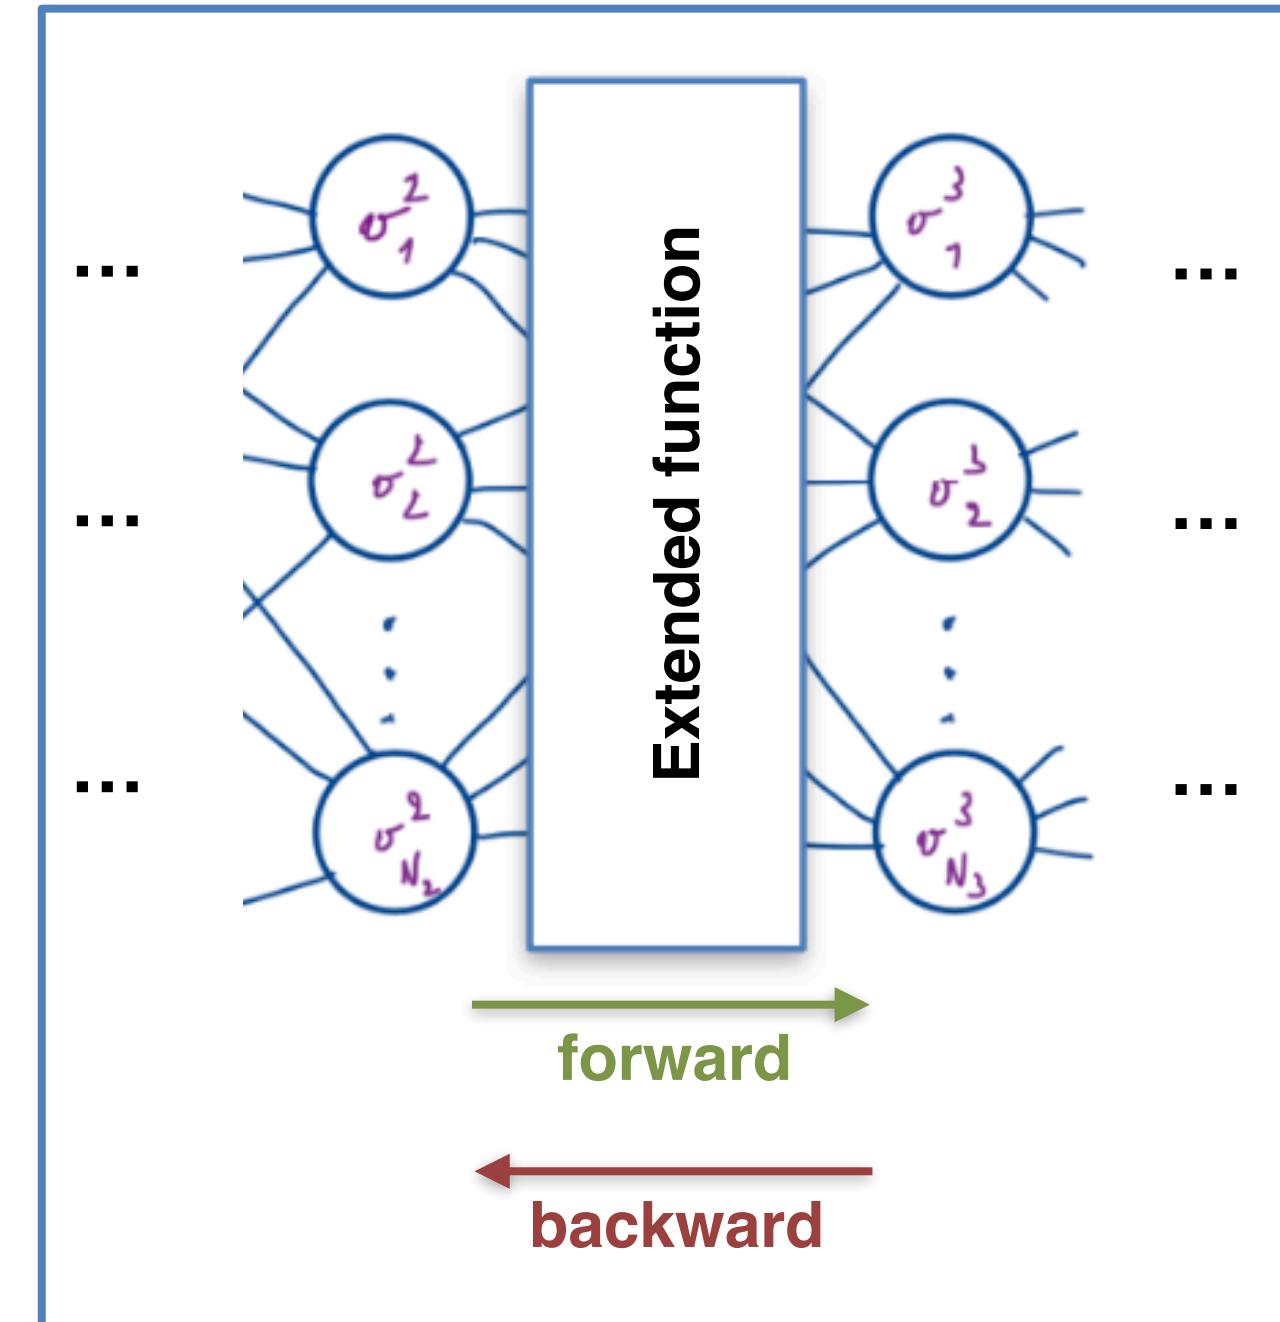
3.4 Aller plus loin – Définir ses propres couches

```
from torch.autograd import Function

class MyReLU(Function):
    @staticmethod
    def forward(ctx, input):
        ctx.save_for_backward(input)
        return input.clamp(min=0)

    @staticmethod
    def backward(ctx, grad_output):
        input, = ctx.saved_tensors
        grad_input = grad_output.clone()
        grad_input[input < 0] = 0
        return grad_input

relu = MyReLU.apply
```



<https://pytorch.org/docs/stable/notes/extending.html>

https://pytorch.org/tutorials/beginner/examples_autograd/two_layer_net_custom_function.html

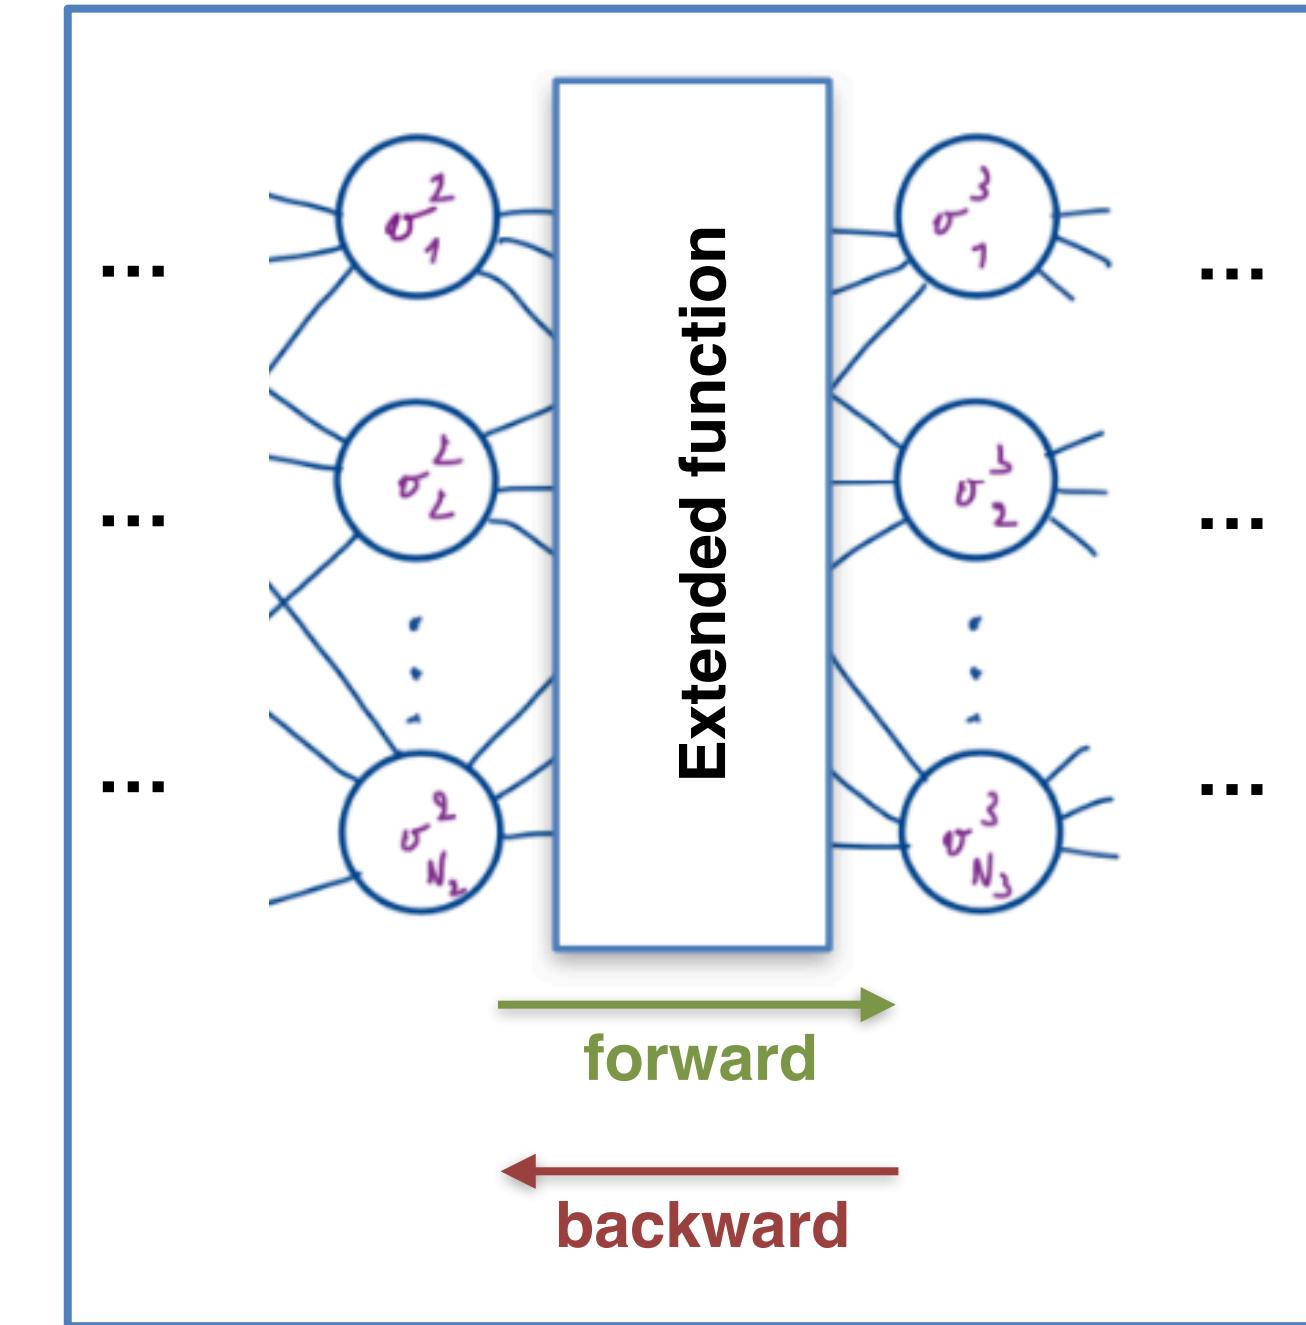
3.4 Aller plus loin – Définir ses propres couches

```
from torch.autograd import Function

class MulConstant(Function):
    @staticmethod
    def forward(ctx, tensor, constant):
        ctx.constant = constant
        return tensor * constant

    @staticmethod
    def backward(ctx, grad_output):
        return grad_output * ctx.constant, None

mulCst = MulConstant.apply
```



<https://pytorch.org/docs/stable/notes/extending.html>

https://pytorch.org/tutorials/beginner/examples_autograd/two_layer_net_custom_function.html

3.4 Aller plus loin – Définir ses propres couches

```
class LinearFunction(Function):

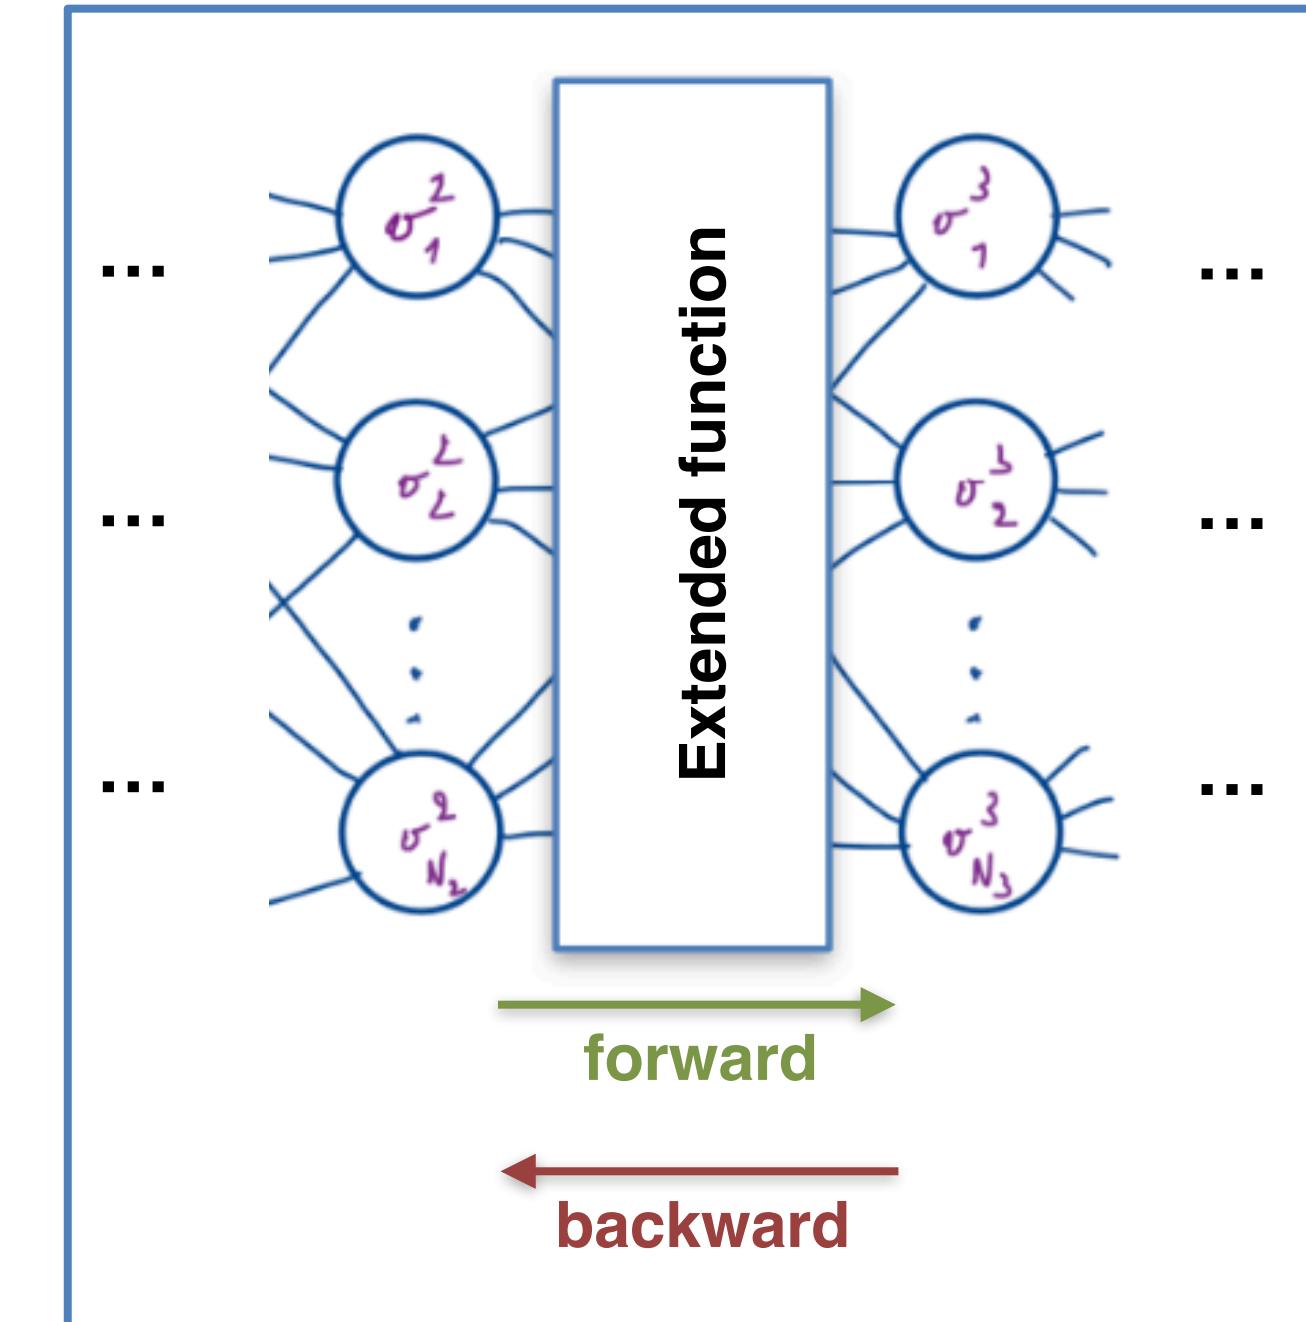
    @staticmethod
    def forward(ctx, input, weight, bias):
        ctx.save_for_backward(input, weight, bias)
        output = input.mm(weight.t())
        output += bias.unsqueeze(0).expand_as(output)
        return output

    @staticmethod
    def backward(ctx, grad_output):
        input, weight, bias = ctx.saved_tensors

        grad_input = grad_output.mm(weight)
        grad_weight = grad_output.t().mm(input)
        grad_bias = grad_output.sum(0)

        return grad_input, grad_weight, grad_bias

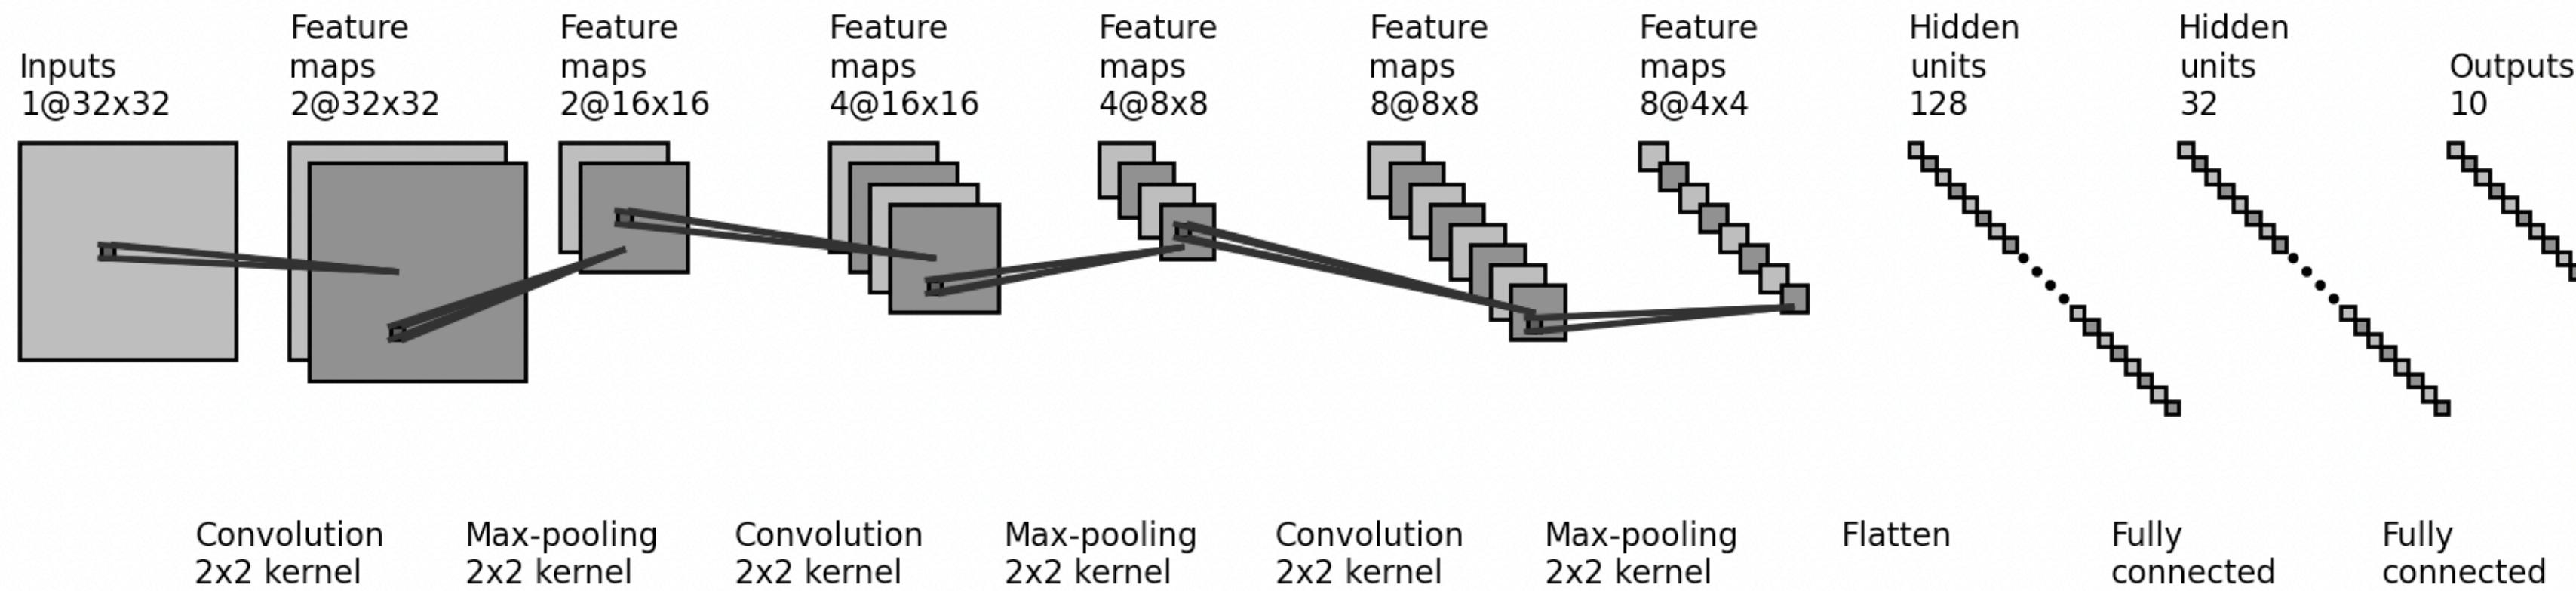
linear = LinearFunction.apply
```



<https://pytorch.org/docs/stable/notes/extending.html>

https://pytorch.org/tutorials/beginner/examples_autograd/two_layer_net_custom_function.html

3.5 Aller plus loin – Récupérer les valeurs des neurones

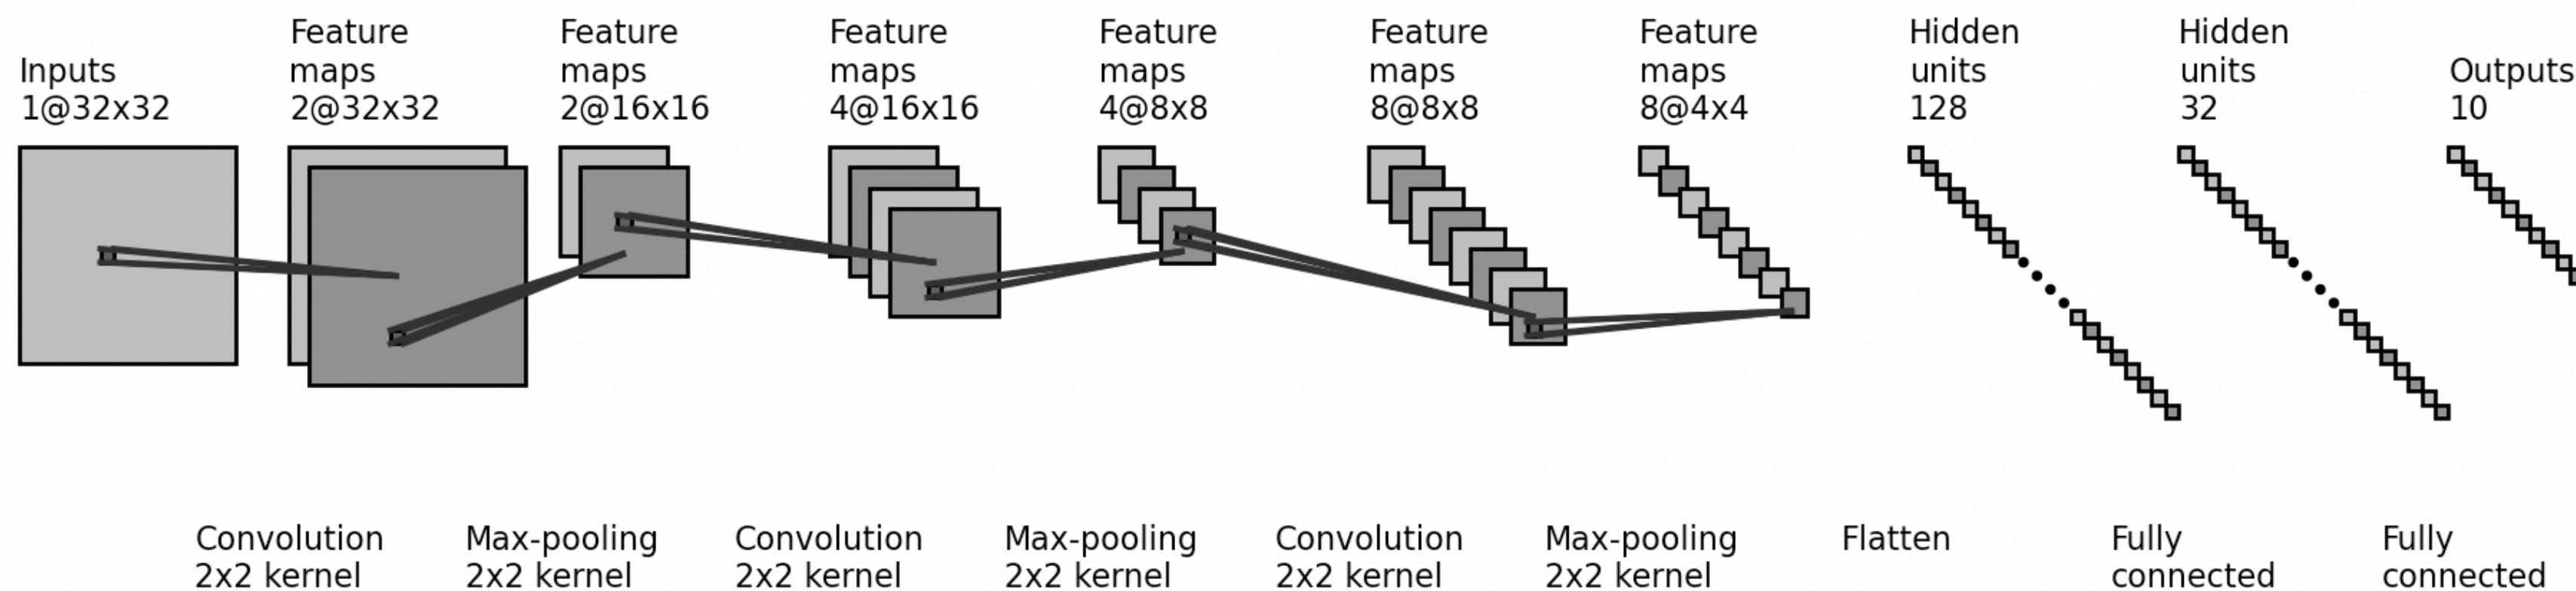


```
class basicCNN(nn.Module):
    def __init__(self):
        super(basicCNN, self).__init__()
        self.conv1 = nn.Conv2d(1, 2, kernel_size=2, stride=1, padding=1)
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv2 = nn.Conv2d(2, 4, kernel_size=2, stride=1, padding=1)
        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.conv3 = nn.Conv2d(4, 8, kernel_size=2, stride=1, padding=1)
        self.pool3 = nn.MaxPool2d(kernel_size=2, stride=2)
        self.fc1 = nn.Linear(8 * 4 * 4, 32)
        self.fc2 = nn.Linear(32, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool1(x)
        x = F.relu(self.conv2(x))
        x = self.pool2(x)
        x = F.relu(self.conv3(x))
        x = self.pool3(x)
        x = x.view(-1, 8*4*4)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```

```
basicCNN_model = basicCNN()
```

3.5 Aller plus loin — Récupérer les valeurs des neurones



```
#get 100 observations
var_X = Variable(torch_X_train[0:100,:,:,:].view(-1,1,32,32)).float()

#create a hook function that will get the transformed data in an hidden
#layer. These data will be stored in the 'activation' dictionary
activation = {}
def get_activation(name):
    def hook(model, input, output):
        activation[name] = output.detach()
    return hook

#Add hooks at the output of different parts of the model.
#-> The hook location is in "[HERE].register_forward_hook(...)"
#-> "get_activation('[HERE]') is the key where the data are saved
basicCNN_model.pool1.register_forward_hook(get_activation('pool1'))
basicCNN_model.pool2.register_forward_hook(get_activation('pool2'))
basicCNN_model.pool3.register_forward_hook(get_activation('pool3'))
basicCNN_model.fc1.register_forward_hook(get_activation('fc1'))

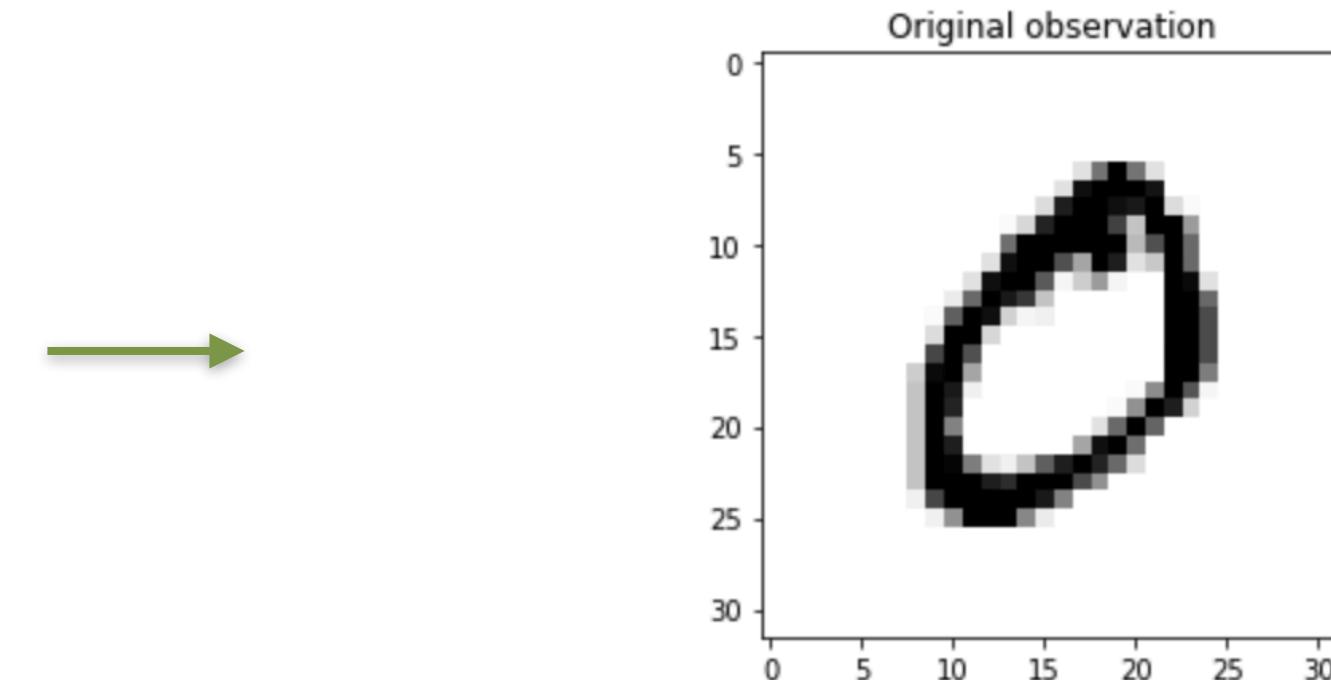
#propagate the model and get the hooks
output = basicCNN_model(var_X)
```

```
print(activation['pool1'].shape)
print(activation['pool2'].shape)
print(activation['pool3'].shape)
print(activation['fc1'].shape)
```

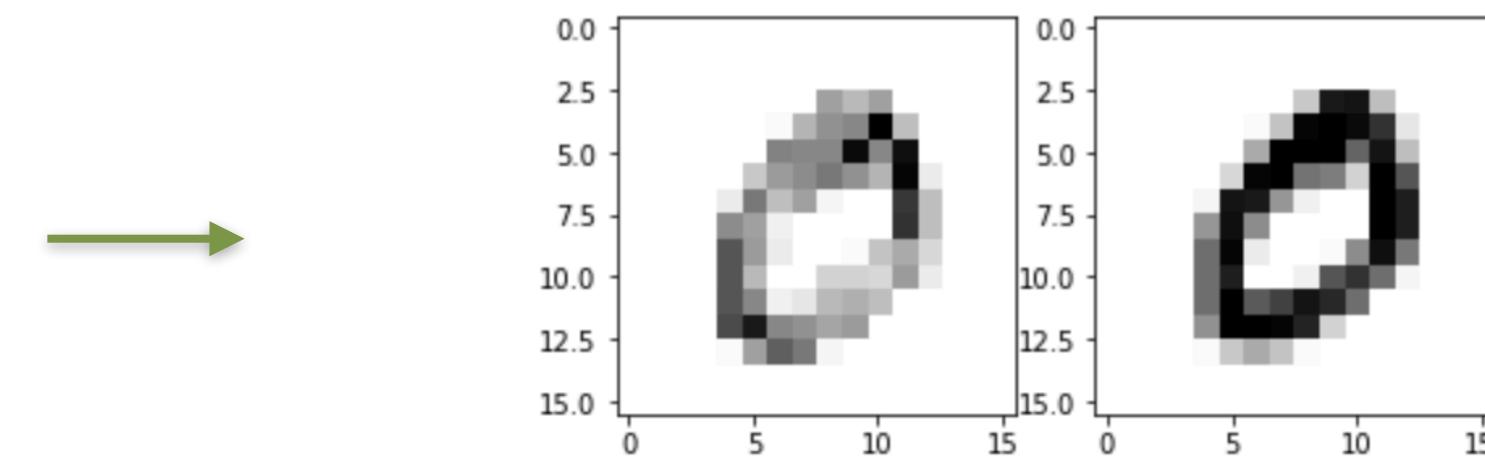
```
torch.Size([100, 2, 16, 16])
torch.Size([100, 4, 8, 8])
torch.Size([100, 8, 4, 4])
torch.Size([100, 32])
```

3.5 Aller plus loin – Récupérer les valeurs des neurones

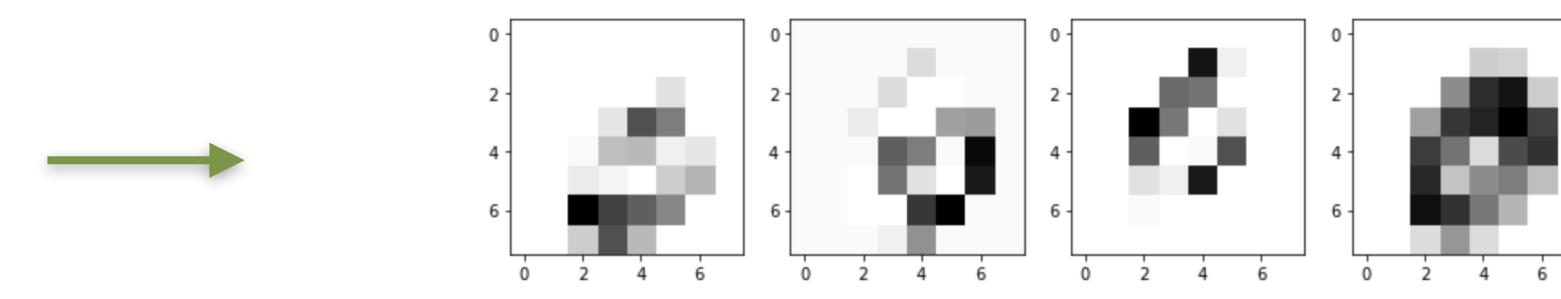
```
image_ID=1  
plt.imshow(torch_X_train[image_ID,:,::], cmap='Greys')  
plt.colorbar()  
plt.title('Original observation')  
plt.show()
```



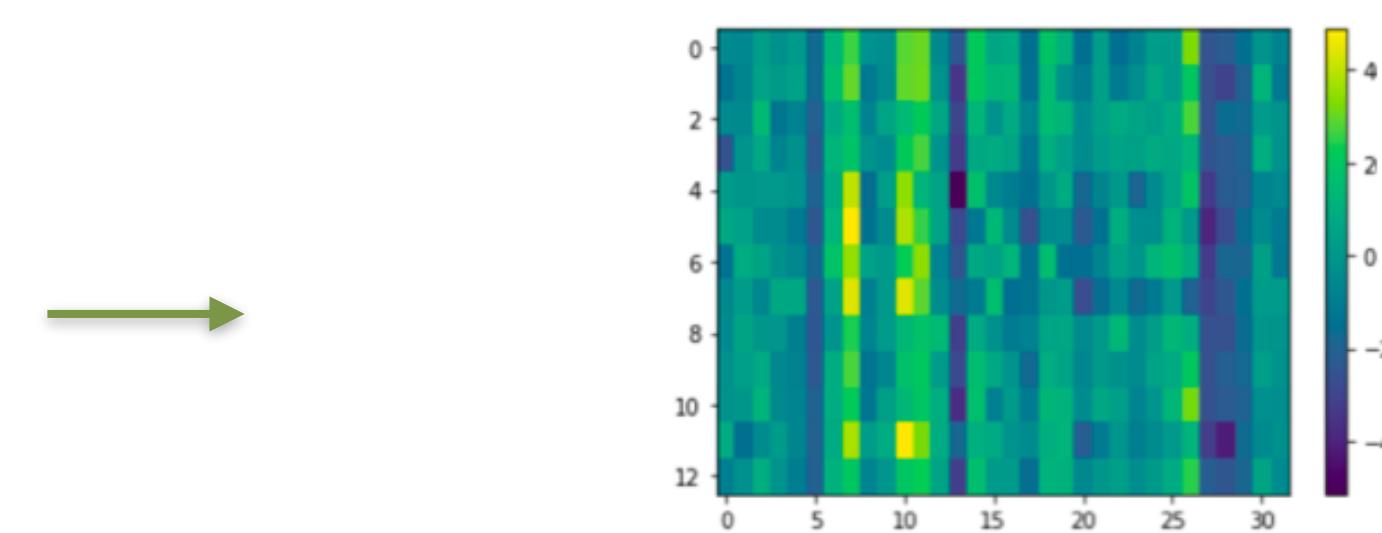
```
fig, (ax0, ax1) = plt.subplots(nrows=1, ncols=2, sharex=True)  
ax0.imshow(activation['pool1'][image_ID,0,:,:], cmap='Greys')  
ax1.imshow(activation['pool1'][image_ID,1,:,:], cmap='Greys')  
plt.show()
```



```
fig, (ax0, ax1, ax2, ax3) = plt.subplots(nrows=1, ncols=4, figsize=(12, 6))  
ax0.imshow(activation['pool2'][image_ID,0,:,:], cmap='Greys')  
ax1.imshow(activation['pool2'][image_ID,1,:,:], cmap='Greys')  
ax2.imshow(activation['pool2'][image_ID,2,:,:], cmap='Greys')  
ax3.imshow(activation['pool2'][image_ID,3,:,:], cmap='Greys')  
plt.show()
```



```
Obs_with_IDeq0=np.where(var_y[:].numpy()==0)[0]  
  
print('Latent space:')  
plt.imshow(activation['fc1'][Obs_with_IDeq0,:], aspect='auto')  
plt.colorbar()  
plt.show()
```

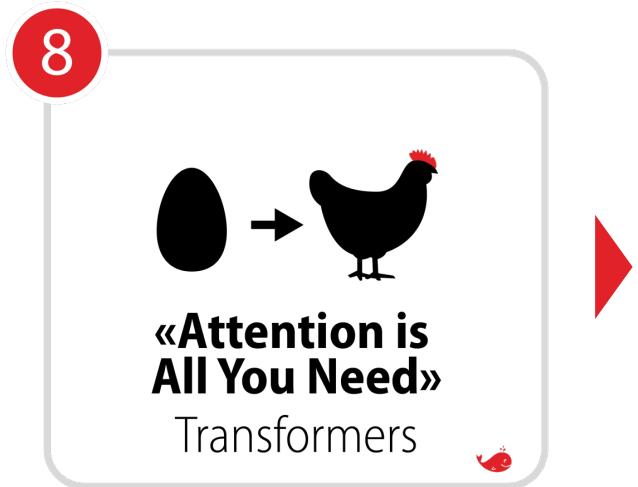




Little things and concepts to **keep in mind**

- PyTorch c'est :
 - puissant
 - flexible
 - basé sur l'objet `torch.tensor` et une série de fonctions
- L'utilisation de la rétropropagation peut y être totalement sous jacente...
- ... mais ce framework permet aussi d'accéder, de controller et de modifier des aspects fins des DNNs !

Next, on Fidle :



Jeudi 17 janvier, 14h

Séquence n°8

Attention Is All You Need", quand les Transformers changent la donne !

Utilisation et l'architecture classique des transformers ,
Principe du mécanisme d'attention et du multi-head attention
Architectures des transformers (auto-regressive, auto-encoding et encoder decoder)
Pré-entraînement (BERT et GPT)
Fine tuning
Utilisation des transformers dans les autres domaines
Durée : 2h

Next, on Fidle :

8



Jeudi 17 janvier, 14h

Séquence n° 8

**Attention Is All You Need", quand les
Transformers changent la donne !**



To be continued...

contact@fidle.cnrs.fr

FIDLE <https://fidle.cnrs.fr>

YouTube <https://fidle.cnrs.fr/youtube>



Attribution-NonCommercial-NoDerivatives 4.0 International (CC BY-NC-ND 4.0)
<https://creativecommons.org/licenses/by-nc-nd/4.0/>