



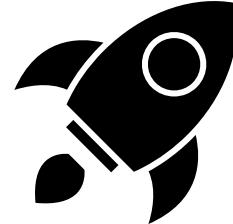
Formation

# Introduction au Deep Learning

Variational Autoencoder (VAE)  
ou comment jouer avec les espaces latents ;-)



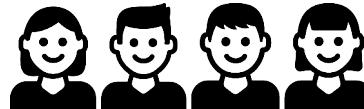
Annonce !



# « Journée Deep Learning pour la Science ! »

Mai 2023, Orsay

In the real world !



# Resources

<https://fidle.cnrs.fr>

Powered by CNRS CRIC, and UGA DGDSI  
of Grenoble, Thanks !



Course materials (pdf)



Practical work environment\*



Corrected notebooks



Videos (YouTube)

(\*) Procedure via Docket or pip  
Remember to get the latest version !



Attribution-NonCommercial-NoDerivatives 4.0 International (CC BY-NC-ND 4.0)  
<https://creativecommons.org/licenses/by-nc-nd/4.0/>

# Resources

You can also subscribe to :



<http://fidle.cnrs.fr/listeinfo>  
Fidle information list

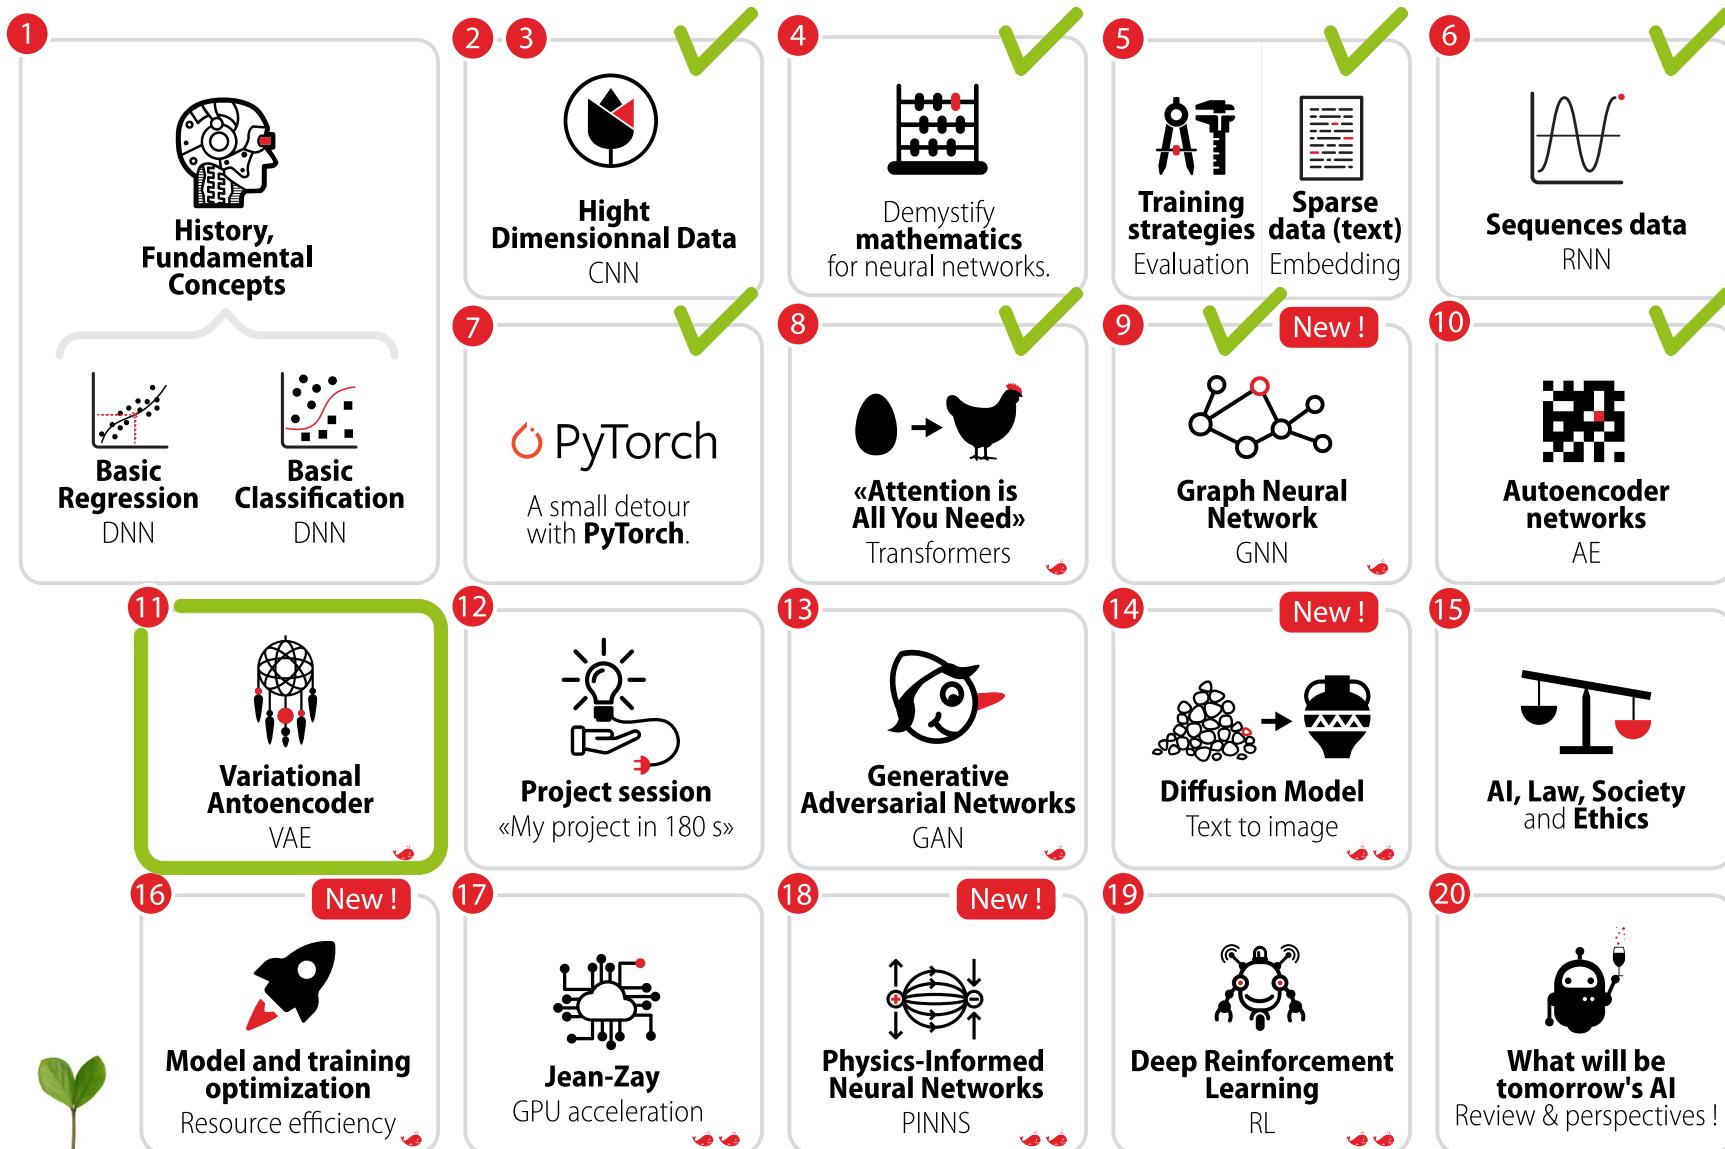


<https://listes.services.cnrs.fr/wws/info/devlog>  
List of ESR\* « Software developers » group

<https://listes.math.cnrs.fr/wws/info/calcul>  
List of ESR\* « Calcul » group

# Program

FIDDLE



SAISON  
22/23



11.1

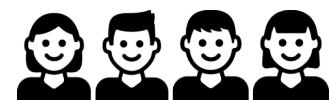
## Variational Autoencoder network

- Concepts and architecture
- Gaussian / probabilistic projections
- Kullback-Leibler divergence
- Morphing in latent space

11.2

## Example 1 : VAE1/3

- Deeper in Keras !
- VAE using Functional API (MNIST)
- VAE using Model subclass (MNIST)



11



Variational  
Autoencoder

VAE

11.1

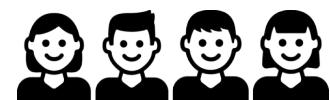
## Variational Autoencoder network

- Concepts and architecture
- Gaussian / probabilistic projections
- Kullback-Leibler divergence
- Morphing in latent space

11.2

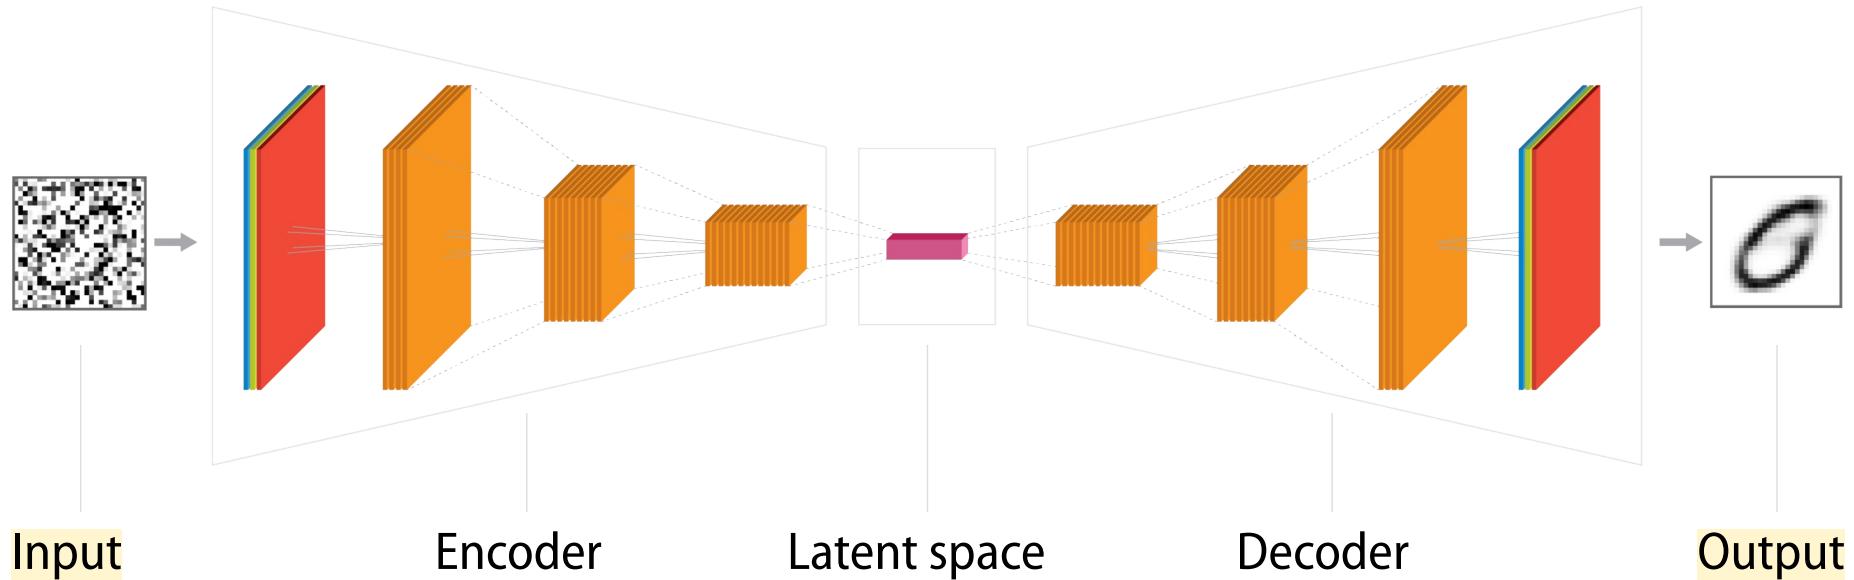
## Example 1 : VAE1/3

- Deeper in Keras !
- VAE using Functional API (MNIST)
- VAE using Model subclass (MNIST)



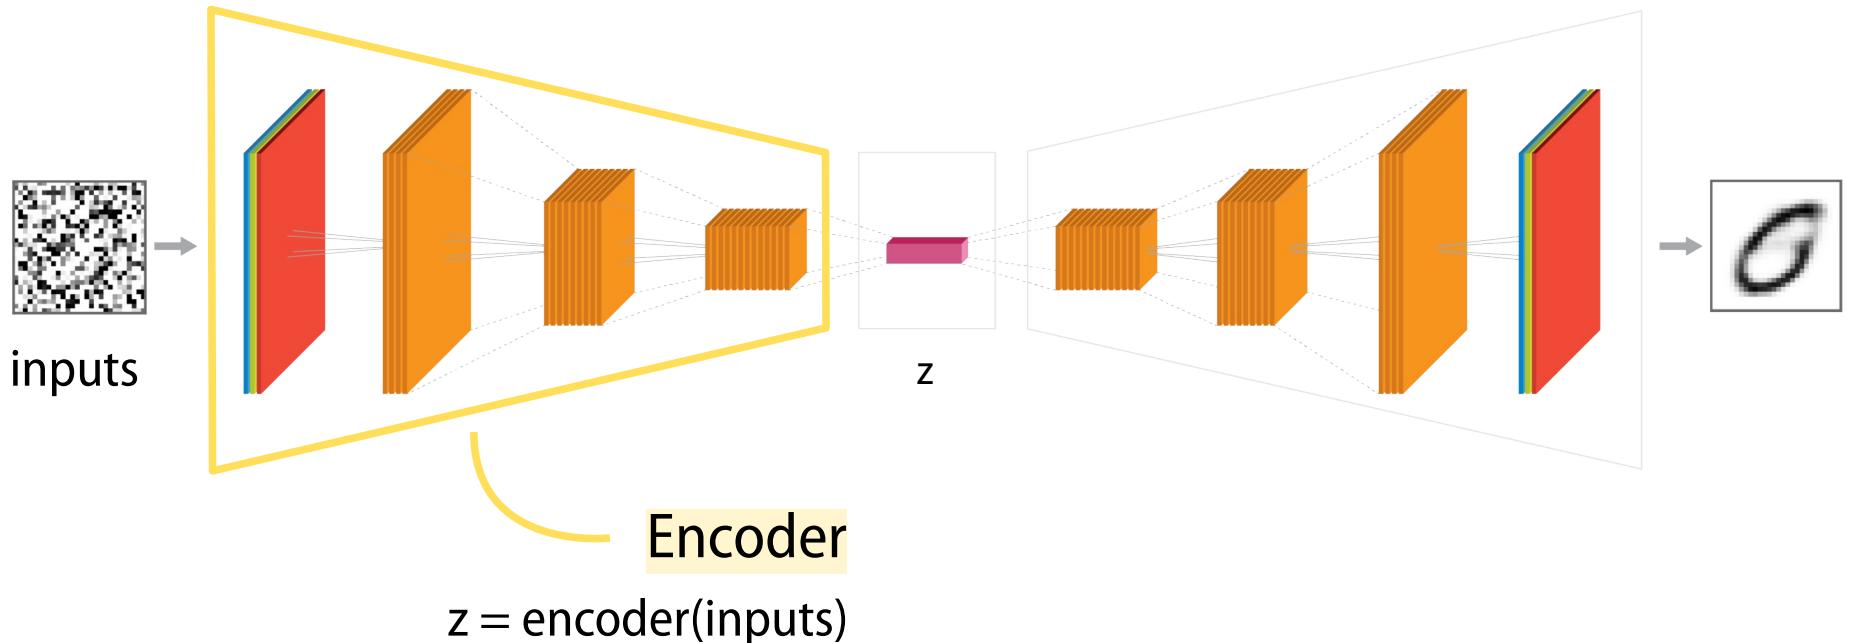
# Concepts & architecture

We have seen that an autoencoder network is trained to minimize a reconstruction error :



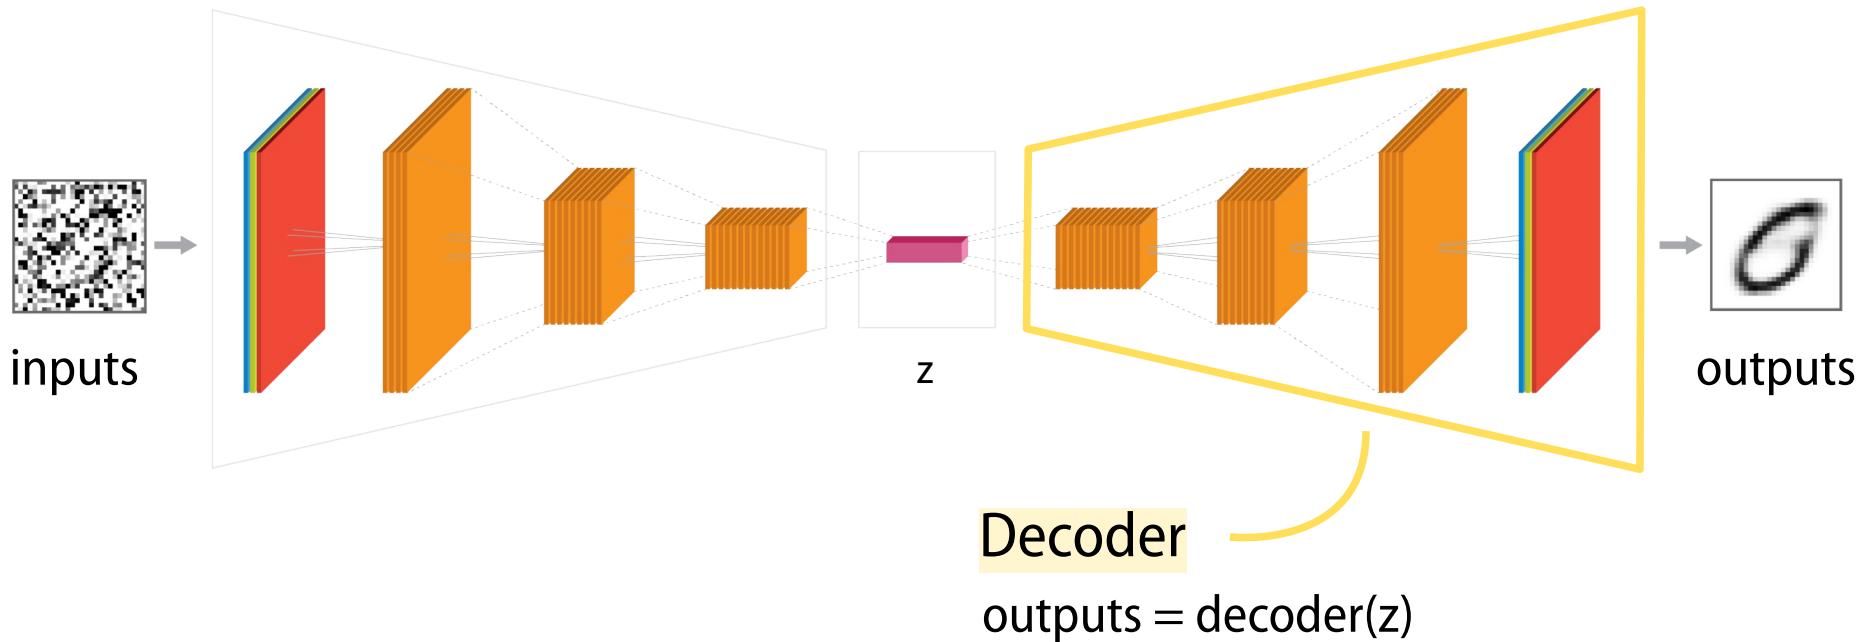
# Autoencoder

An autoencoders network have 2 parts :



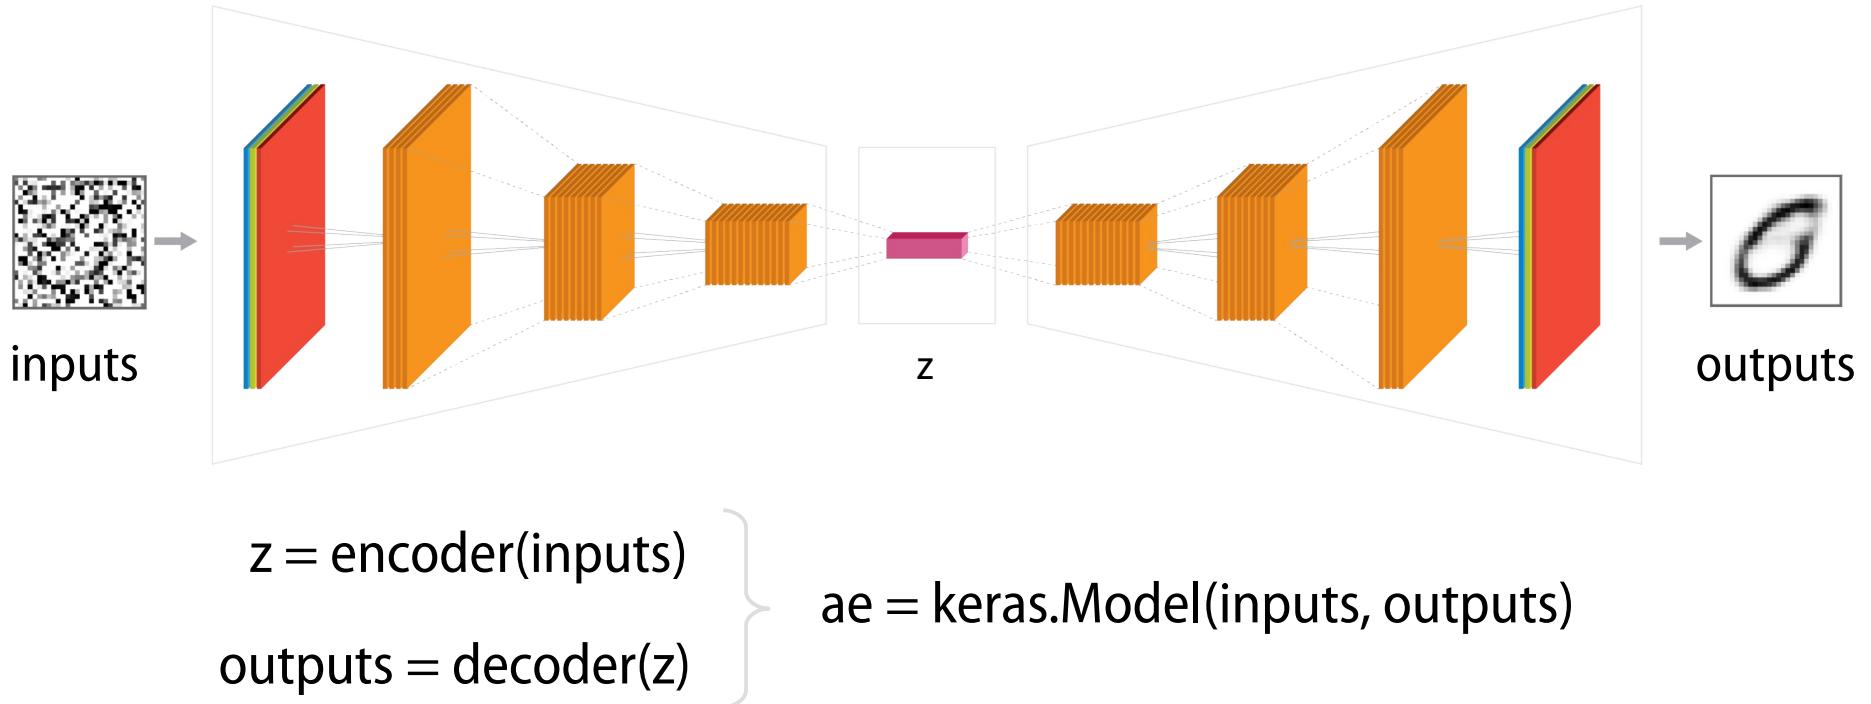
# Autoencoder

An autoencoders network have 2 parts :



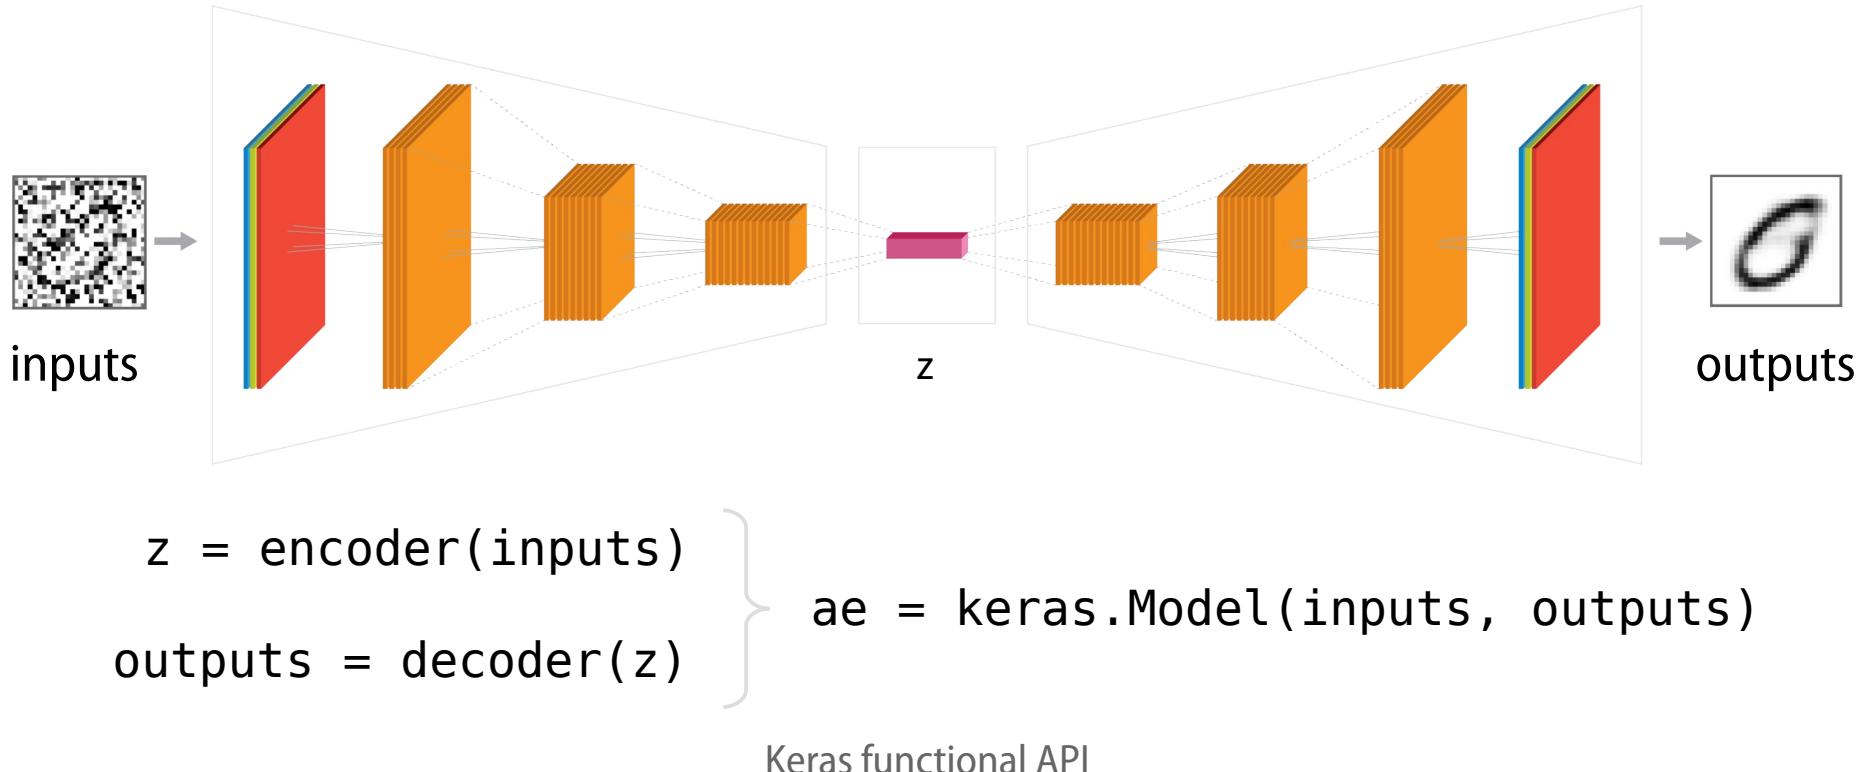
# Autoencoder

An autoencoders network have 2 parts :



# Autoencoder

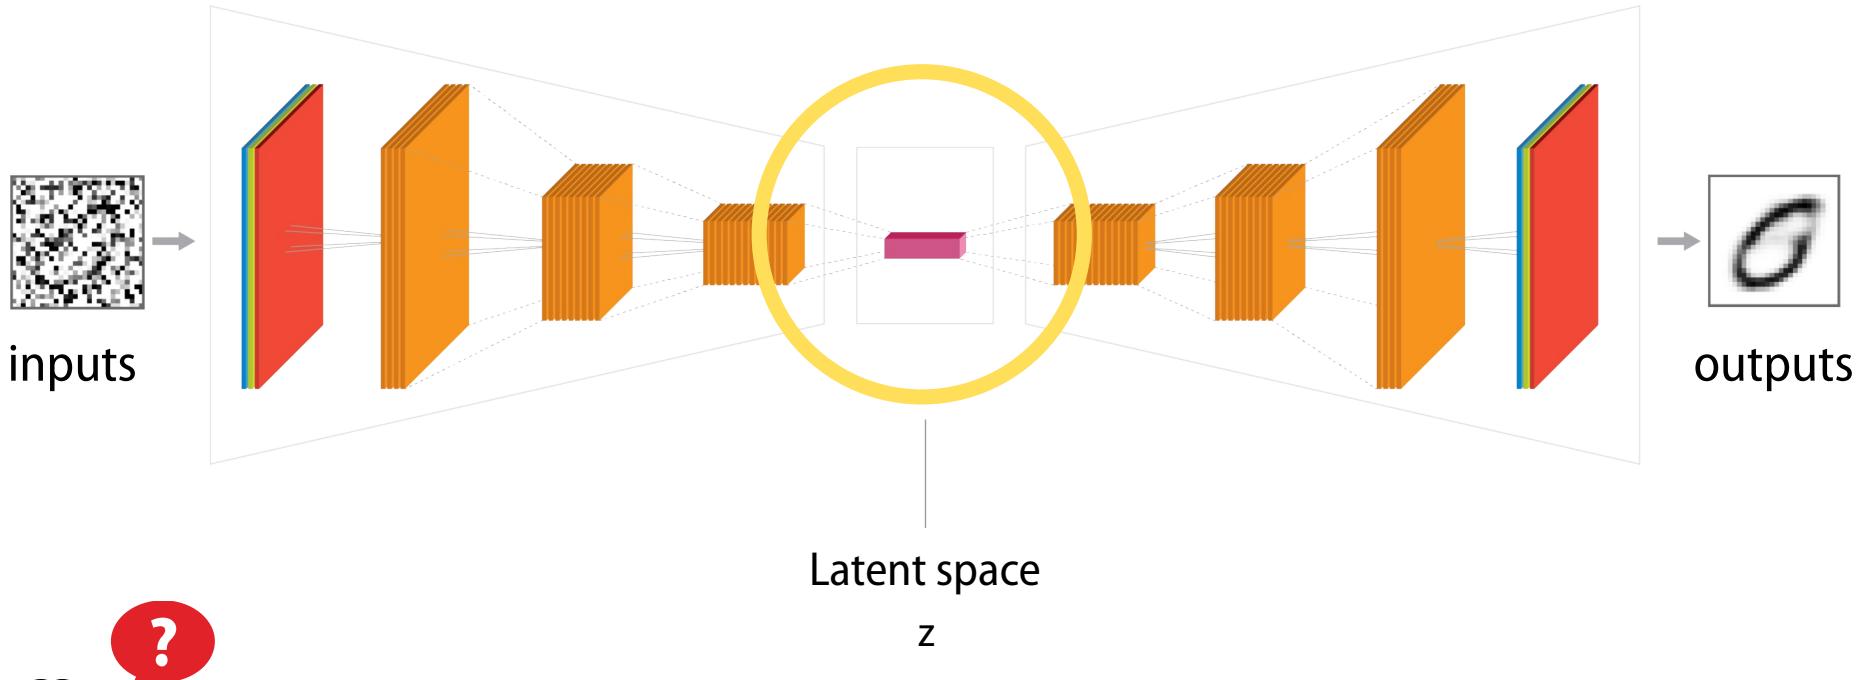
An autoencoders network have 2 parts :



# Autoencoder

3

An autoencoders network have ~~2~~ parts :



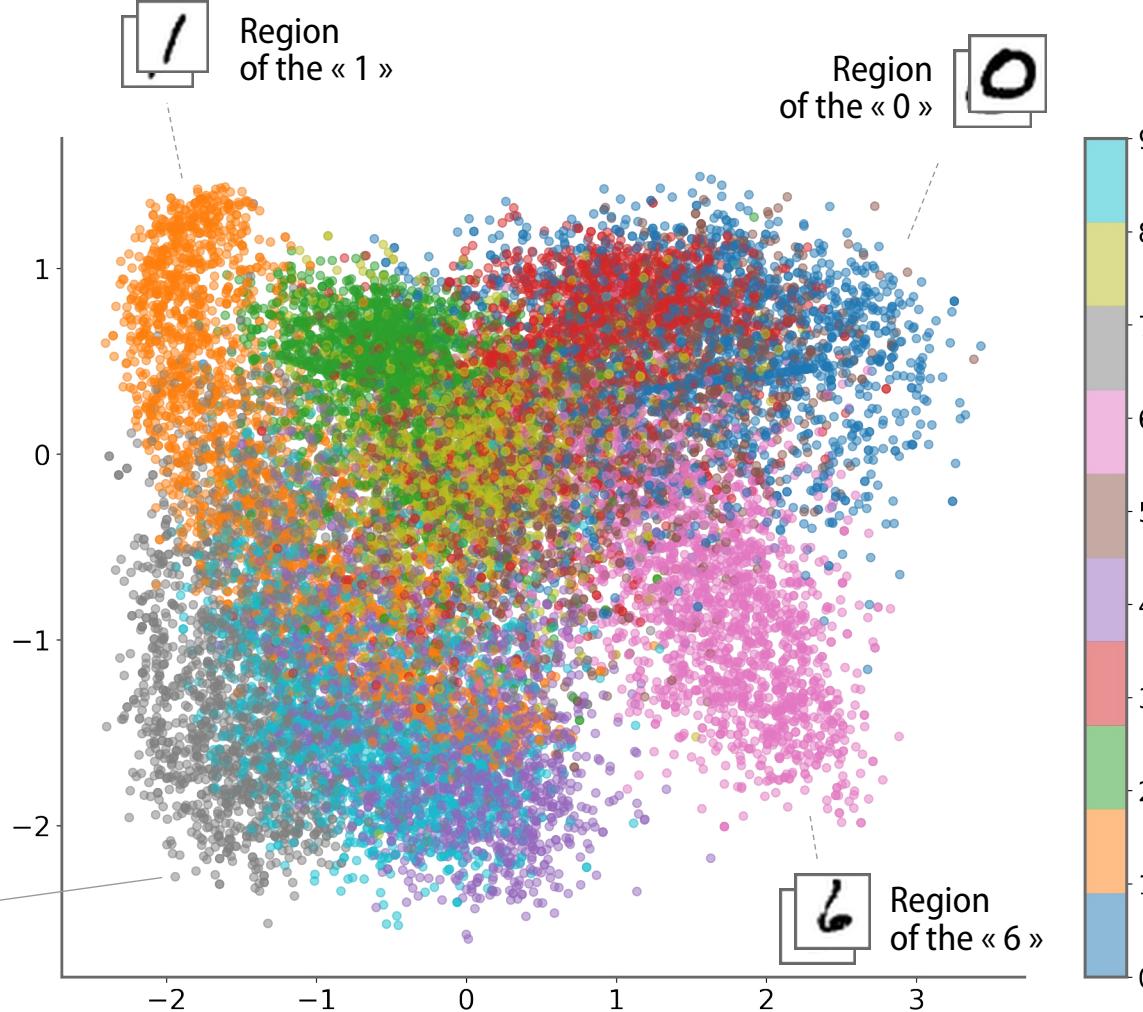
# Autoencoder

Example of MNIST dataset distribution in its latent space



Only two dimensions are represented in abscissa and ordinate

$z = \text{encoder}(\text{inputs})$



Clusters appear, but many of them are nested or very spread out



How can we make our network better separate the different clusters?



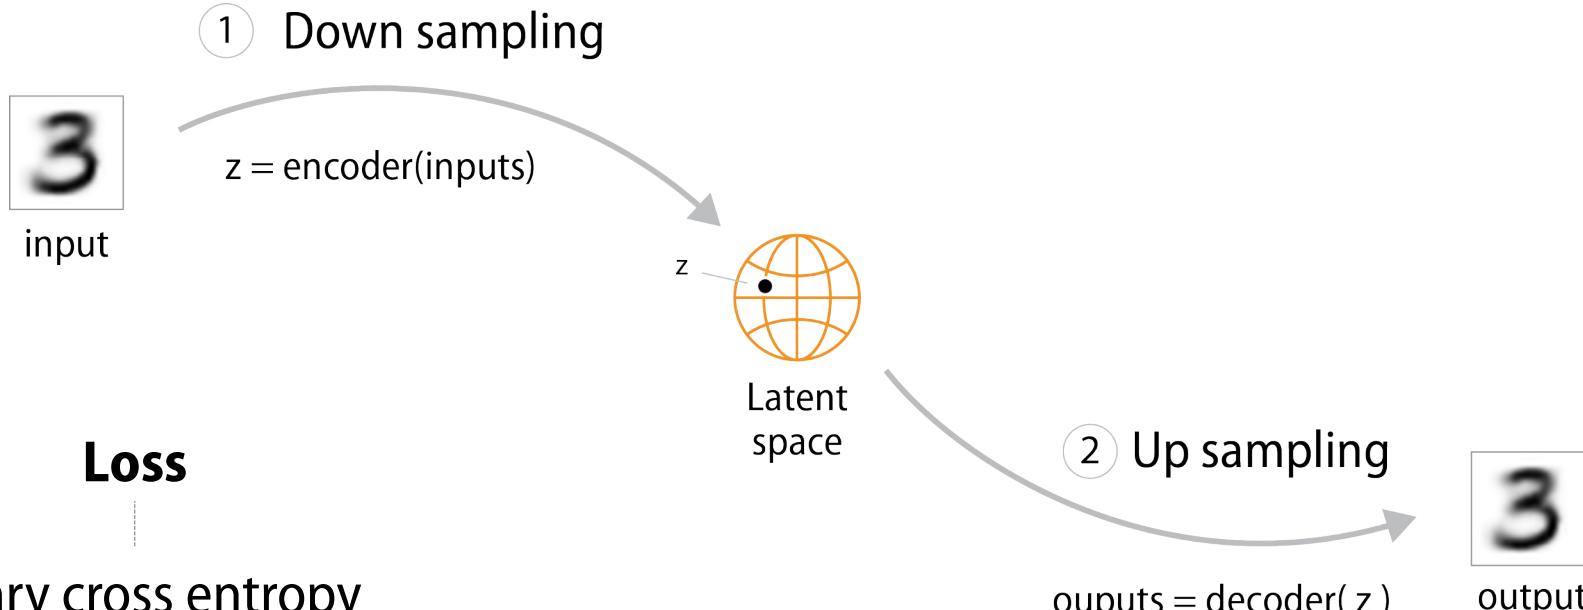
→ Just train it !

# Autoencoder (AE)

$z$  vector in latent space



An autoencoder performs a direct projection into the latent space and an upsampling from this latent space.



## Loss

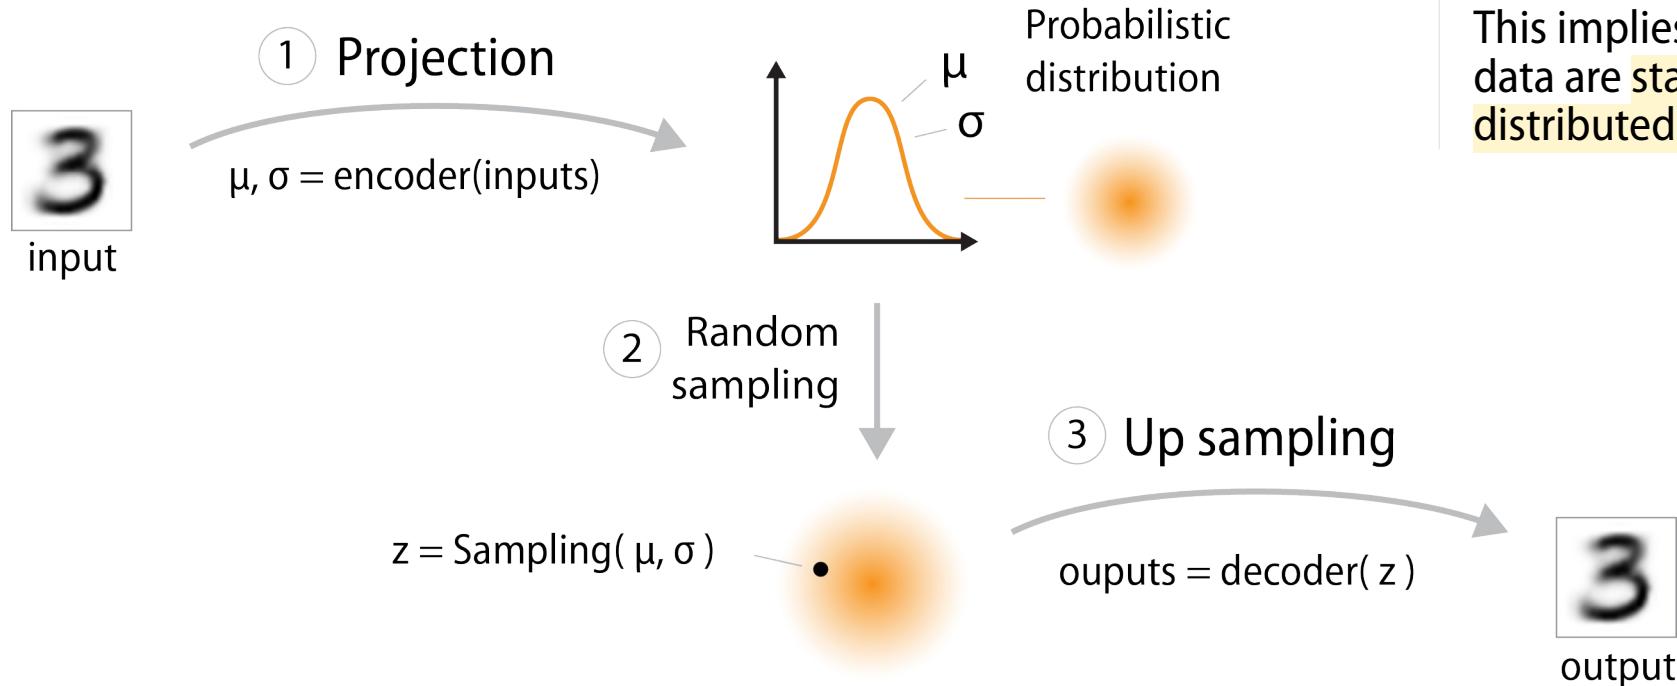
Binary cross entropy

Measures the difference  
between input and output

# Variational Autoencoder (VAE)



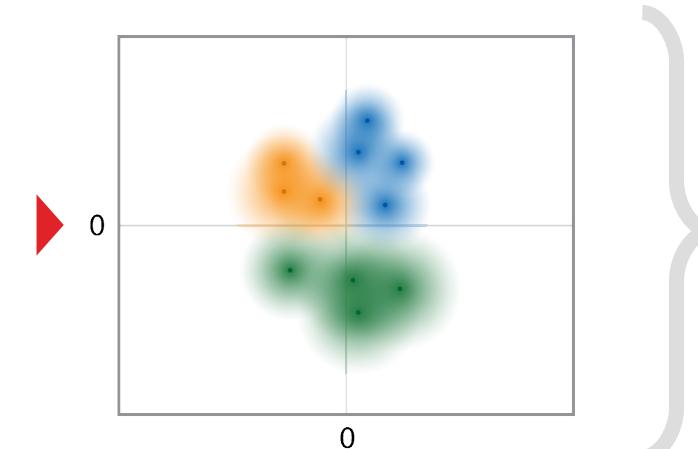
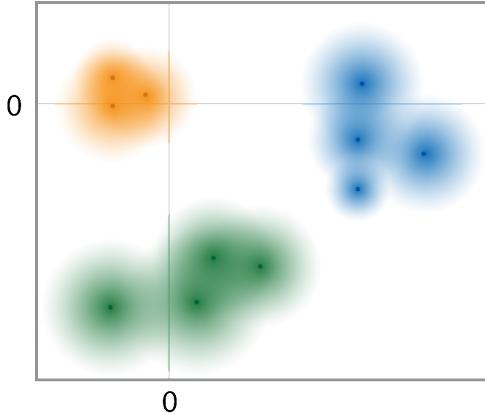
$\mu$  is a mean  
 $\sigma$  is a variance  
 $z$  vector in latent space



The stochastic approach encourages the network to generate relevant output throughout the distribution space.

This implies that the input data are statistically distributed.

# Variational Autoencoder (VAE)



## Reconstruction loss

Binary cross entropy

Measures the difference between input and output

## Kullback-Leibler divergence\*

$$D_{KL} = \frac{1}{2} \sum_{i=1}^k (\sigma_i^2 + \mu_i^2 - 1 - \ln(\sigma_i^2)).$$

KL divergence between two (probability) distributions measures how much they diverge from each other

## KL loss



## Total loss

$$\text{total}_{\text{loss}} = k_1 \cdot r_{\text{loss}} + k_2 \cdot \text{kl}_{\text{loss}}$$



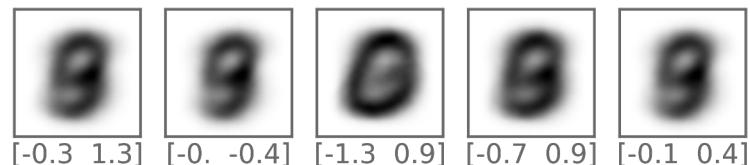
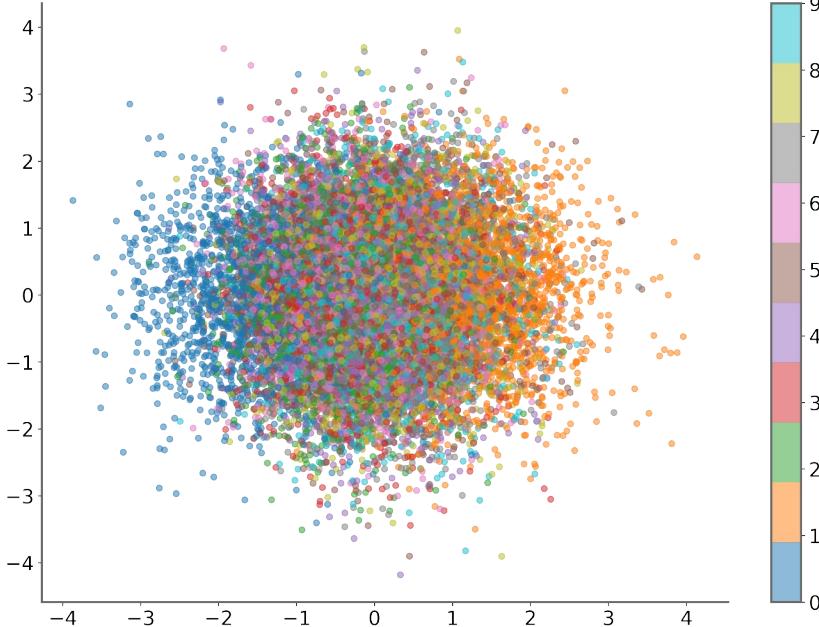
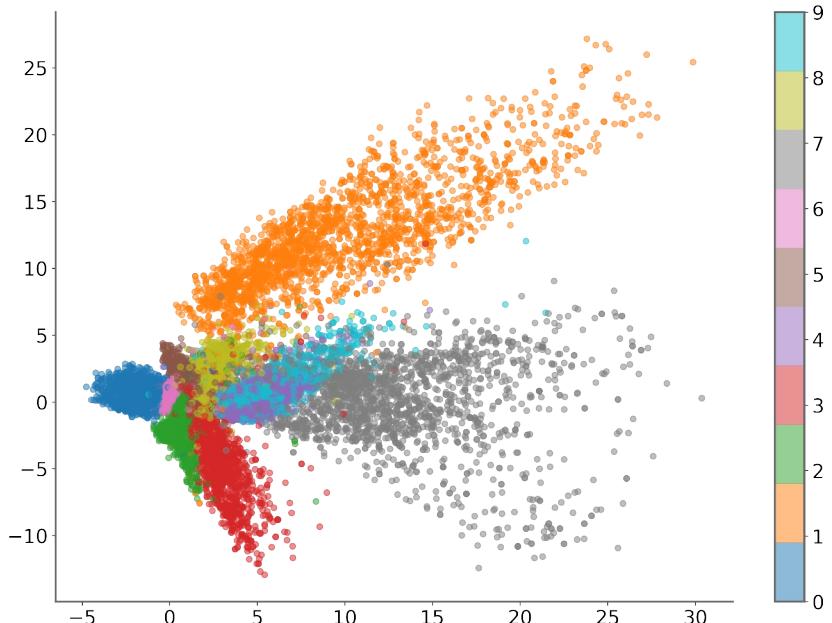
The trick is to find a nice compromise between these two components of the loss function.

(\*) Special case for a standard normal distribution

# Variational Autoencoder (VAE)

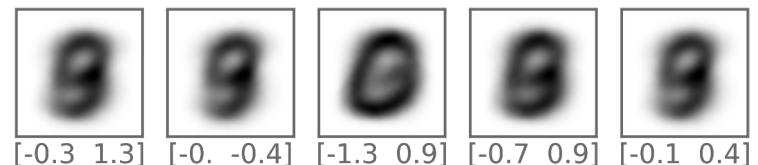
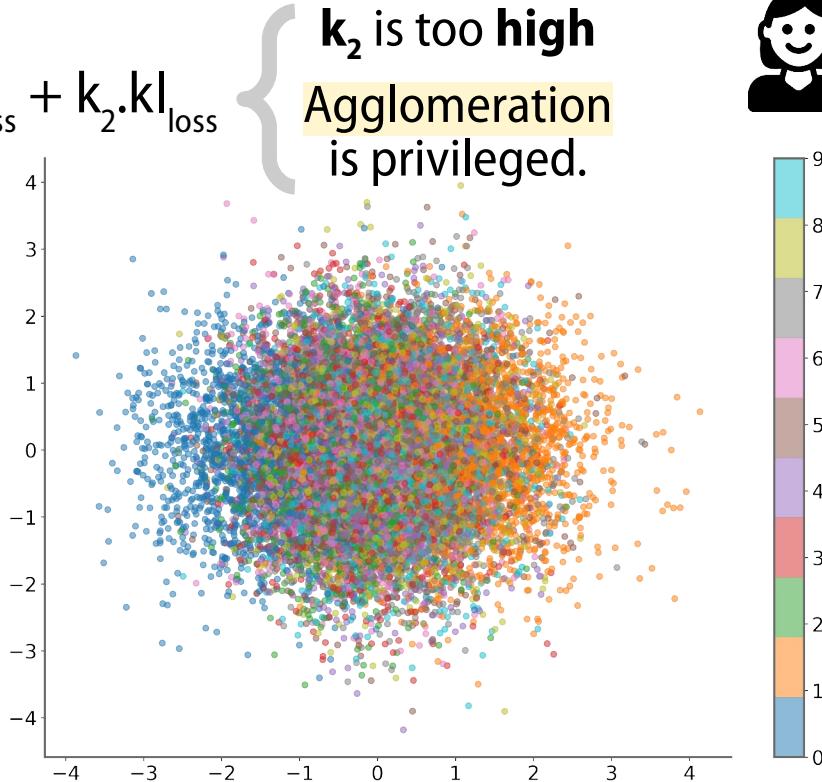
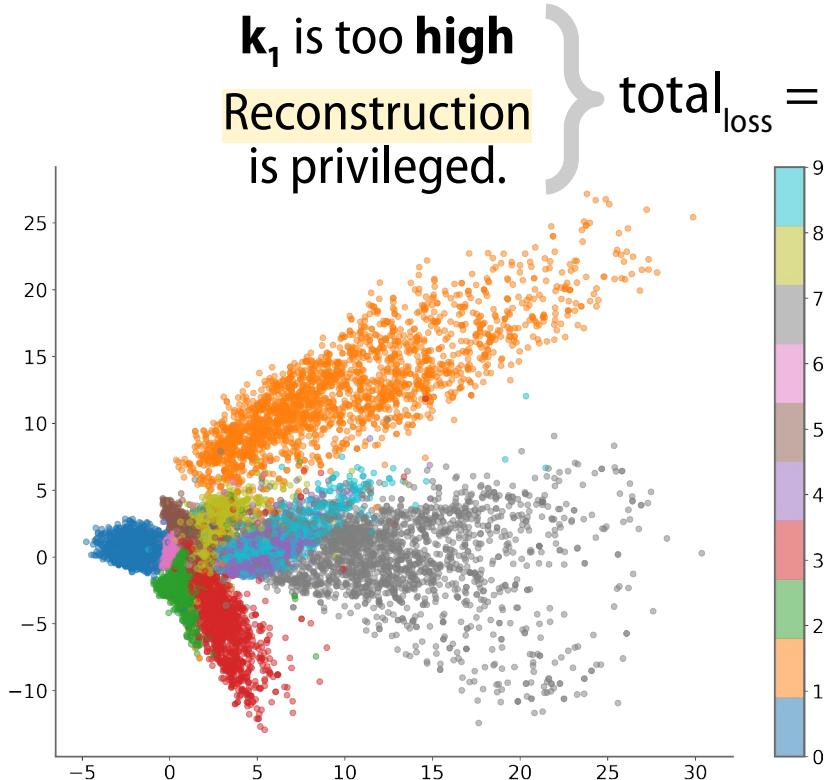


$$\text{total loss} = k_1 \cdot r_{\text{loss}} + k_2 \cdot \text{kl}_{\text{loss}}$$



# Variational Autoencoder (VAE)

Easy !



# Variational Autoencoder (VAE)

Nice !



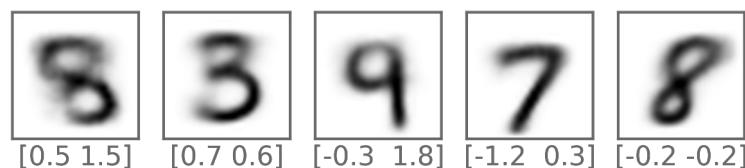
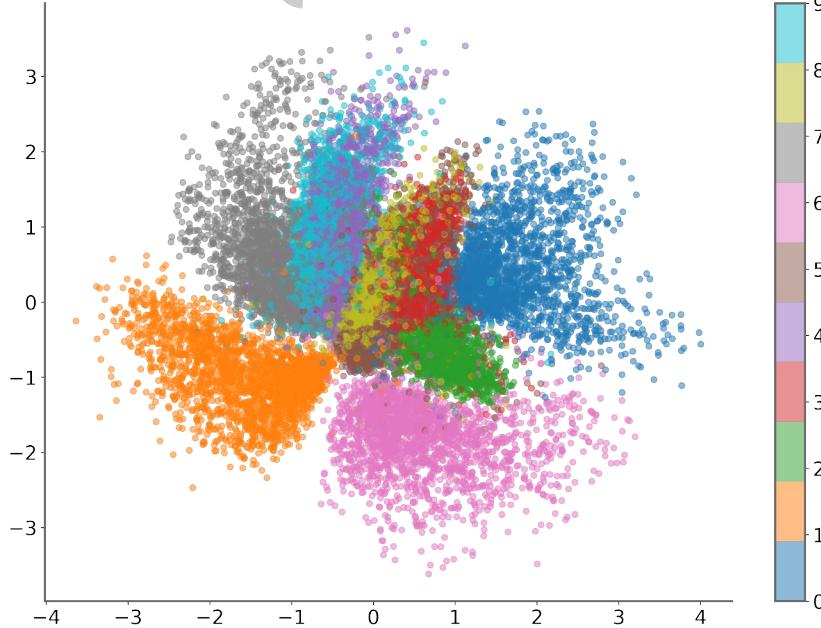
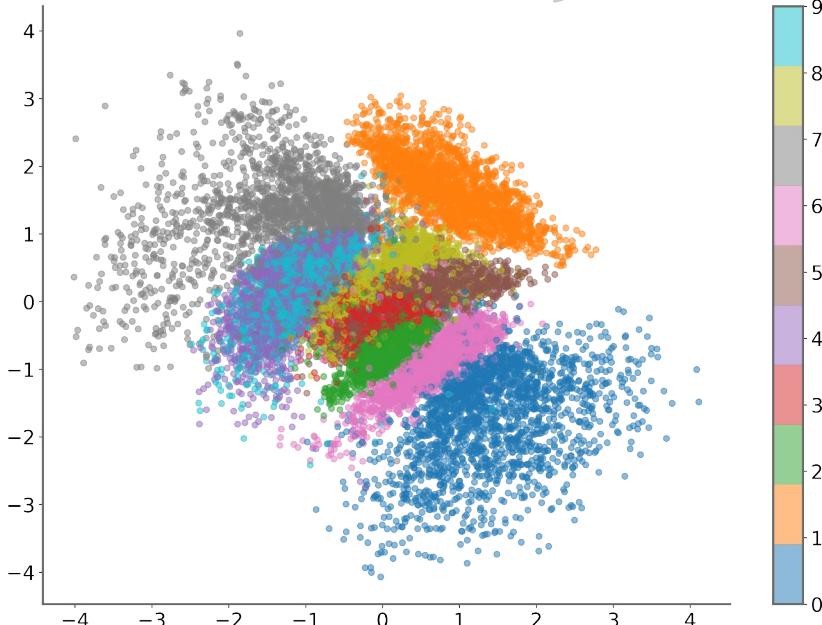
VAE2, with :  
 $k_1, k_2 = [1, 5e-4]$

}

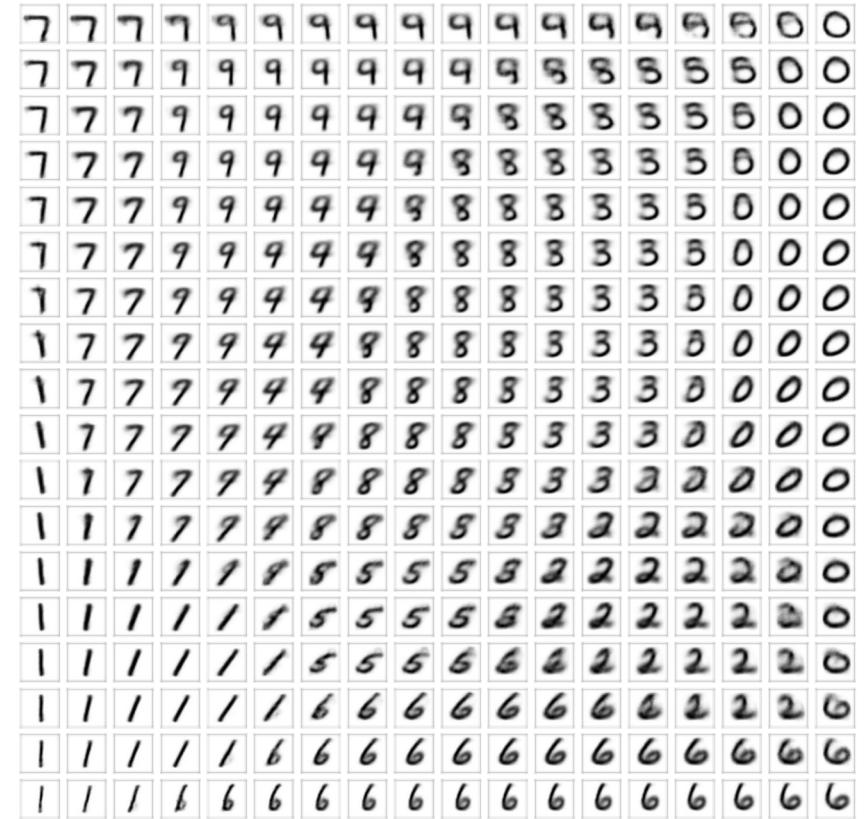
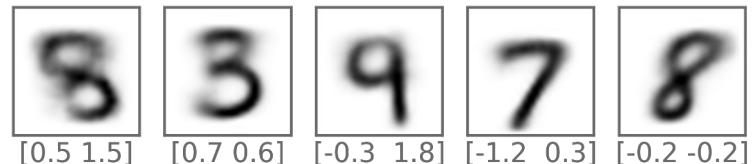
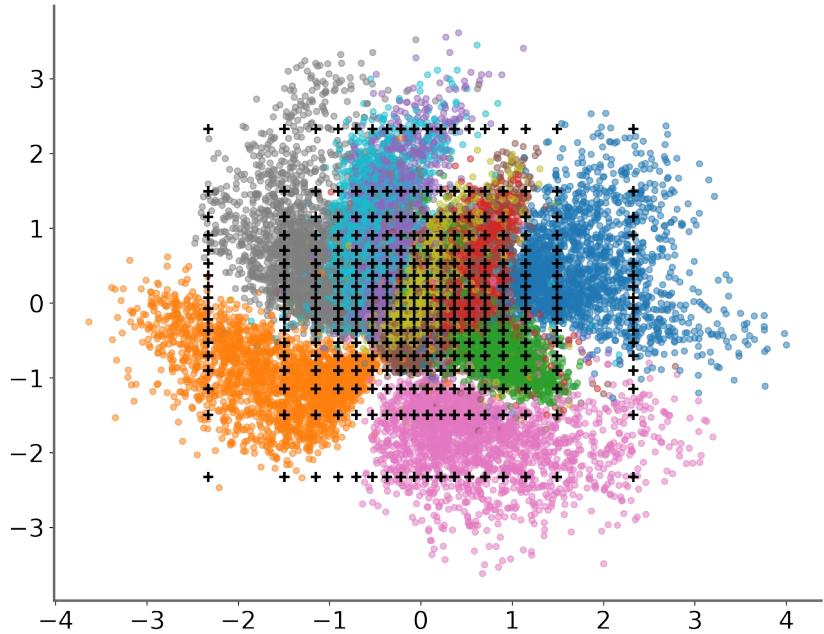
$$\text{total}_{\text{loss}} = k_1 \cdot r_{\text{loss}} + k_2 \cdot \text{kl}_{\text{loss}}$$

VAE2, with :  
 $k_1, k_2 = [1, 1e-3]$

}



# Variational Autoencoder (VAE)



Tha's great ! We can generate data in profusion !!!

11



Variational  
Autoencoder

VAE

11.1

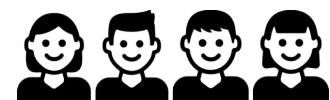
## Variational Autoencoder network

- Concepts and architecture
- Gaussian / probabilistic projections
- Kullback-Leibler divergence
- Morphing in latent space

11.2

## Example 1 : VAE1/3

- Deeper in Keras !
- VAE using Functional API (MNIST)
- VAE using Model subclass (MNIST)



11



## Variational Autoencoder

VAE



11.1

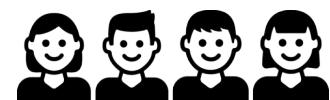
### Variational Autoencoder network

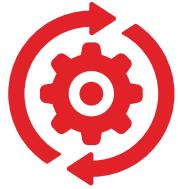
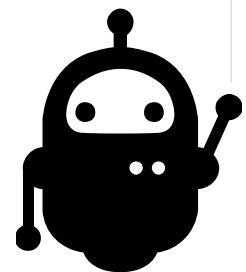
- Concepts and architecture
- Gaussian / probabilistic projections
- Kullback-Leibler divergence
- Morphing in latent space

11.2

### Example 1 : VAE1/3

- Deeper in Keras !
- VAE using Functional API (MNIST)
- VAE using Model subclass (MNIST)





# Our first VAE

Notebook : [\[VAE1/3\]](#)

## **Objectives :**

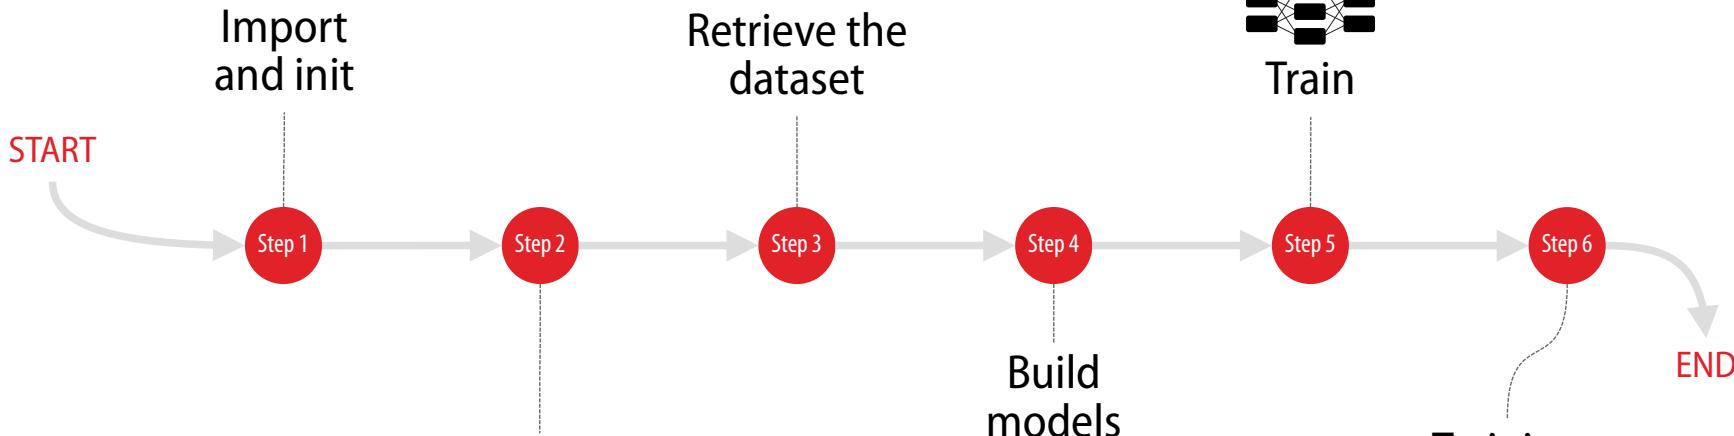
implementing a VAE, using Keras functionel API  
and model subclass, using real Python !

## **Dataset :**

MNIST



20' on a CPU  
With scale=1



## Objectives :

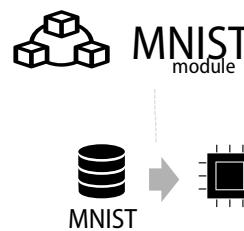
VAE using  
Keras  
functional API  
and MNIST  
dataset

### Parameters

```

latent_dim = 2
loss_weights = [1., .001]

scale       = .1
seed        = 123
  
```



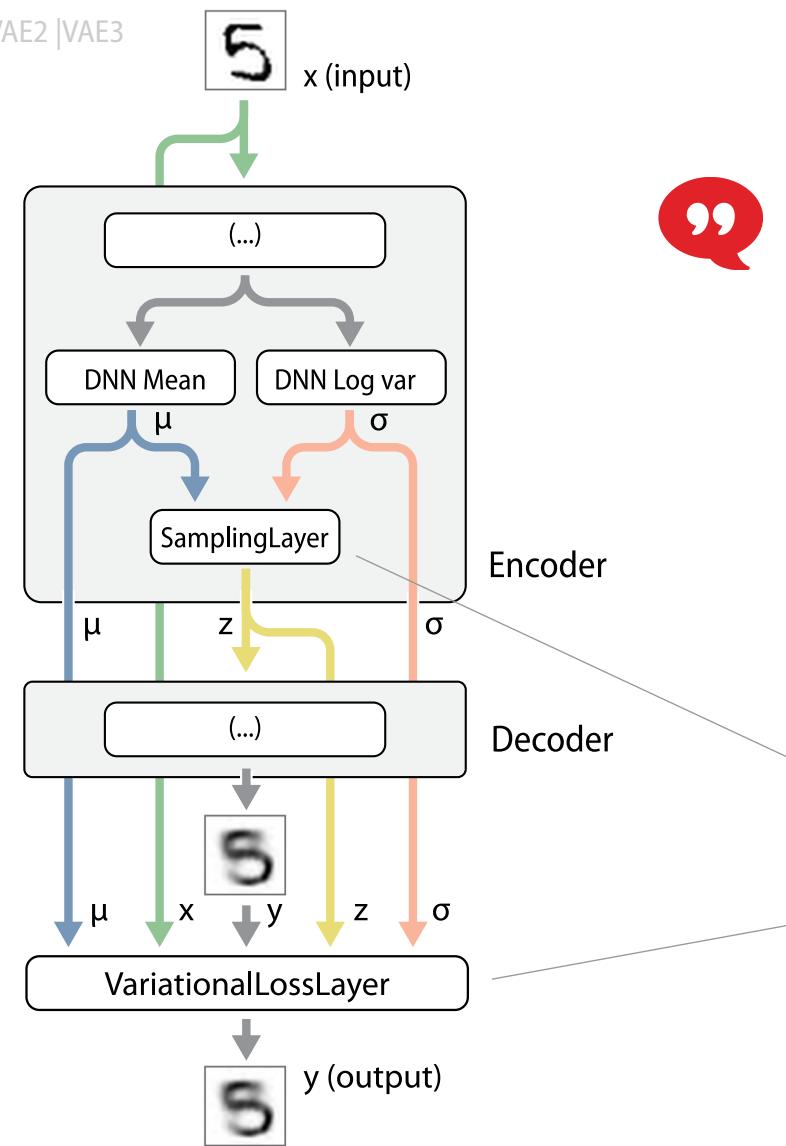
MNIST module

ImagesCallBack module

BestModelCallBack module

SamplingLayer module

VariationalLossLayer module



The classical loss functions use only the **inputs** and the **outputs** tensors.

The calculation of the loss is here more complex and will be done via a **specific layer**.

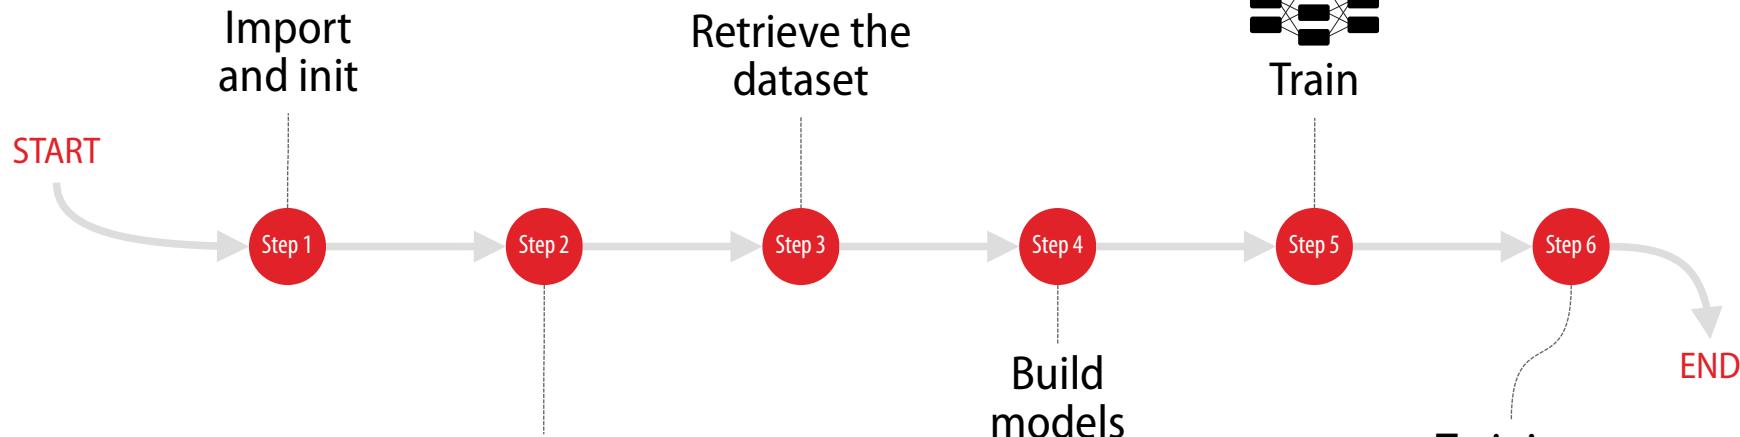
So we have two custom layers in this example :

 **SamplingLayer**

 **VariationalLossLayer**



20' on a CPU  
With scale=1



## Objectives :

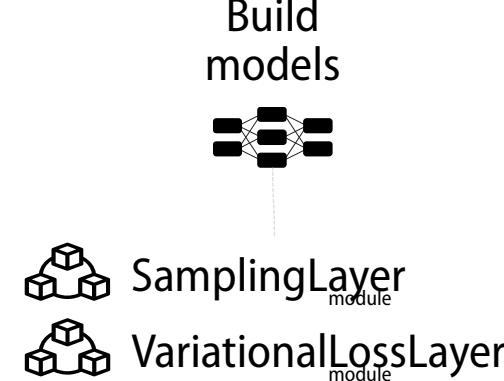
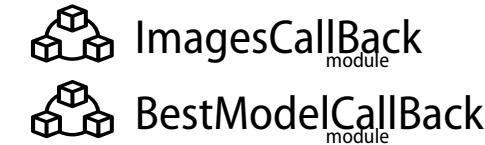
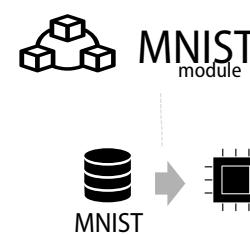
VAE using  
Keras  
functional API  
and MNIST  
dataset

### Parameters

```

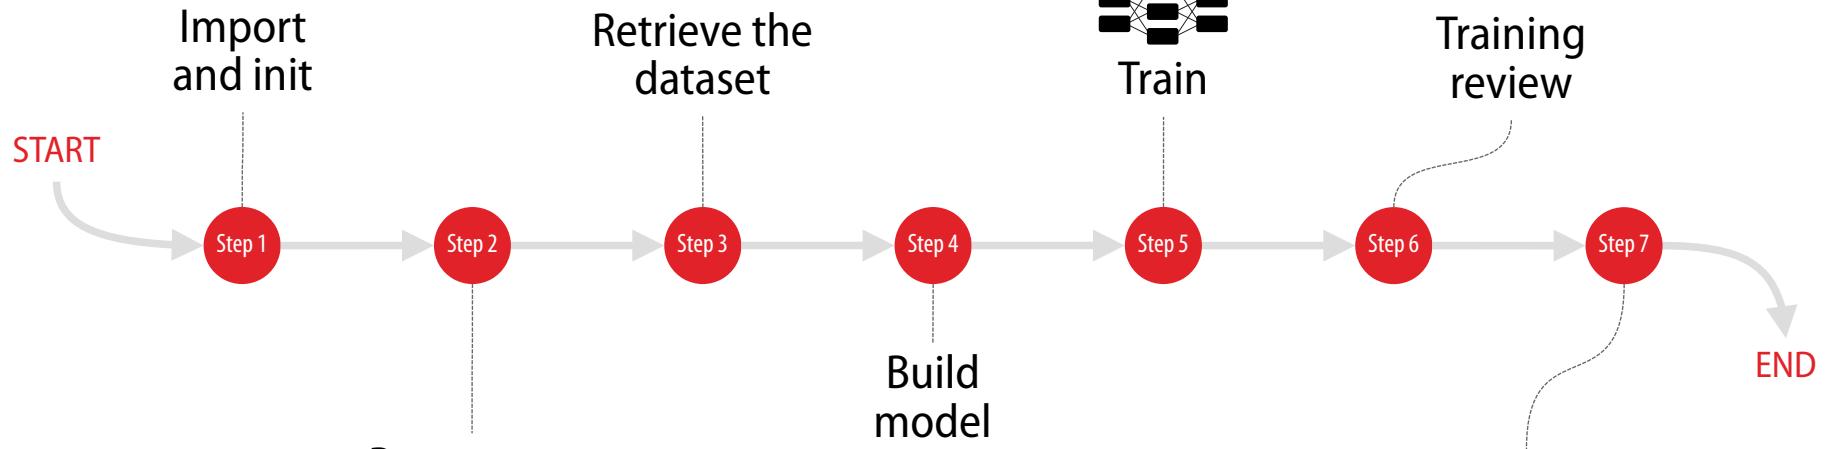
latent_dim = 2
loss_weights = [1., .001]

scale = .1
seed = 123
  
```





20' on a CPU  
With scale=1



## Objectives :

VAE using  
Keras subclass  
API and MNIST  
dataset

```
latent_dim = 2
loss_weights = [1, .001]
```

```
scale = .1
seed = 123
```



MNIST  
module



Retrieve the  
dataset



SamplingLayer  
module



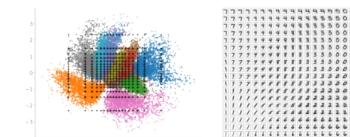
VAE  
module

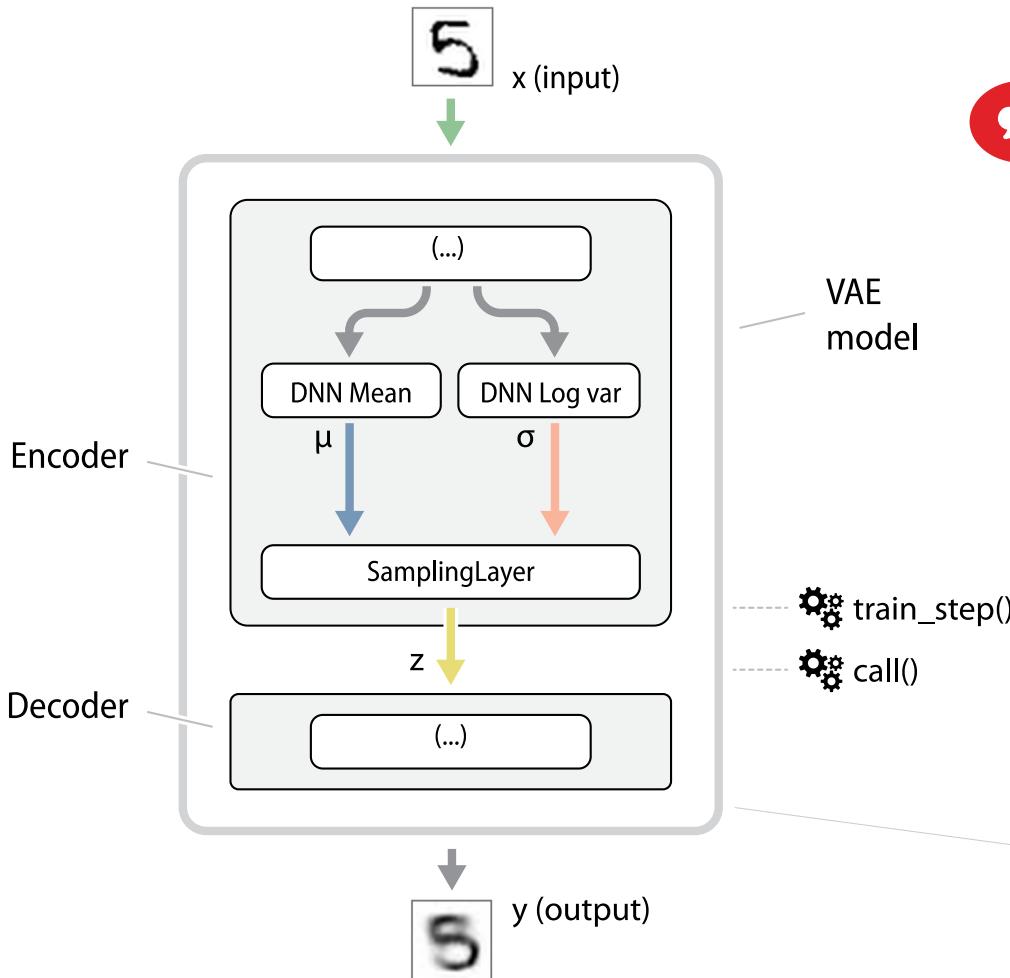


ImagesCallBack  
module



BestModelCallBack  
module





The encoder and decoder are always built via the functional API and a custom layer.

The model is defined by a custom class in real Python :-)

The gradient descent and the loss function are defined via a method.

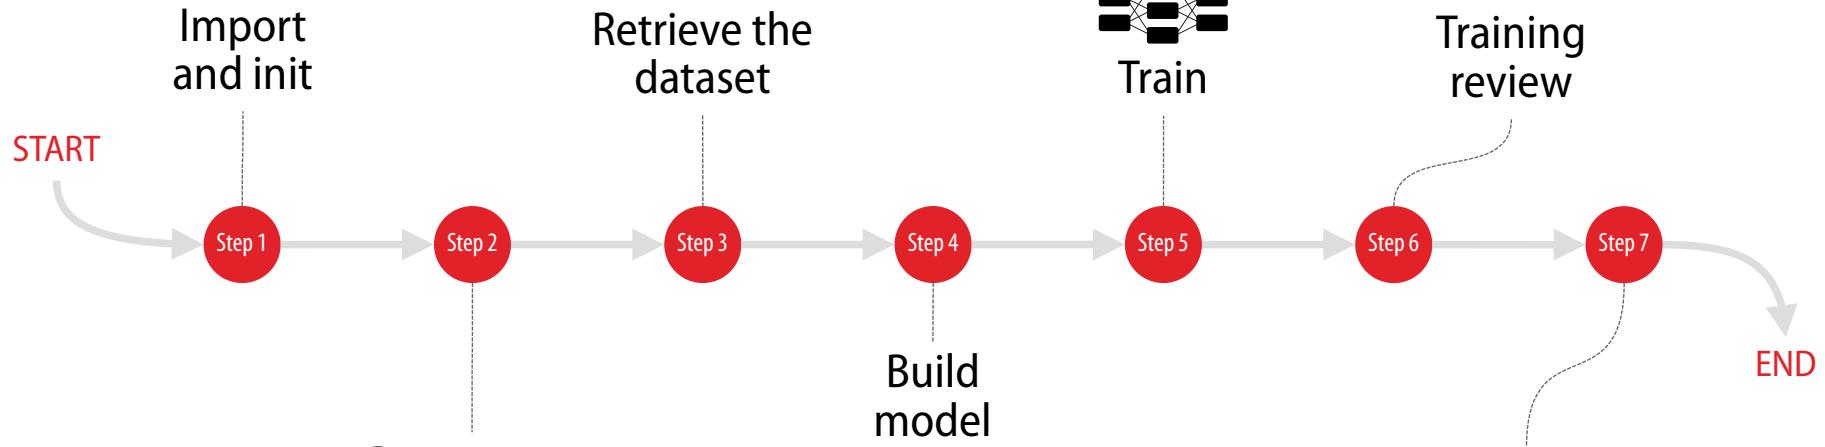
We have one custom layer and one custom model class :

 SamplingLayer

 VAE



20' on a CPU  
With scale=1

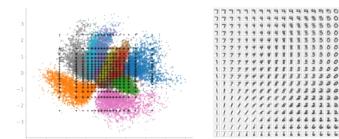


## Objectives :

VAE using  
Keras subclass  
API and MNIST  
dataset

```
latent_dim = 2
loss_weights = [1, .001]
```

```
scale = .1
seed = 123
```

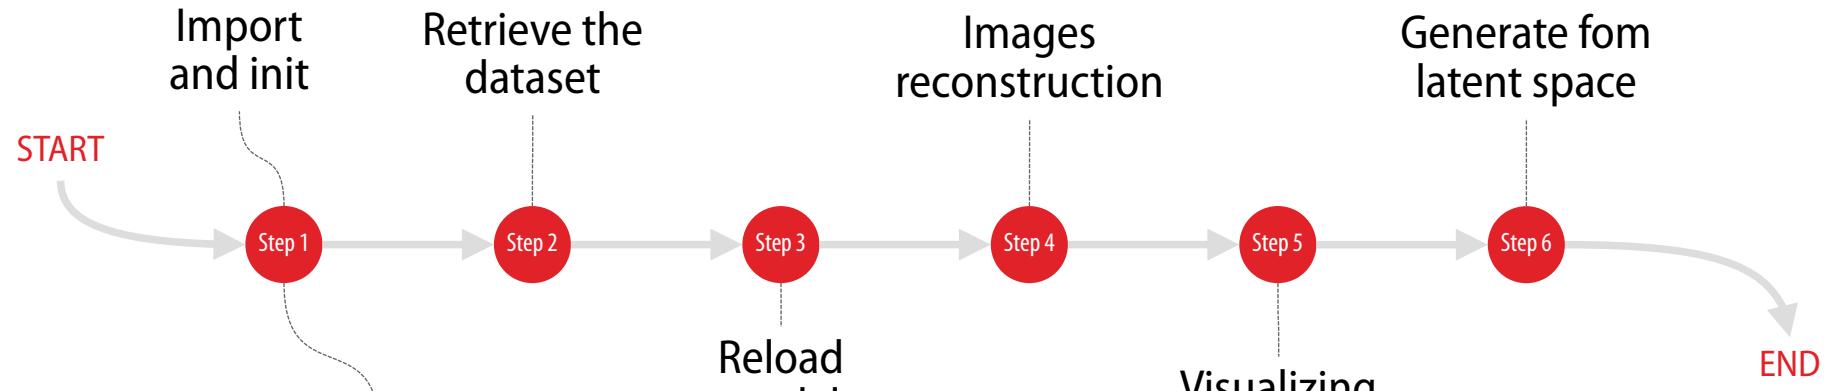




VAE1 | VAE2 | **VAE3**



Few seconds !



Parameters

scale = .1  
seed = 123

**Objectives :**  
Reload a saved model and visualize latent space



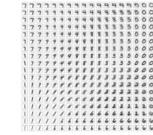
MNIST  
module



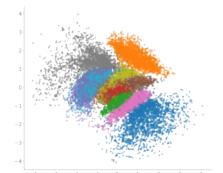
MNIST



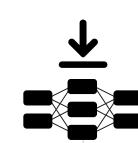
Images  
reconstruction



Generate from  
latent space



Reload  
model



SamplingLayer  
module



VAE  
module



35

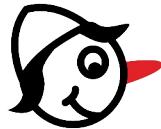


## Little things and concepts to **keep in mind**

- Unsupervised learning is great!
- The use cases of the VAE are very wide
  - Anomaly detection
  - Clustering
  - Semi-supervised classification
  - ...
- VAEs are often coupled with other tools
- Balancing the loss function is not always easy to find
- GPUs are great!
- Other (and better?) generative networks exist
- ...

Next, on Fidle :

13



**Generative  
Adversarial Networks**  
GAN



**Jeudi 9 mars,**

Épisode 13 :

## **Generative Adversarial Networks (GAN)**

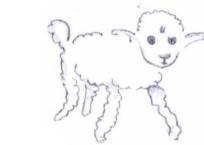
Principes et architecture d'un GAN

Generator - Discriminator - Apprentissage

WGAN et WGAN-GP

Exemple proposé : Dessine moi un mouton !

Durée : 2h



# « Journée Deep Learning pour la Science ! »



Next, on Fidle :

13



**Generative  
Adversarial Networks**  
GAN

**Jeudi 9 mars,**

Épisode 13 :

## **Generative Adversarial Networks (GAN)**

Principes et architecture d'un GAN

Generator - Discriminator - Apprentissage

WGAN et WGAN-GP

Exemple proposé : Dessine moi un mouton !

Durée : 2h

👉 Mon projet en 180 secondes !



Next on Fidle :



Jeudi 9 mars,

Séquence 13 :

**Generative Adversarial Networks (GAN)**  
**Principes et architecture d'un GAN**



To be continued...