



Formation
Introduction au
Deep Learning

Generative Adversarial Networks (GANs)





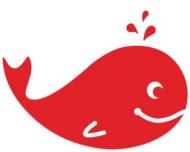
Cette session va être enregistrée.
Retrouvez-nous sur notre chaîne YouTube :-)
This session will be recorded.
Find us on our YouTube channel :-)

<https://fidle.cnrs.fr/youtube>



Formation
Introduction au
Deep Learning

Generative Adversarial Networks (GANs)



FIDLE



MIAI
Grenoble Alpes



UGA
Université
Grenoble Alpes



Resources

<https://fidle.cnrs.fr>

Powered by CNRS CRIC, and UGA DGDSI
of Grenoble, Thanks !



Course materials (pdf)



Practical work environment*



Corrected notebooks



Videos (YouTube)

(*) Procedure via Docket or pip
Remember to get the latest version !



Attribution-NonCommercial-NoDerivatives 4.0 International (CC BY-NC-ND 4.0)
<https://creativecommons.org/licenses/by-nc-nd/4.0/>

Resources

You can also subscribe to :



<http://fidle.cnrs.fr/listeinfo>
Fidle information list



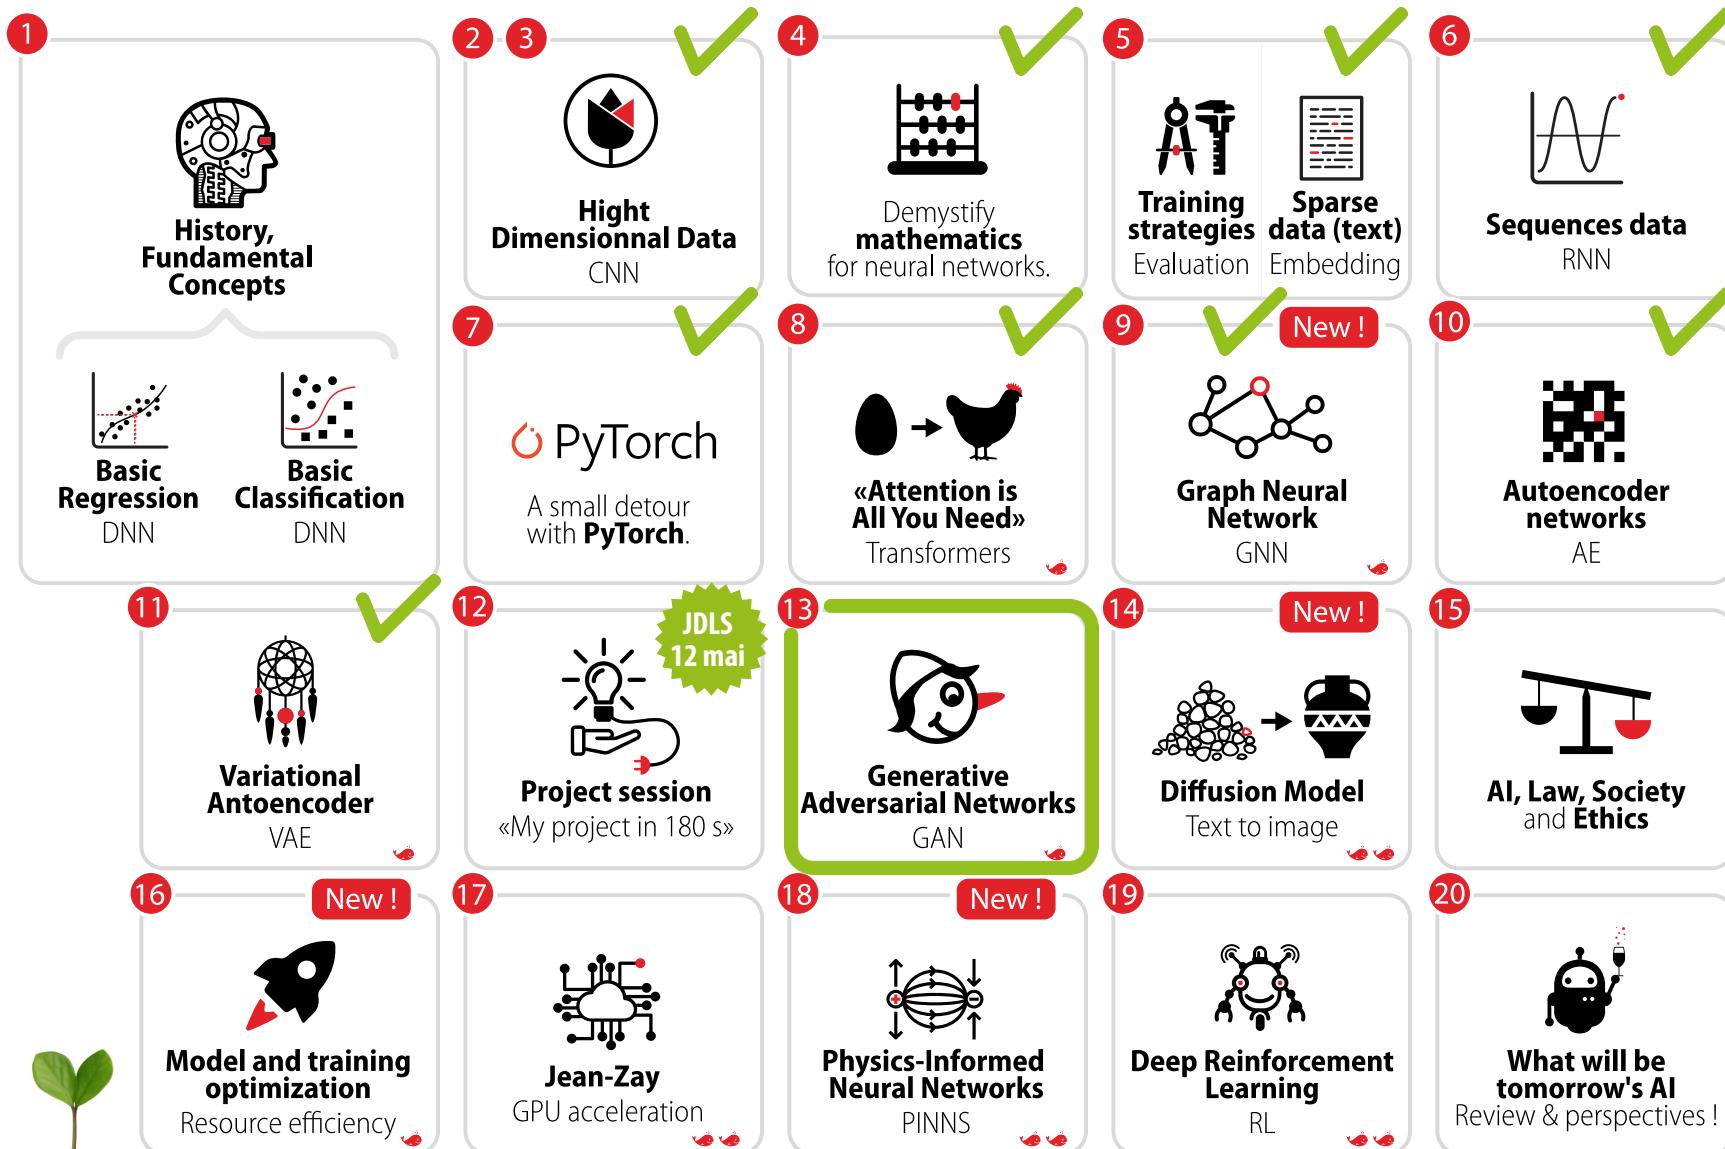
<https://listes.services.cnrs.fr/wws/info/devlog>
List of ESR* « Software developers » group



<https://listes.math.cnrs.fr/wws/info/calcul>
List of ESR* « Calcul » group

Program

FIDDLE



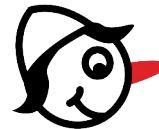
20 Séquences
du 17 novembre
au 14 mai 2023



SAISON
22/23

Roadmap

13



Generative
Adversarial Networks
GAN

1

About GAN

- Principles
- Training

2

Example 1 : Sheep1

- Draw me a sheep !

3

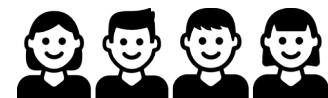
From GANs to *GANs

- Wasserstein and Earth Moving Distance
- From GAN to WGAN
- From WGAN to WGAN-GP

4

Example 2 : Sheep2

- Again, but with a WGAN-GP :-)



Roadmap

13



**Generative
Adversarial Networks**
GAN

Illustration : <https://www.saintexupery-domainepublic.be/>

1

About GAN

- Principles
- Training

2

Example 1 : Sheep1

- Draw me a sheep !

3

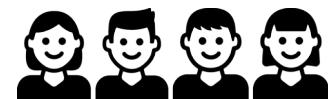
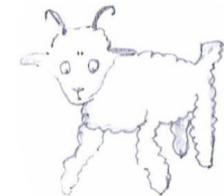
From GANs to *GANs

- Wasserstein and Earth Moving Distance
- From GAN to WGAN
- From WGAN to WGAN-GP

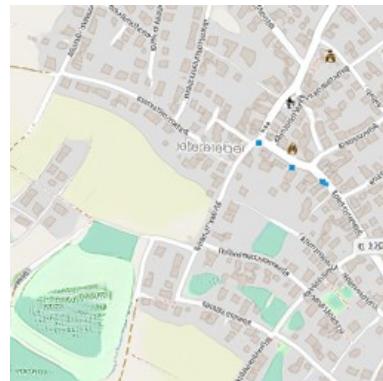
4

Example 2 : Sheep2

- Again, but with a WGAN-GP :-)



About GAN



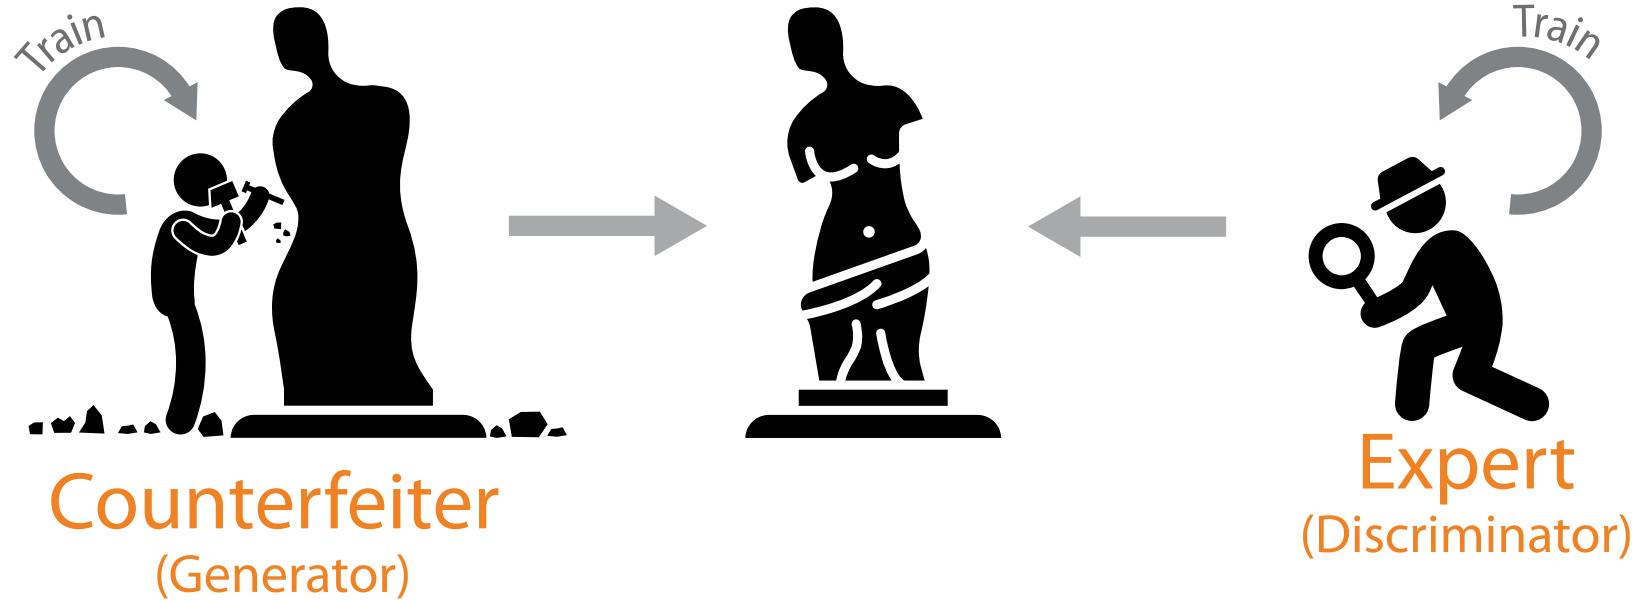
Which Face Is Real ?



<https://www.whichfaceisreal.com>

Sophie J. Nightingale and Hany Farid, « AI-synthesized faces are indistinguishable from real faces and more trustworthy », PNAS, <https://doi.org/10.1073/pnas.2120481119>, February 2022

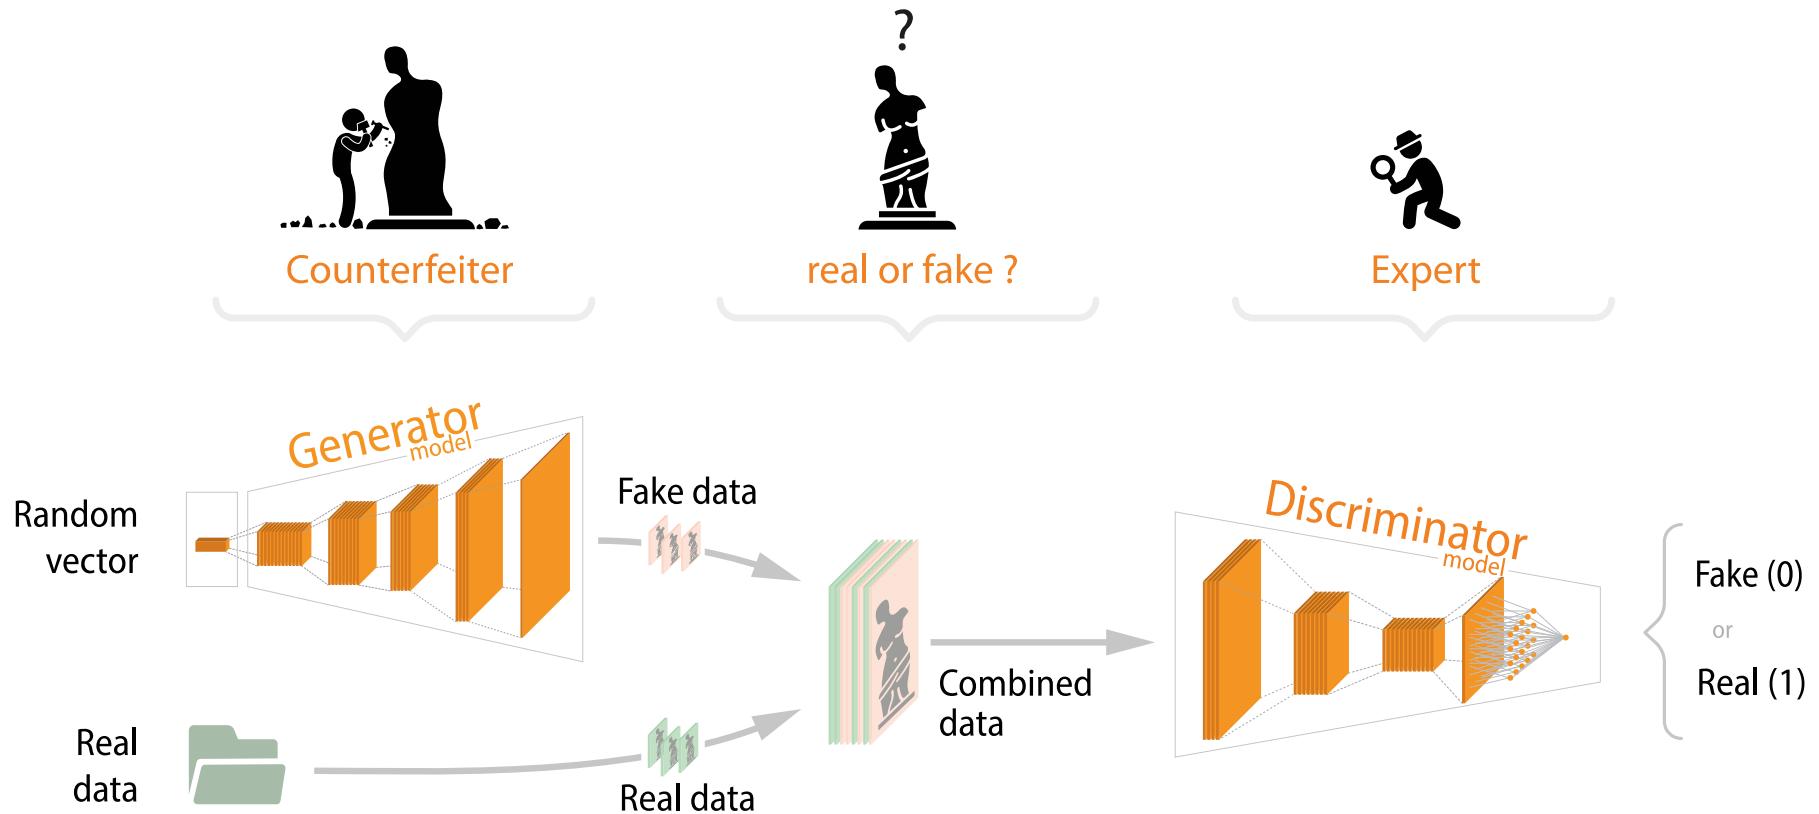
{ « (...) synthetically generated faces are more trustworthy than real faces. This may be because synthesized faces tend to look more like average faces which themselves are deemed more trustworthy. (...) »



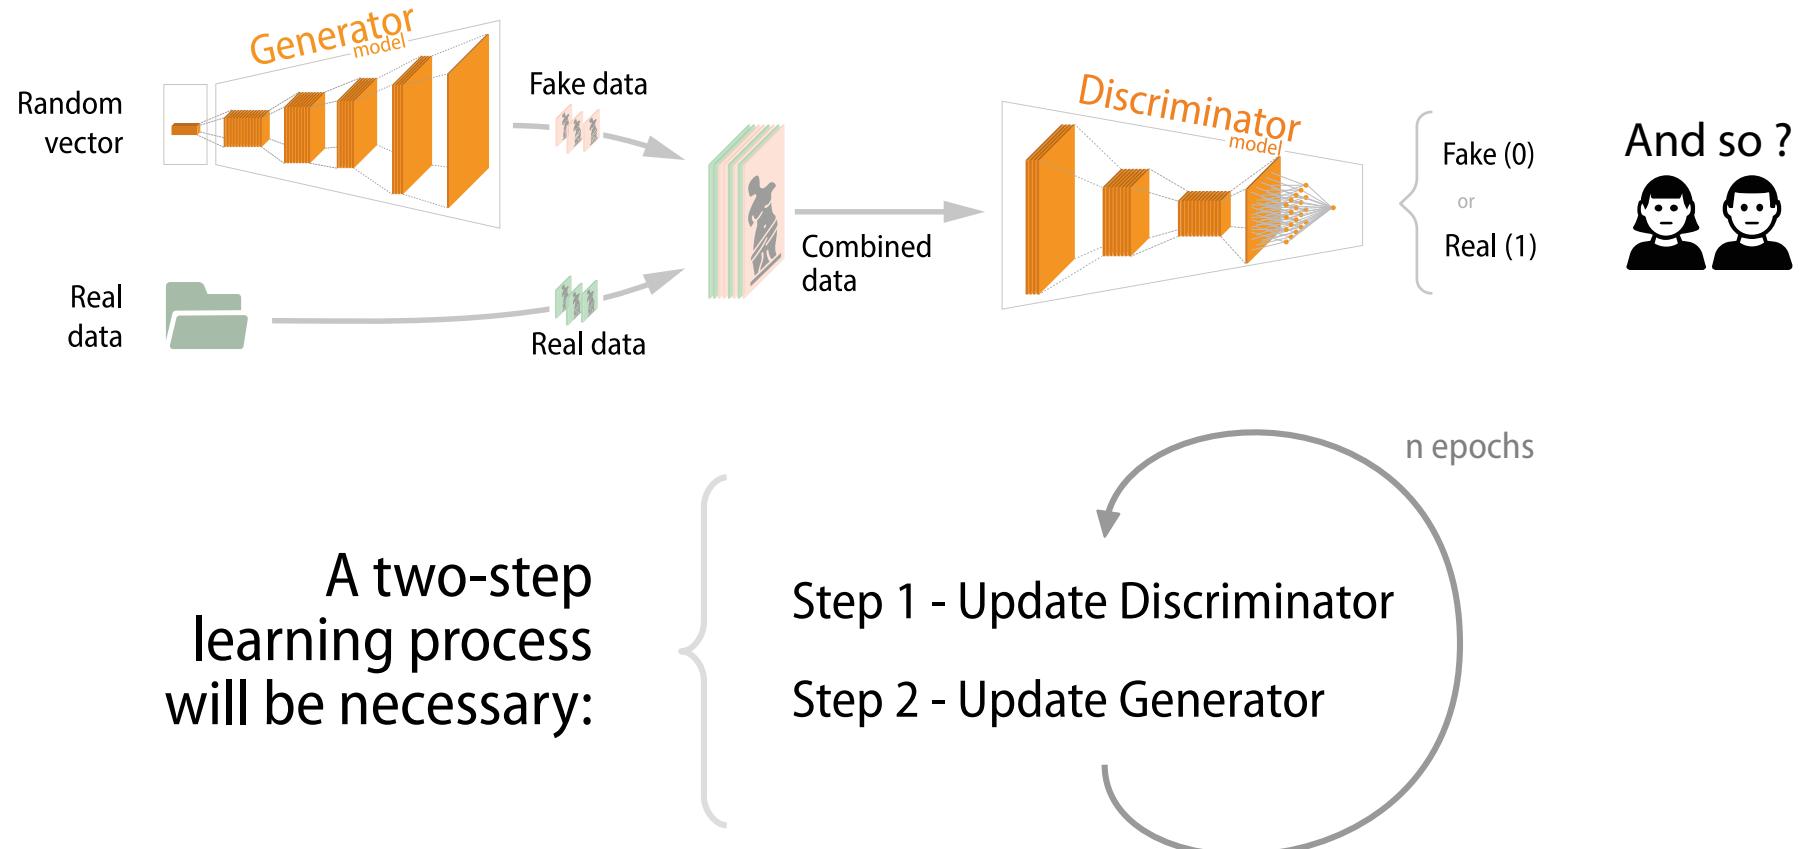
Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville et Yoshua Bengio, « Generative Adversarial Networks », in Advances in Neural Information Processing Systems 27, 2014

<https://arxiv.org/abs/1406.2661>

Principle



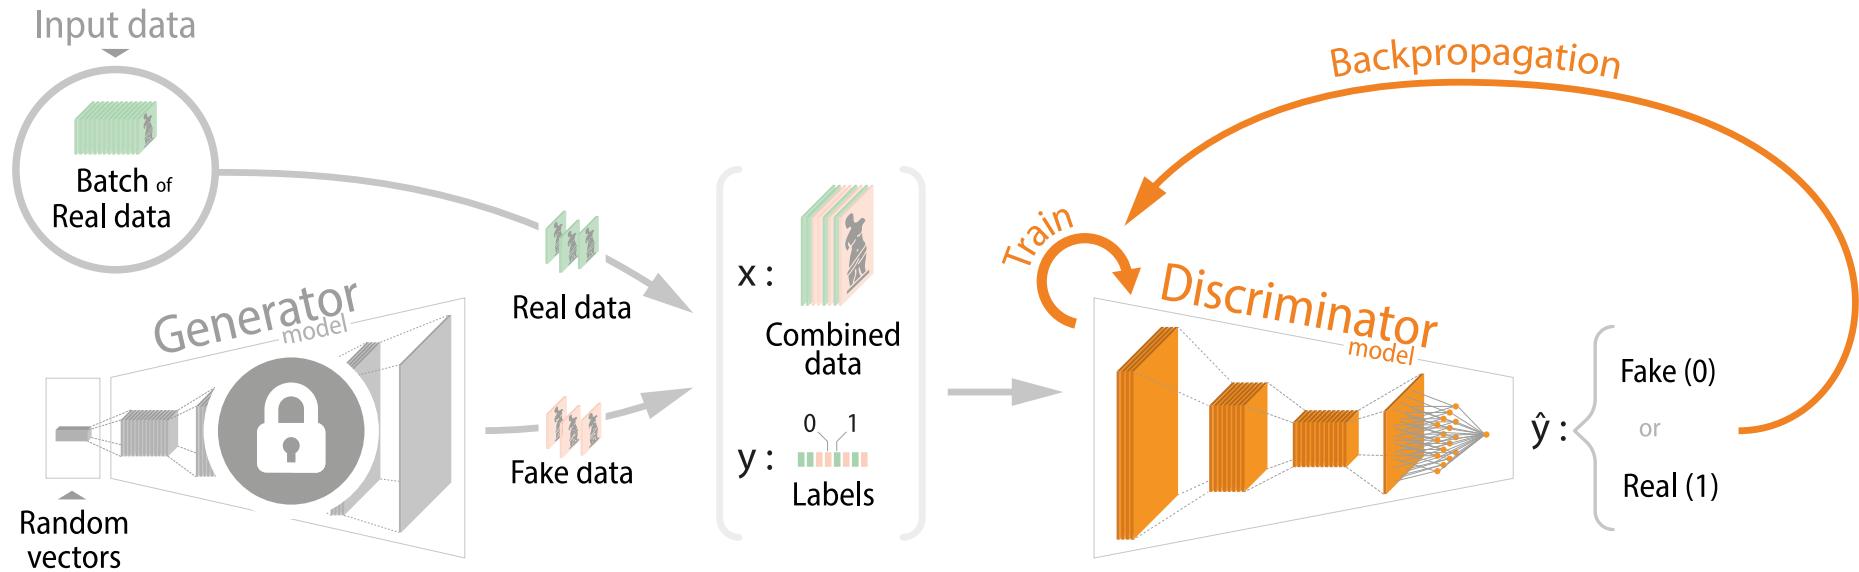
How to train a GAN ?



How to train a GAN ?

Step1 - Update Discriminator

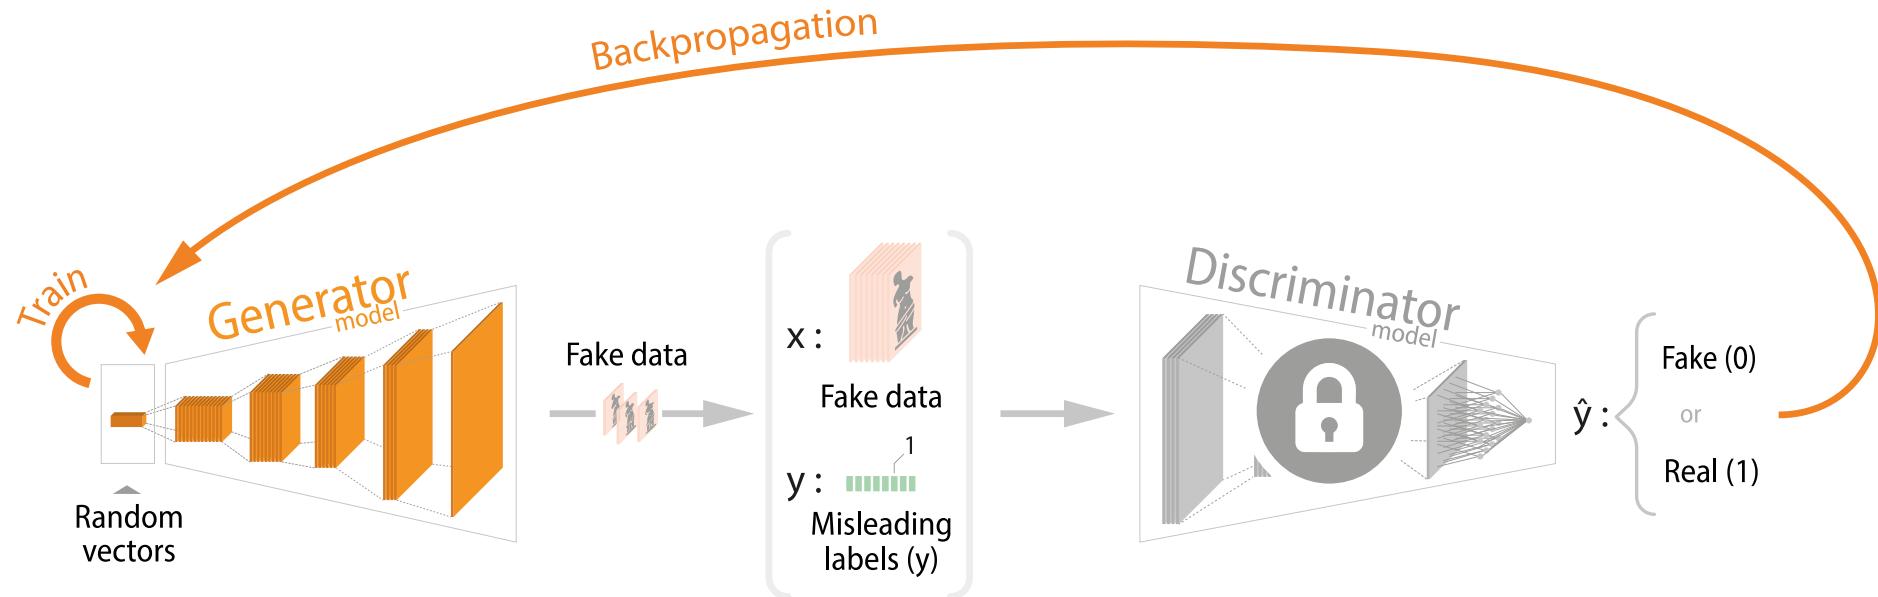
Generator is locked



How to train a GAN ?

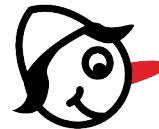
Step2 - Update Generator

Discriminator is locked



Roadmap

13



Generative
Adversarial Networks
GAN

1

About GAN

- Principles
- Training

2

Example 1 : Sheep1

- Draw me a sheep !

3

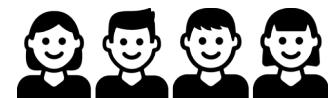
From GANs to *GANs

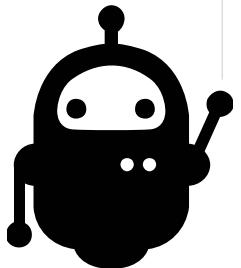
- Wasserstein and Earth Moving Distance
- From GAN to WGAN
- From WGAN to WGAN-GP

4

Example 2 : Sheep2

- Again, but with a WGAN-GP :-)





Draw me a sheep !

Notebook : [SHEEP1]



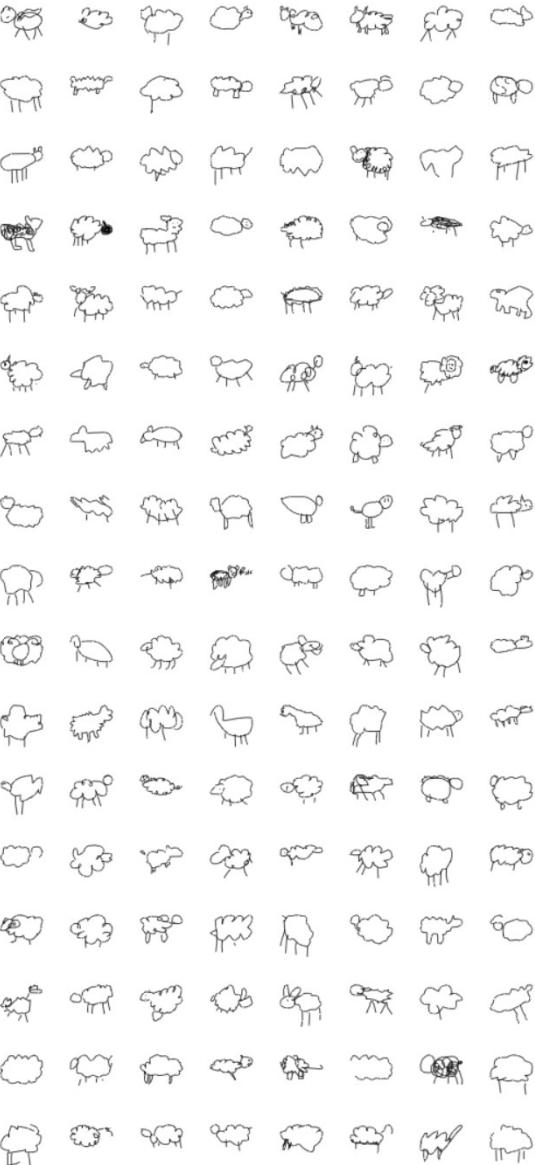
Objective :

Learn to draw or generate sheep.

Dataset :

Quick draw dataset

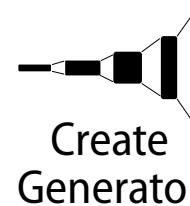
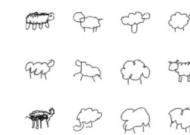
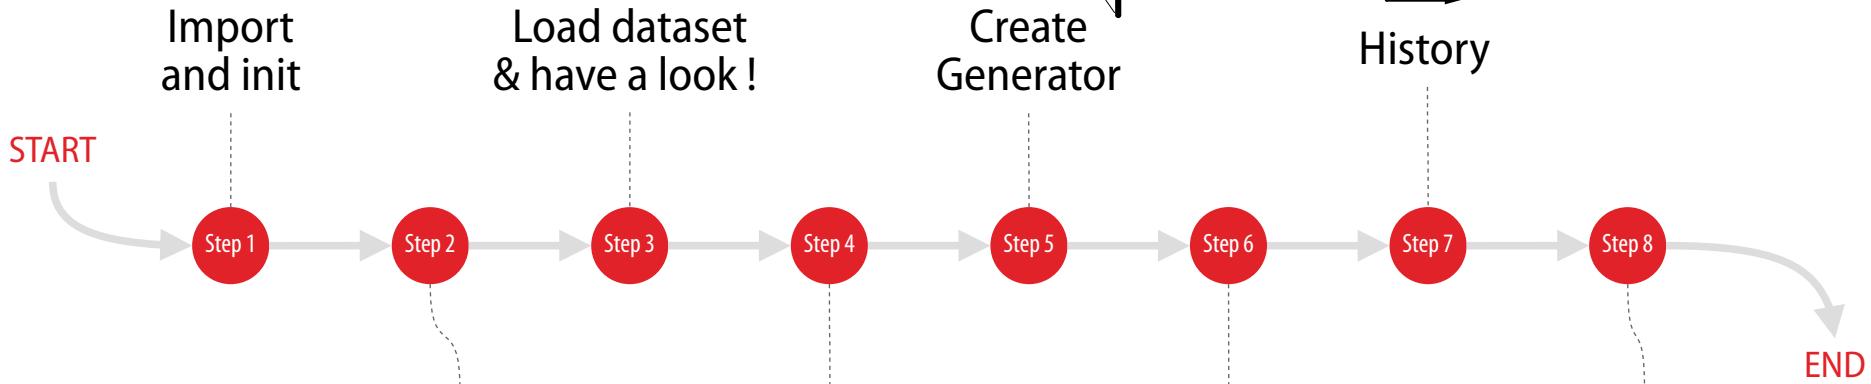
<https://quickdraw.withgoogle.com/data>



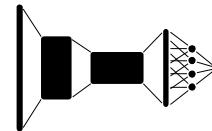


Objectives :

Learn to
draw
sheep !



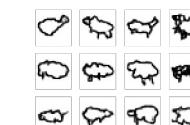
Create
Generator



Create
Discriminator

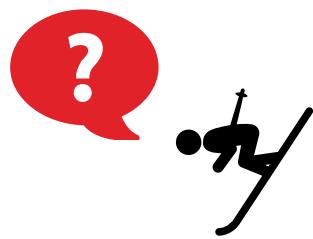


Build, compile
& train



6m on a V100 !
(10 epochs)





Fine, but...
How to calculate a
gradient with
Tensorflow/Keras ?

```
import numpy as np
import tensorflow as tf

# ---- My function f
def f(x):
    y = x**2 + 4*x - 5
    return y

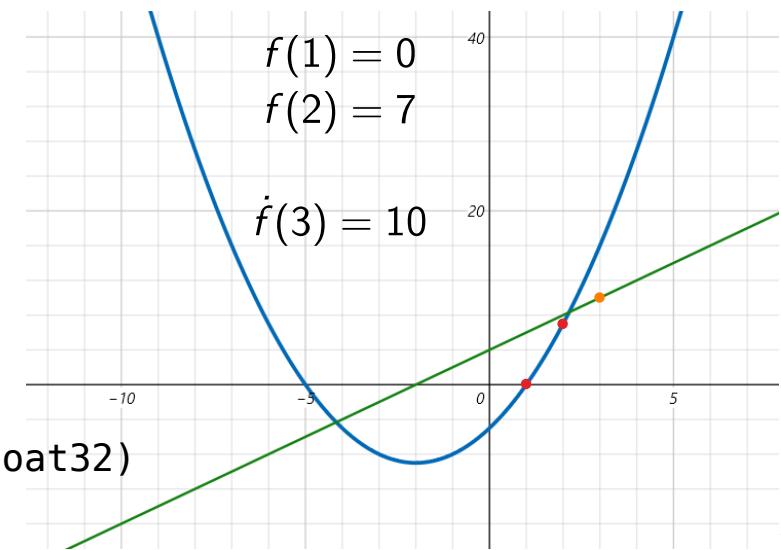
# ---- Examples :
print('f(1) is : ', f(1))
print('f(2) is : ', f(2))

x = tf.Variable(2.0)
print('f(x) is : ', f(x))
```

```
f(1) is : 0
f(2) is : 7
f(x) is : tf.Tensor(7.0, shape=(), dtype=float32)
```

$$\left\{ \begin{array}{l} f(x) = x^2 + 4x - 5 \\ \dot{f}(x) = \frac{\delta f}{\delta x} = 2x + 4 \end{array} \right.$$

We have :



```
# ----- Derivative using tf
#
with tf.GradientTape() as tape:
    x = tf.Variable(3.0)
    y = f(x)

dy = tape.gradient(y, x) # dy/dx

print(dy)
```

```
tf.Tensor(10.0, shape=(), dtype=float32)
```

$$\left\{ \begin{array}{l} f(x) = x^2 + 4x - 5 \\ \dot{f}(x) = \frac{\delta f}{\delta x} = 2x + 4 \end{array} \right.$$

We have :

$$\begin{aligned} f(1) &= 0 \\ f(2) &= 7 \end{aligned}$$

$$\dot{f}(3) = 10$$



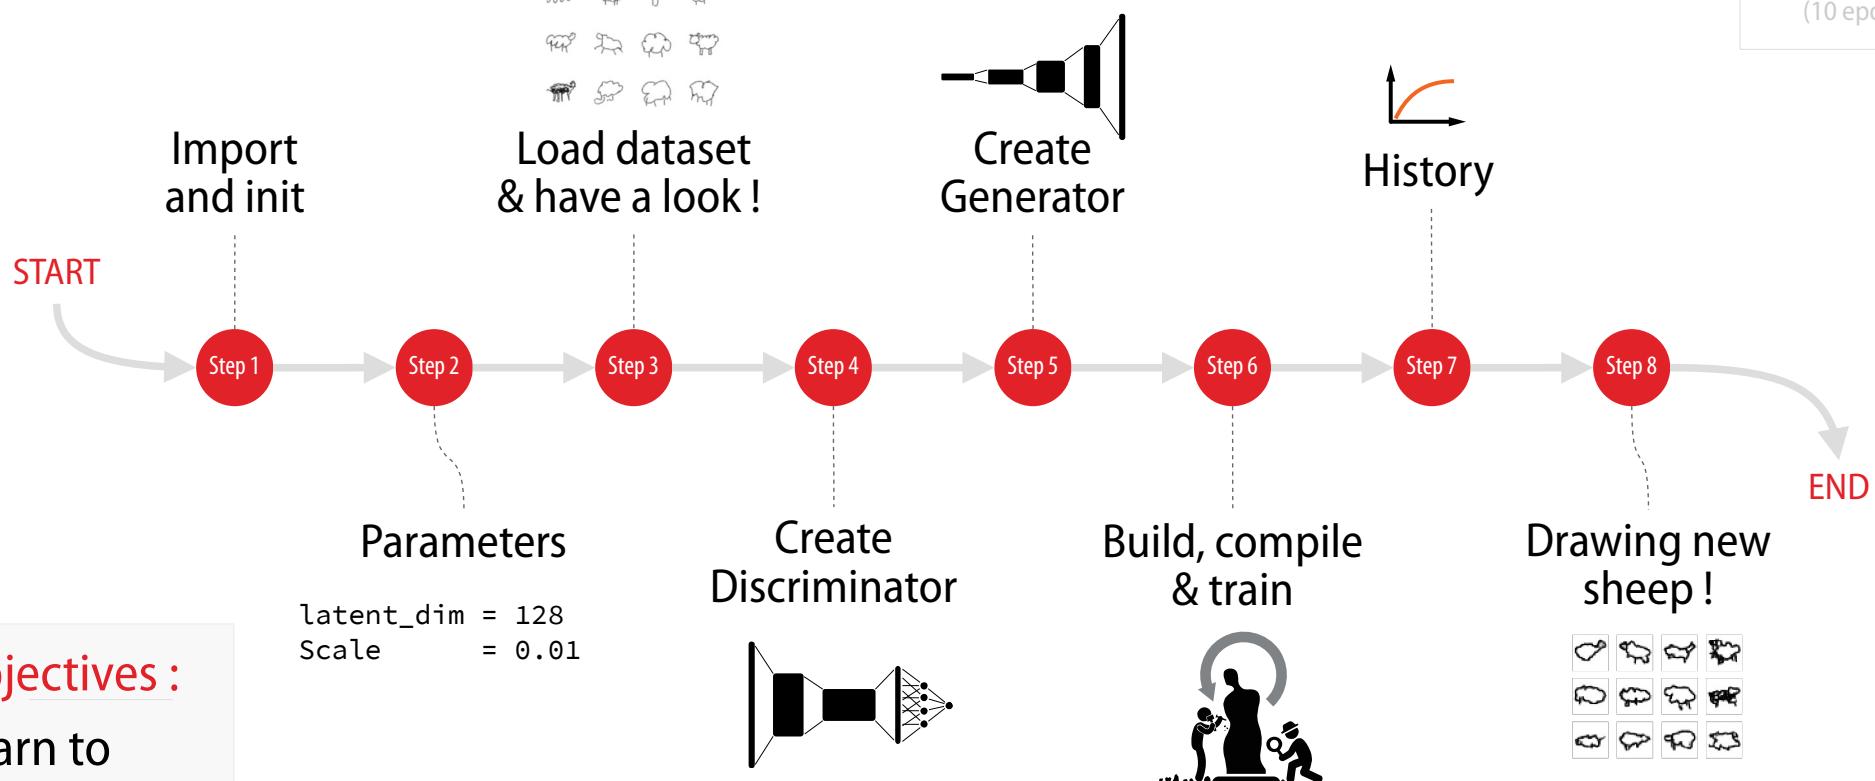
More about Automatic Differentiation and Gradients :

<https://www.tensorflow.org/guide/autodiff>

https://www.tensorflow.org/api_docs/python/tf/GradientTape



6m on a V100 !
(10 epochs)



Roadmap

13



**Generative
Adversarial Networks**
GAN

1

About GAN

- Principles
- Training

2

Example 1 : Sheep1

- Draw me a sheep !

3

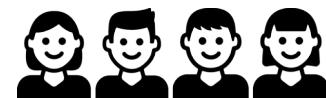
From GANs to *GANs

- Wasserstein and Earth Moving Distance
- From GAN to WGAN
- From WGAN to WGAN-GP

4

Example 2 : Sheep2

- Again, but with a WGAN-GP :-)



Original GAN

2014

→ Original paper [1] :

$$L_D = \frac{1}{m} \sum_{i=1}^m \left[\log(D(x^{(i)})) + \log(1 - D(G(z^{(i)}))) \right]$$

$$L_G = \frac{1}{m} \sum_{i=1}^m \left[\log(1 - D(G(z^{(i)}))) \right]$$

GAN loss functions

$x^{(i)}$

is a real data from a set of m values

$z^{(i)}$

is a latent vector from a set of m values

$D(x)$

Output of the discriminator for a real image x
ie : probability that a real image is considered as real.

$G(z^{(i)})$

Output of the generator from an input z, from latent space
so, $G(z)$ look like an X image, but is really fake...

$D(G(z^{(i)}))$

Output of the discriminator for a fake image.
ie: probability that a fake image is considered as real.

Categorical crossentropy :

$$H(y, \hat{y}) = -\frac{1}{n_{\text{obs}}} \sum_{i=1}^{n_{\text{obs}}} \sum_{c=1}^{n_{\text{class}}} y_c \cdot \log \hat{y}_c$$

Discriminator will try to
maximize L_D :

$$\forall x \in X, D(x) \approx 1$$

$$\forall z \in Z, D(G(z)) \approx 0$$

Generator will try to
minimize L_G :

$$\forall z \in Z, D(G(z)) \approx 1$$

Where : $\begin{cases} \text{fake} : 0 \\ \text{real} : 1 \end{cases}$

[1] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, Yoshua Bengio, « Generative Adversarial Networks »,
<https://arxiv.org/abs/1406.2661>, 2014

Original GAN

2014

→ Original paper [1] :

$$L_D = \frac{1}{m} \sum_{i=1}^m \left[\log(D(x^{(i)})) + \log(1 - D(G(z^{(i)}))) \right]$$

$$L_G = \frac{1}{m} \sum_{i=1}^m \left[\log(1 - D(G(z^{(i)}))) \right]$$

GAN loss functions



Limitations :

Gradient vanishing problem.
Risk of blocking in the early stages of
GAN learning, when the discriminator
work is very easy.

Categorical crossentropy :

$$H(y, \hat{y}) = -\frac{1}{n_{\text{obs}}} \sum_{i=1}^{n_{\text{obs}}} \sum_{c=1}^{n_{\text{class}}} y_c \cdot \log \hat{y}_c$$

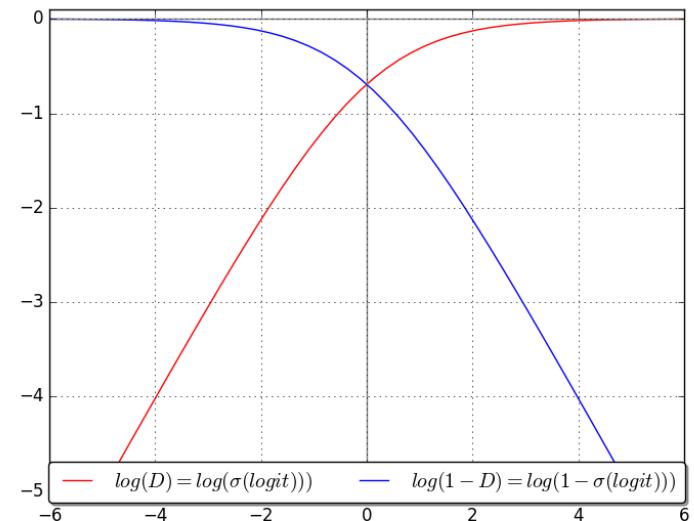


Figure 6: Illustration of vanishing gradient in the negative quadrant when using the original loss formulation $\log(1-D)$ (blue curve). The x-axis is D's logit (output of last layer before sigmoid activation). Minimising $\log(1-D)$ yields the same solution as maximising $\log(D)$ (red curve) but the red curve exhibits stronger gradients.

<https://developer.nvidia.com/blog/photo-editing-generative-adversarial-networks-1/>

[1] Ian J. Goodfellow, Jean Pouget-Abadie, Mehdi Mirza, Bing Xu, David Warde-Farley, Sherjil Ozair, Aaron Courville, Yoshua Bengio, « Generative Adversarial Networks », <https://arxiv.org/abs/1406.2661>, 2014

→ Proposition of a more theoretical approach to the learning strategy [2]:



What is the best way to measure the distance between two distributions ?

Optimization problem [3]

{ Total Variation distance
Kullback-Leibler divergence
Jensen-Shannon divergence
Earth-Mover Distance* / Wasserstein-1

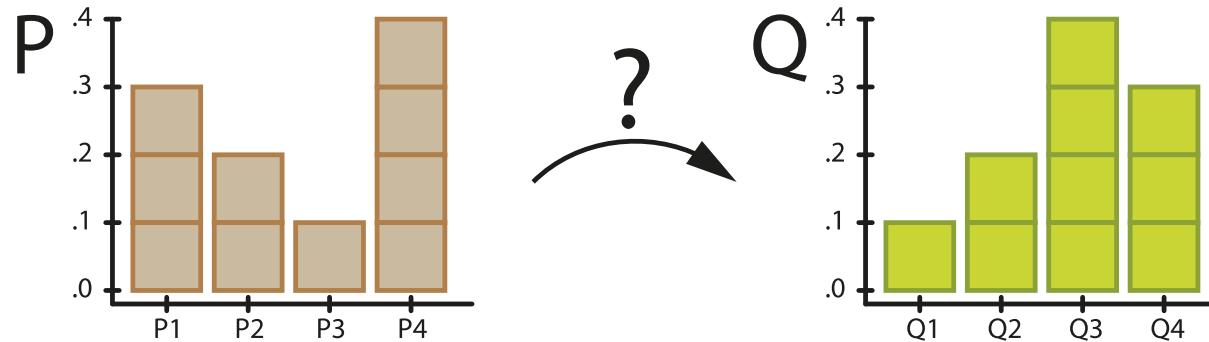
(*) Distance du terrassier

[2] Martin Arjovsky, Soumith Chintala, Léon Bottou, « Wasserstein GAN »,
<https://arxiv.org/abs/1701.07875>, 2017

[3] Cédric Villani, « Optimal transport, old and new », Springer, 2008
https://cedricvillani.org/sites/dev/files/old_images/2012/08/preprint-1.pdf
<https://www.youtube.com/watch?v=zo46TEp6FB8>

Objective is to measuring the distance between 2 distributions

This distance can be interpreted by the minimal energy needed to move from one box distribution to another:



This energy can be calculated as:

$$W = \text{Number of boxes moved} \times \text{Moving distance of the boxes}$$

We can write :

$$W = \sum_i^{\text{class}} |\delta_i|$$

$$\text{with } \begin{cases} \delta_0 = 0 \\ \delta_{i+1} = \delta_i + P_i - Q_i \end{cases}$$

$$\delta_0 = 0$$

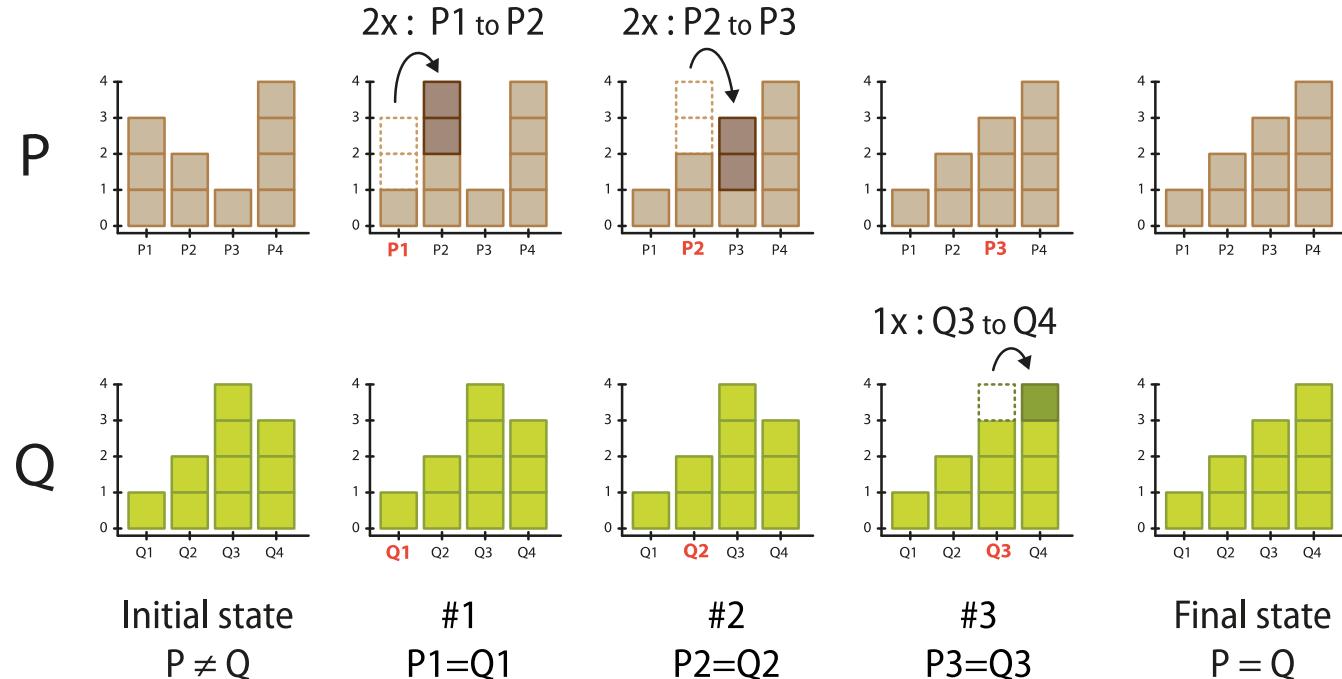
$$\delta_1 = 0 + 3 - 1 = 2$$

$$\delta_2 = 2 + 2 - 2 = 2$$

$$\delta_3 = 2 + 1 - 4 = -1$$

$$\delta_4 = -1 + 4 - 3 = 0$$

$$W = \sum_{i=1}^4 |\delta_i| = 5$$



With a continuous probability domain, the formula can be written :

$$\text{EMD}(P_r, P_g) = \inf_{\gamma \in \Pi(P_r, P_g)} \sum_{x,y} \gamma(x, y) \|x - y\| = \inf_{\gamma \in \Pi} \mathbb{E}_{(x,y) \sim \gamma} \|x - y\|$$

Earth Mover Distance

With : $\Pi(P_r, P_g)$ Set of all possible joint probability distributions between P_r and P_g

$\gamma(x, y)$ Percentage of « sand » that must be transported from point x to point y so that x follows the same probability distribution as y .

$\|x - y\|$ Distance between x and y

A little more on the subject:

[4] Lilian Weng, « From GAN to WGAN », <https://arxiv.org/abs/1904.08994>, 2019

[5] Ludovic Platon , « GAN : Vers une meilleure estimation des distributions » (fr),
<https://www.aquiladata.fr/insights/gan-vers-une-meilleure-estimation-des-distributions/>

Wasserstein GAN

2017, dec .6

→ Wasserstein GAN [2] :

$$L_{\text{critic}} = \frac{1}{m} \sum_{i=1}^m \left[D(x^{(i)}) - D(G(z^{(i)})) \right]$$

$$L_{\text{generator}} = \frac{1}{m} \sum_{i=1}^m \left[D(G(z^{(i)})) \right]$$

WGAN loss functions

$x^{(i)}$

is a real data from a set of m values

$z^{(i)}$

is a latent vector from a set of m values

$D(x)$

Output of the discriminator for a real image x
ie : probability that a real image is considered as real.

$G(z^{(i)})$

Output of the generator from an input z, from latent space
so, $G(z)$ look like an X image, but is really fake...

$D(G(z^{(i)}))$

Output of the discriminator for a fake image.
ie: probability that a fake image is considered as real.



Critic / Discriminator
try to **maximize** L_{critic}

$\forall x \in X, \forall z \in Z$
 $D(x) \gg D(G(z))$

Generator try
to **maximize** $L_{\text{generator}}$

Note : $\begin{cases} \text{fake} : \text{low score} \\ \text{real} : \text{high score} \end{cases}$

Wasserstein GAN

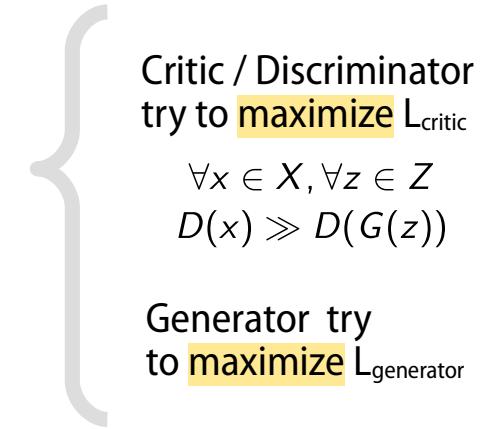
2017, dec .6

→ Wasserstein GAN [2] :

$$L_{\text{critic}} = \frac{1}{m} \sum_{i=1}^m \left[D(x^{(i)}) - D(G(z^{(i)})) \right]$$

$$L_{\text{generator}} = \frac{1}{m} \sum_{i=1}^m \left[D(G(z^{(i)})) \right]$$

WGAN loss functions



! According to the authors, the discriminator should no longer be considered as a classifier but more as a **critic**. The output of the discriminator will be greater for true instances than for false instances. Concretely, a critic has **no sigmoid** function at its output

! The calculation of the EMD distance is complex. A simplified version is however possible if the discriminator is a **k-Lipschitz function**, implying a constraint on it. The proposed solution is to clip the weights to contain them in a **limited interval** $[-c, c]$, with $c=0.01$

Note : $\begin{cases} \text{fake} : \text{low score} \\ \text{real} : \text{high score} \end{cases}$

2017, dec .6

→ Wasserstein GAN [2] :

$$L_{\text{critic}} = \frac{1}{m} \sum_{i=1}^m \left[D(x^{(i)}) - D(G(z^{(i)})) \right]$$

$$L_{\text{generator}} = \frac{1}{m} \sum_{i=1}^m \left[D(G(z^{(i)})) \right]$$

WGAN loss functions



Limitations :

« Weight clipping is a clearly terrible way to enforce a Lipschitz constraint (...) and we actively encourage interested researchers to improve on this method. » [2]

According to the authors, the discriminator should no longer be considered as a classifier but more as a critic. The output of the discriminator will be greater for true instances than for false instances.

Concretely, a critic has no sigmoid function at its output



The calculation of the EMD distance is complex. A simplified version is however possible if the discriminator is a **k-Lipschitz function**, implying a constraint on it. The proposed solution is to clip the weights to contain them in a **limited interval $[-c, c]$** , with $c=0.01$

Note : $\begin{cases} \text{fake} : \text{low score} \\ \text{real} : \text{high score} \end{cases}$

Wasserstein GAN - GP

2017, dec 25

→ Wasserstein GAN with Gradient Penalty [6] :

Here is the gradient
penalty

$$L_{\text{critic}} = \mathbb{E}_{\tilde{x} \sim \mathbb{P}_g} [D(\tilde{x})] - \mathbb{E}_{x \sim \mathbb{P}_r} [D(x)] + \lambda \mathbb{E}_{\hat{x} \sim \mathbb{P}_{\hat{x}}} [(\|\nabla_{\hat{x}} D(\hat{x})\|_2 - 1)^2]$$

$$L_{\text{generator}} = \frac{1}{m} \sum_{i=1}^m [D(G(z^{(i)}))]$$

WGANGP loss functions

x	is a real data from a set of real values
$\tilde{x} = G(z)$	is a fake data, from the generator
\hat{x}	mix fake / real data
λ	penalty coefficient, proposed as $\lambda=10$
$z^{(i)}$	is a latent vector from a set of m values
$D(x)$	Output of the discriminator for a real image x
$G(z^{(i)})$	Output of the generator from an input z, from latent space
$D(G(z^{(i)}))$	Output of the discriminator for a fake image.

Critic / Discriminator
try to minimize L_{critic}

$$\forall x \in X, \forall z \in Z \\ D(x) \gg D(G(z))$$

Generator try
to maximize $L_{\text{generator}}$

$$\hat{x} = \varepsilon x + (1 - \varepsilon)\tilde{x} \\ \text{with } t \sim U[0, 1]$$



The paper demonstrates that
the addition of the penalty
gradient guarantees that the
loss function is 1-Lipschitz.

[6] Ishaan Gulrajani, Faruk Ahmed, Martin Arjovsky, Vincent Dumoulin, Aaron Courville,
« Improved Training of Wasserstein GANs », <https://arxiv.org/abs/1704.00028>, 2017

Algorithm 1 WGAN with gradient penalty. We use default values of $\lambda = 10$, $n_{\text{critic}} = 5$, $\alpha = 0.0001$, $\beta_1 = 0$, $\beta_2 = 0.9$.

Require: The gradient penalty coefficient λ , the number of critic iterations per generator iteration n_{critic} , the batch size m , Adam hyperparameters α, β_1, β_2 .

Require: initial critic parameters w_0 , initial generator parameters θ_0 .

```

1: while  $\theta$  has not converged do
2:   for  $t = 1, \dots, n_{\text{critic}}$  do
3:     for  $i = 1, \dots, m$  do
4:       Sample real data  $\mathbf{x} \sim \mathbb{P}_r$ , latent variable  $\mathbf{z} \sim p(\mathbf{z})$ , a random number  $\epsilon \sim U[0, 1]$ .
5:        $\tilde{\mathbf{x}} \leftarrow G_\theta(\mathbf{z})$ 
6:        $\hat{\mathbf{x}} \leftarrow \epsilon \mathbf{x} + (1 - \epsilon) \tilde{\mathbf{x}}$ 
7:        $L^{(i)} \leftarrow D_w(\tilde{\mathbf{x}}) - D_w(\mathbf{x}) + \lambda(\|\nabla_{\hat{\mathbf{x}}} D_w(\hat{\mathbf{x}})\|_2 - 1)^2$ 
8:     end for
9:      $w \leftarrow \text{Adam}(\nabla_w \frac{1}{m} \sum_{i=1}^m L^{(i)}, w, \alpha, \beta_1, \beta_2)$ 
10:   end for
11:   Sample a batch of latent variables  $\{\mathbf{z}^{(i)}\}_{i=1}^m \sim p(\mathbf{z})$ .
12:    $\theta \leftarrow \text{Adam}(\nabla_\theta \frac{1}{m} \sum_{i=1}^m -D_w(G_\theta(\mathbf{z})), \theta, \alpha, \beta_1, \beta_2)$ 
13: end while
```

Roadmap

13



Generative
Adversarial Networks
GAN

1

About GAN

- Principles
- Training

2

Example 1 : Sheep1

- Draw me a sheep !

3

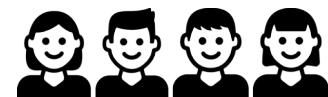
From GANs to *GANs

- Wasserstein and Earth Moving Distance
- From GAN to WGAN
- From WGAN to WGAN-GP

4

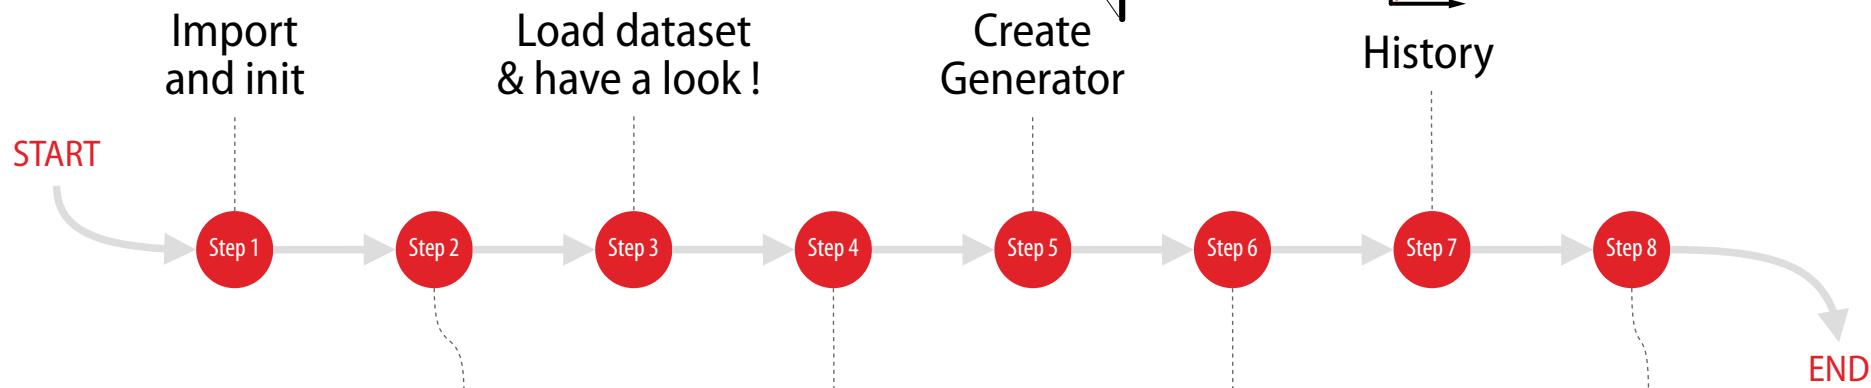
Example 2 : Sheep2

- Again, but with a WGAN-GP :-)





6m on a V100 !
(10 epochs)



Objectives :

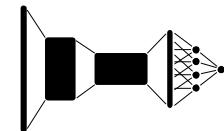
Still learning
to draw a
sheep, but
with a
WGAN-GP !

Parameters

```

latent_dim = 128
Scale      = 0.01
...
n_critic = 2
  
```

Create Discriminator



Build, compile & train



Drawing new sheep !



...and what else ?

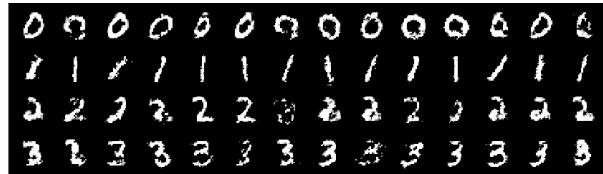
It is time to conclude ! ;-)



GANs, GANs, more GANs...

Conditional GANs

Mehdi Mirza, Simon Osindero,
« Conditional Generative Adversarial Nets »,
<https://arxiv.org/abs/1411.1784>, 2014



Progressive GANs

Tero Karras, Timo Aila, Samuli Laine, Jaakko Lehtinen, « Progressive Growing of GANs for Improved Quality, Stability, and Variation »,
<https://arxiv.org/abs/1710.10196>, 2017



Image-to-Image Translation

Phillip Isola, Jun-Yan Zhu, Tinghui Zhou, Alexei A. Efros,
« Image-to-Image Translation with Conditional Adversarial Networks »,
<https://arxiv.org/abs/1611.07004>, 2016



CycleGAN

Jun-Yan Zhu Taesung Park Phillip Isola Alexei A. Efros
« Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks »,
<https://arxiv.org/abs/1703.10593>, 2017



Text-to-Image Synthesis

Han Zhang, Tao Xu, Hongsheng Li, Shaoting Zhang, Xiaogang Wang, Xiaolei Huang, Dimitris Metaxas,
« StackGAN: Text to Photo-realistic Image Synthesis with Stacked Generative Adversarial Networks »,
<https://arxiv.org/abs/1612.03242>, 2016

« This smaller brown bird has white stripes on the coverts, wingbars and secondaries »



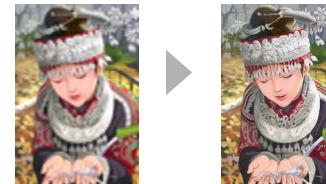
Semantic Image Inpainting

Raymond A. Yeh, Chen Chen, Teck Yian Lim, Alexander G. Schwing, Mark Hasegawa-Johnson, Minh N. Do,
« Semantic Image Inpainting with Deep Generative Models »,
<https://arxiv.org/abs/1607.07539>, 2016



Super-Resolution

Christian Ledig, Lucas Theis, Ferenc Huszar, Jose Caballero, Andrew Cunningham, Alejandro Acosta, Andrew Aitken, Alykhan Tejani, Johannes Totz, Zehan Wang, Wenzhe Shi,
« Photo-Realistic Single Image Super-Resolution Using a Generative Adversarial Network »,
<https://arxiv.org/abs/1609.04802>, 2016



Face Frontal View Generation

Rui Huang, Shu Zhang, Tianyu Li, Ran He,
« Beyond Face Rotation: Global and Local Perception GAN for Photorealistic and Identity Preserving Frontal View Synthesis »,
<https://arxiv.org/abs/1704.04086>, 2017





2014



2015



2016



2017

?



Séquence 14 - **AI, droit, société et éthique**
Jeudi 23 mars, 14h

...and then ?



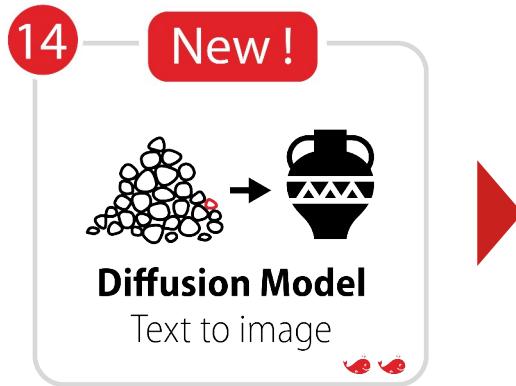
...and then ?



AE → VAE → GAN → Diffusion model

New !

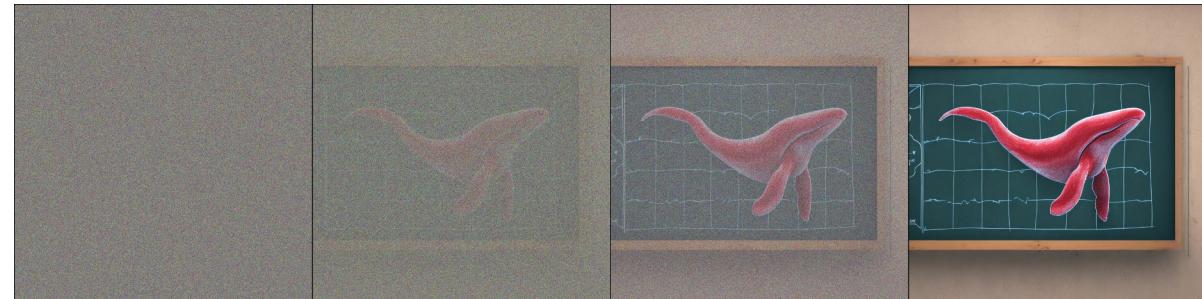
Next, on Fidle :



Jeudi 16 mars,

Épisode 14 :

Diffusion model, text to image



- Principe des Diffusion Model
- Le processus et l'architecture du DDPM
- Améliorations et optimisations des DDPM
- Utilisation des Diffusion Model

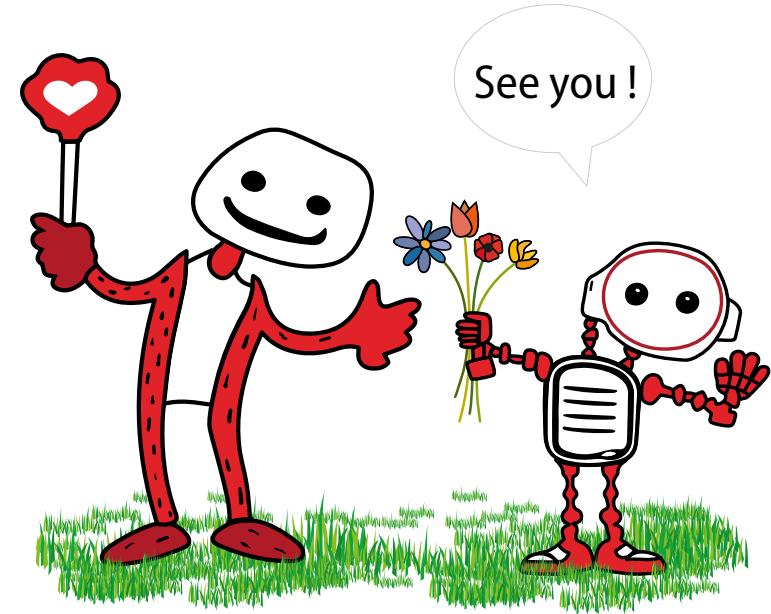
Exemple proposé : Génération d'une garde robe :-)

Durée : 2h

Next on Fidle :



Jeudi 16 mars, Séquence 9: **Diffusion model, text to image**



« A funny person who says hello with his harm »
« A funny robot with a big brain who says hello with his harm. »
by Stable Diffusion 2.1-768

To be continued...