


Formation

Introduction au Deep Learning

19

Des neurones pour la physique
Les Physics-informed neural networks



FIDLE

You can also subscribe to :



<http://fidle.cnrs.fr/listeinfo>
Fidle information list



<https://listes.services.cnrs.fr/wws/info/devlog>
List of ESR* « Software developers » group



GROUPE **CALCUL**

<https://listes.math.cnrs.fr/wws/info/calcul>
List of ESR* « Calcul » group

<https://fidle.cnrs.fr>

Powered by CNRS CRIC, and UGA DGDSI
of Grenoble, Thanks !



Course materials (pdf)









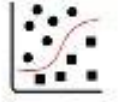







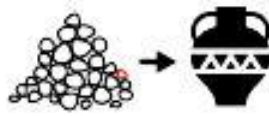




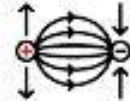

Practical work environment*



Corrected notebooks




Videos (YouTube)

1	 History, Fundamental Concepts	2	3	 Hight Dimensionnal Data CNN	4	 Demystify mathematics for neural networks.	5	 Training strategies Evaluation Sparse data (text) Embedding	6	 Sequences data RNN
 Basic Regression DNN  Basic Classification DNN		7	 PyTorch A small detour with PyTorch .		8	 «Attention is All You Need» Transformers	9	 Graph Neural Network GNN	10	 Autoencoder networks AE
11	 Variational Antoencoder VAE	12	 Project session «My project in 180 s»		13	 Generative Adversarial Networks GAN	14	 Diffusion Model Text to image	15	 AI, Law, Society and Ethics
16	 Model and training optimization Resource efficiency	17	 Jean-Zay GPU acceleration		18	 Deep Reinforcement Learning RL	19	 Physics-Informed Neural Networks PINNS	20	 JDLS 2023 Deep Learning for Science!

20 Séquences
du 17 novembre
au 14 mai 2023

SAISON
22/23



19.1

What are PINNs ?

- Introduction
- Proof of concept

19.2

Loss balancing

- Manual balancing, it's ugly but it works
- Learning rate annealing
- Self adaptive, soft attention like
- Augmented Lagrangian Methods

19.3

Sampling & other tricks


- Non adaptive sampling
- Adaptive sampling
- Tricks you already know



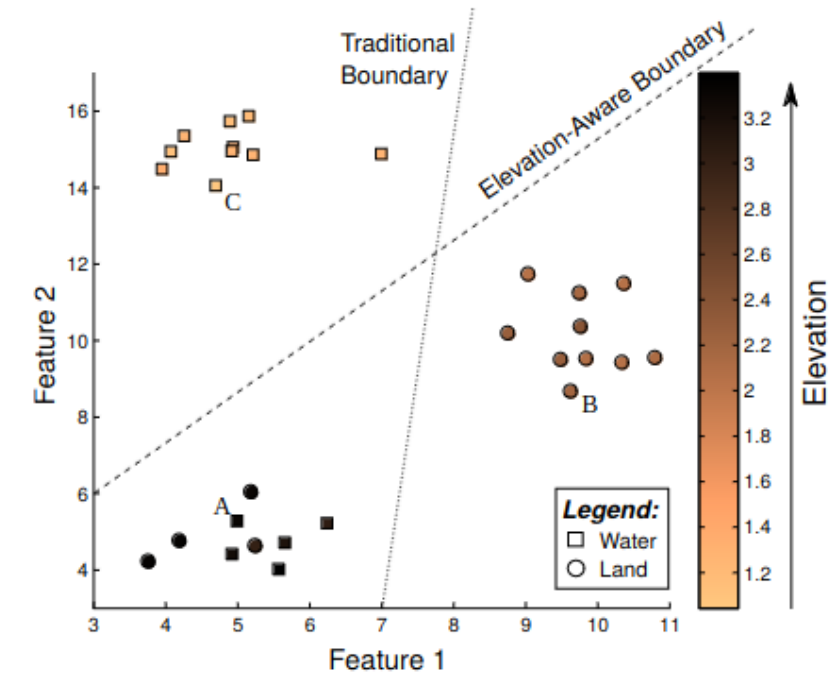
What are PINNs ?



Why theory guided machine learning ?



(b) Lake cross-section.



<https://doi.org/10.1109/TKDE.2017.2720168>


How to inform NN of relations between physical quantities ?

ODE $\frac{dN}{dt} - kN = 0$

PDE $\frac{\partial \rho}{\partial t} + \vec{\nabla} \cdot (\rho \vec{V}) = 0$

IDE $\frac{\partial u}{\partial x} + \int_{x_1}^x f(t, u) dt = g(x, u)$

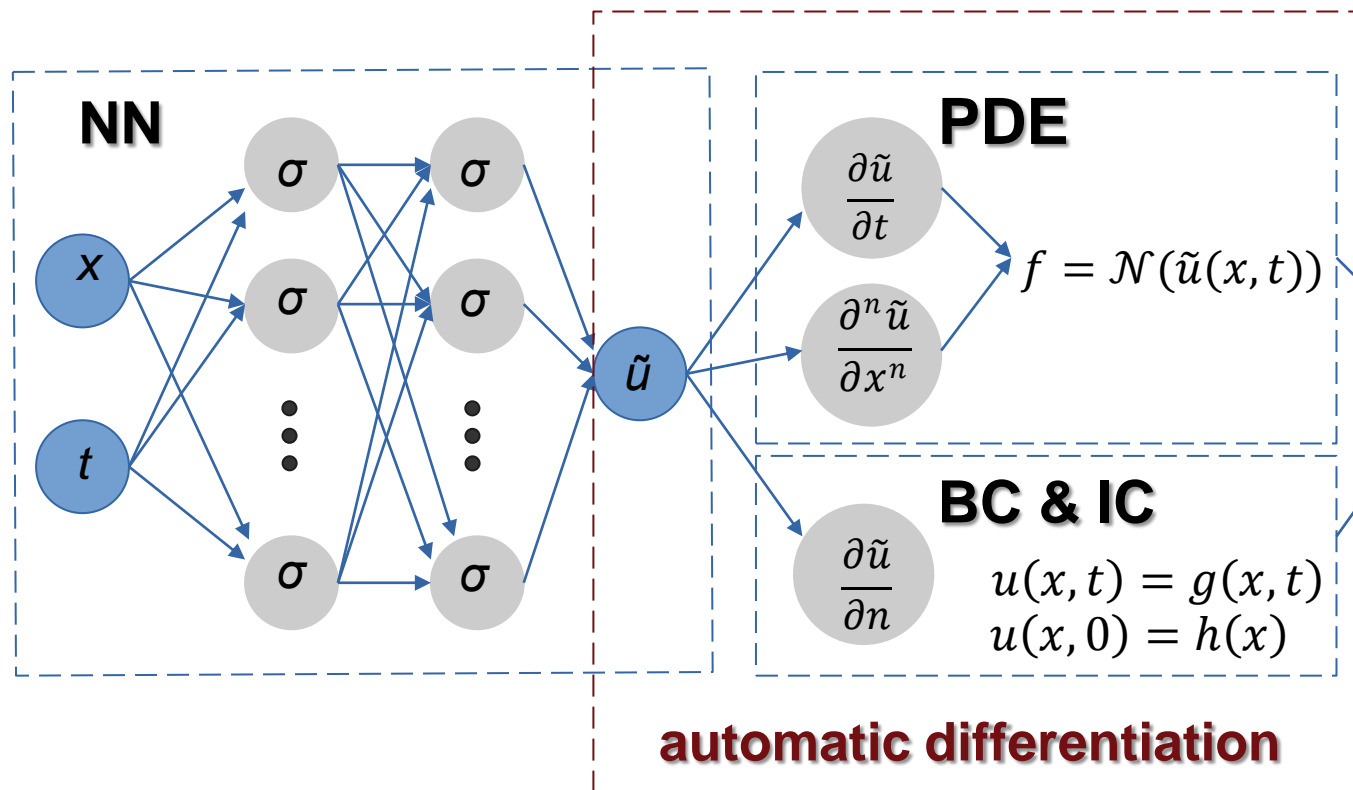
- ❖ Physics Informed NNs
- ❖ Deep Galerkin Method
- ❖ Feynman-Kac formula
- ❖ Deep Ritz Method



$$\mathcal{N}[u] = f, \quad x \in \Omega, \quad t \in [0, T]$$

$$u(x, 0) = h(x), \quad x \in \Omega$$

$$u(x, t) = g(x, t), \quad t \in [0, T], x \in \partial\Omega$$



$$\mathcal{L} = \mathcal{L}_u + \mathcal{L}_{BC} + \mathcal{L}_{IC}$$

Loss functions:

$$\mathcal{L}_u = \frac{1}{N_r} \sum_{i=1}^{N_r} [r(x_r^i, t_r^i)]^2$$

$$\mathcal{L}_{BC} = \frac{1}{N_b} \sum_{i=1}^{N_b} [u(x_b^i, t_b^i) - g_b^i]^2$$

$$\mathcal{L}_{IC} = \frac{1}{N_0} \sum_{i=1}^{N_0} [u(x_b^0, 0) - h_0^i]^2$$

Proof of concept, an example

$$\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} + u = s, \quad x \in [-1,1], \quad y \in [-1,1]$$


$$s = (1 - \pi^2(a_1^2 + a_2^2)) \sin(a_1 \pi x) \sin(a_2 \pi y)$$

$$u(x, y) = 0 \text{ over } \partial\Omega$$

$$a_1 = 2, a_2 = 1$$

Exact

$$u_{exact}(x, y) = \sin(a_1 \pi x) \sin(a_2 \pi y)$$





TP classical PINN


Notebook : [19_classical_pinns.ipynb](#)

Objective :




Implement a simple PINN

Dataset :

No thanks





Proof of concept, an example



Summary

- Outperformed by classical methods to solve PDE
- No mesh, no labeled data
- Ability to solve
 - ❑ ODE, PDE, IDE
 - ❑ ill posed problems, inverse problems, parametric PDE ...



Images from Nvidia Modulus

Many libraries: DeepXDE (<https://doi.org/10.1137/19M1274067>), Nvidia Modulus (<https://docs.nvidia.com/deeplearning/modulus/index.html>), NeuralPDE.jl, ...


Case of study ill posed: inverse problem

Retrieve PDE from 'experimental points' (assumed PDE shape and given initial guess)


$$\frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} + u = s, \quad x \in [-1,1], \quad y \in [-1,1]$$

$$s = (1 - \pi^2(a_1^2 + a_2^2)) \sin(a_1 \pi x) \sin(a_2 \pi y), \quad a_1 \text{ and } a_2 \text{ unknown}$$


$$u(x, y) = 0 \text{ over } \partial\Omega$$



Case of study ill posed: inverse problem



	a_1	a_2
correct	2	1
estimated	2.00027	1.00041



Try varying the initial guess for a_1 and a_2 , the number of points (collocation, boundaries, experimental), add noise, add terms to the PDE (and new unknown variables) ...

Loss regularization




It's naive but it works !

$$\mathcal{L} = \mathcal{L}_u + \omega_{BC} \mathcal{L}_{BC} + \omega_{IC} \mathcal{L}_{IC}$$

classic


Error



$\omega_{BC} = 1$

naive


Error



$\omega_{BC} = 100$

Unbalanced gradients during backpropagation are common failure for PINNs



$$\mathcal{L} = \mathcal{L}_u + \sum_i \lambda_i \mathcal{L}_i$$

 **collocation** **supervised (experiment, boundaries, initial ...)**

$$\lambda_i \leftarrow (1 - \alpha)\lambda_i + \alpha \frac{\max_{\theta}(|\nabla_{\theta} \mathcal{L}_u|)}{|\nabla_{\theta} \mathcal{L}_i|}$$

$$\theta_{n+1} = \theta_n - \eta \nabla_{\theta} \mathcal{L}_u - \eta \sum_i \lambda_i \nabla_{\theta} \mathcal{L}_i$$

Learning rate annealing





$$\mathcal{L}(\theta, \lambda_u, \lambda_b, \lambda_0) = \mathcal{L}_u(\theta, \lambda_u) + \mathcal{L}_{BC}(\theta, \lambda_b) + \mathcal{L}_{IC}(\theta, \lambda_0)$$



$$\mathcal{L}_u = \frac{1}{N_r} \sum_{i=1}^{N_r} g(\lambda_r^i) [r(x_r^i, t_r^i, \theta)]^2$$

$$\mathcal{L}_{BC} = \frac{1}{N_b} \sum_{i=1}^{N_b} g(\lambda_b^i) [(u(x_b^i, t_b^i, \theta) - g_b^i)]^2$$

$$\mathcal{L}_{IC} = \frac{1}{N_0} \sum_{i=1}^{N_0} g(\lambda_0^i) [(u(x_b^0, 0, \theta) - h_0^i)]^2$$

$$\min_{\theta} \max_{\lambda_u, \lambda_b, \lambda_0} \mathcal{L}(\theta, \lambda_u, \lambda_b, \lambda_0)$$






Augmented Lagrangian relaxation

$$\mathcal{L}_{BC} = \frac{\beta}{N_b} \sum_{i=1}^{N_b} [u(x_b^i, t_b^i) - g_b^i]^2 + \frac{1}{N_b} \sum_{i=1}^{N_b} \lambda_i (u(x_b^i) - g_b^i)$$

$$\min_{\theta} \max_{\lambda_i} \mathcal{L}(\theta, \lambda_i)$$



Without modification of the descent method

LPINN: Lagrangian PINNs: A causality-conforming solution to failure modes of physics-informed neural networks, R. Mojdani et al,
<https://doi.org/10.48550/arXiv.2205.02902>

hPINNs: Physics-Informed Neural Networks with Hard Constraints for Inverse Design.
L. Lu et al, <https://doi.org/10.1137/21M1397908>

PECANN: Physics and Equality Constrained Artificial Neural Networks:
Application to Forward and Inverse Problems with Multi-fidelity Data Fusion,
Shamsulhaq Basir et al <https://doi.org/10.1016/j.jcp.2022.111301>

Sampling

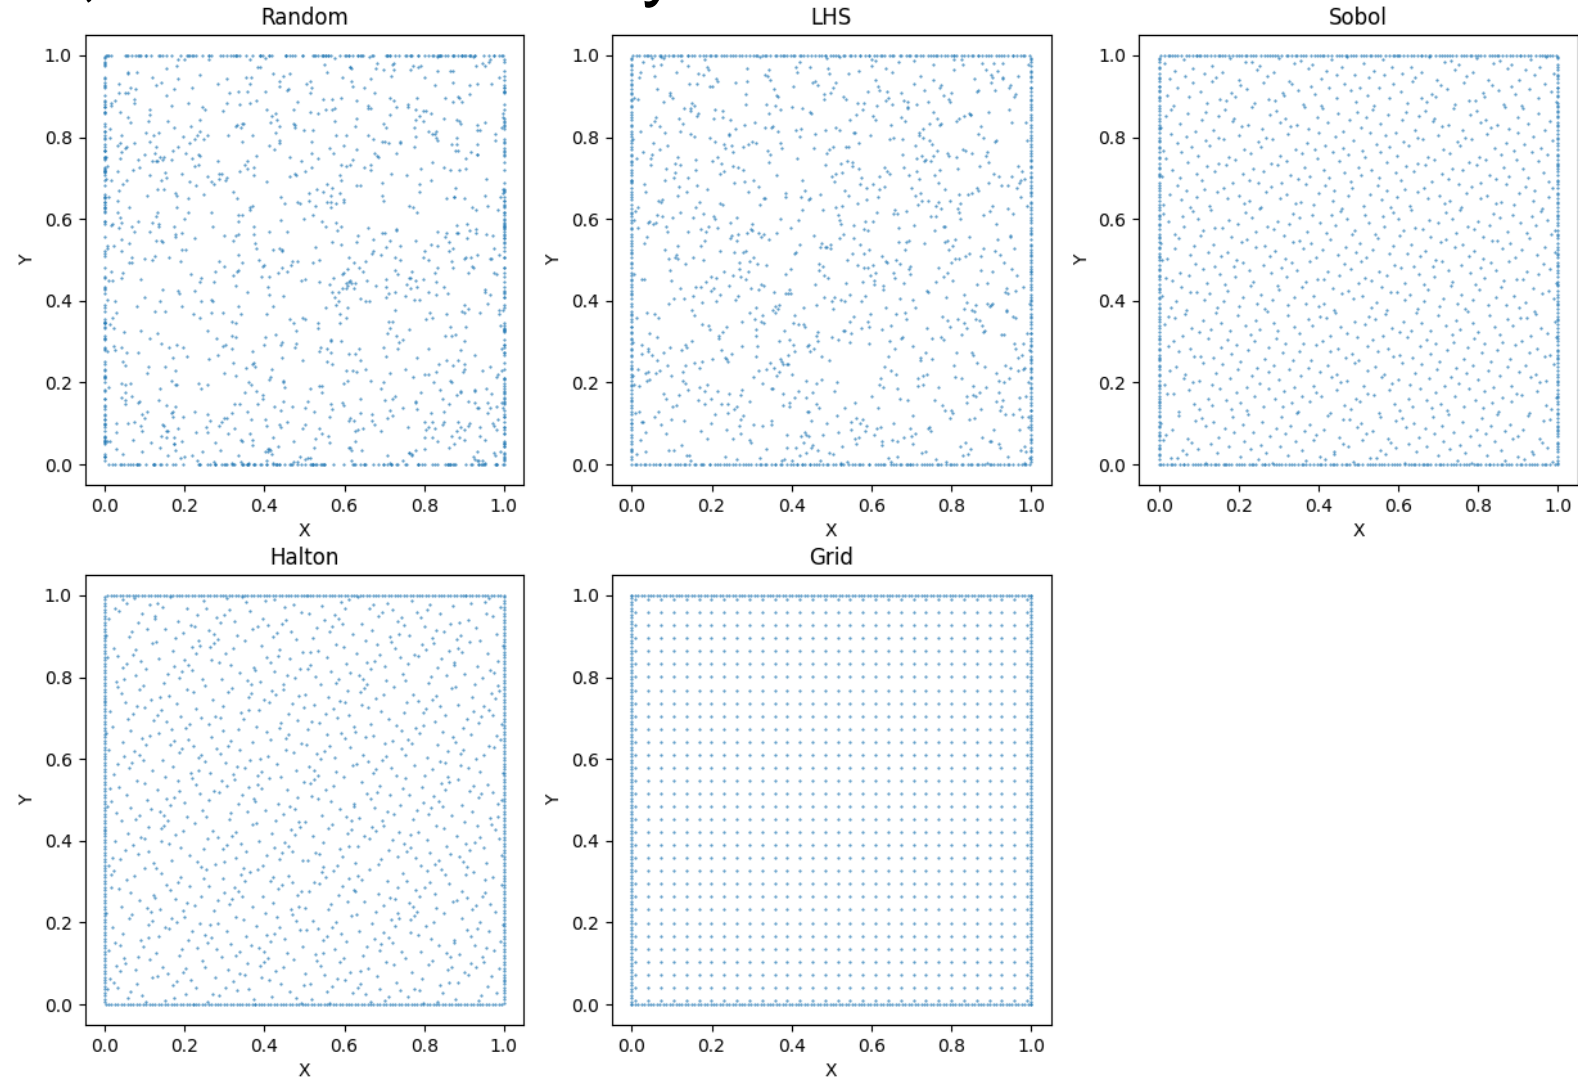


FIDLE

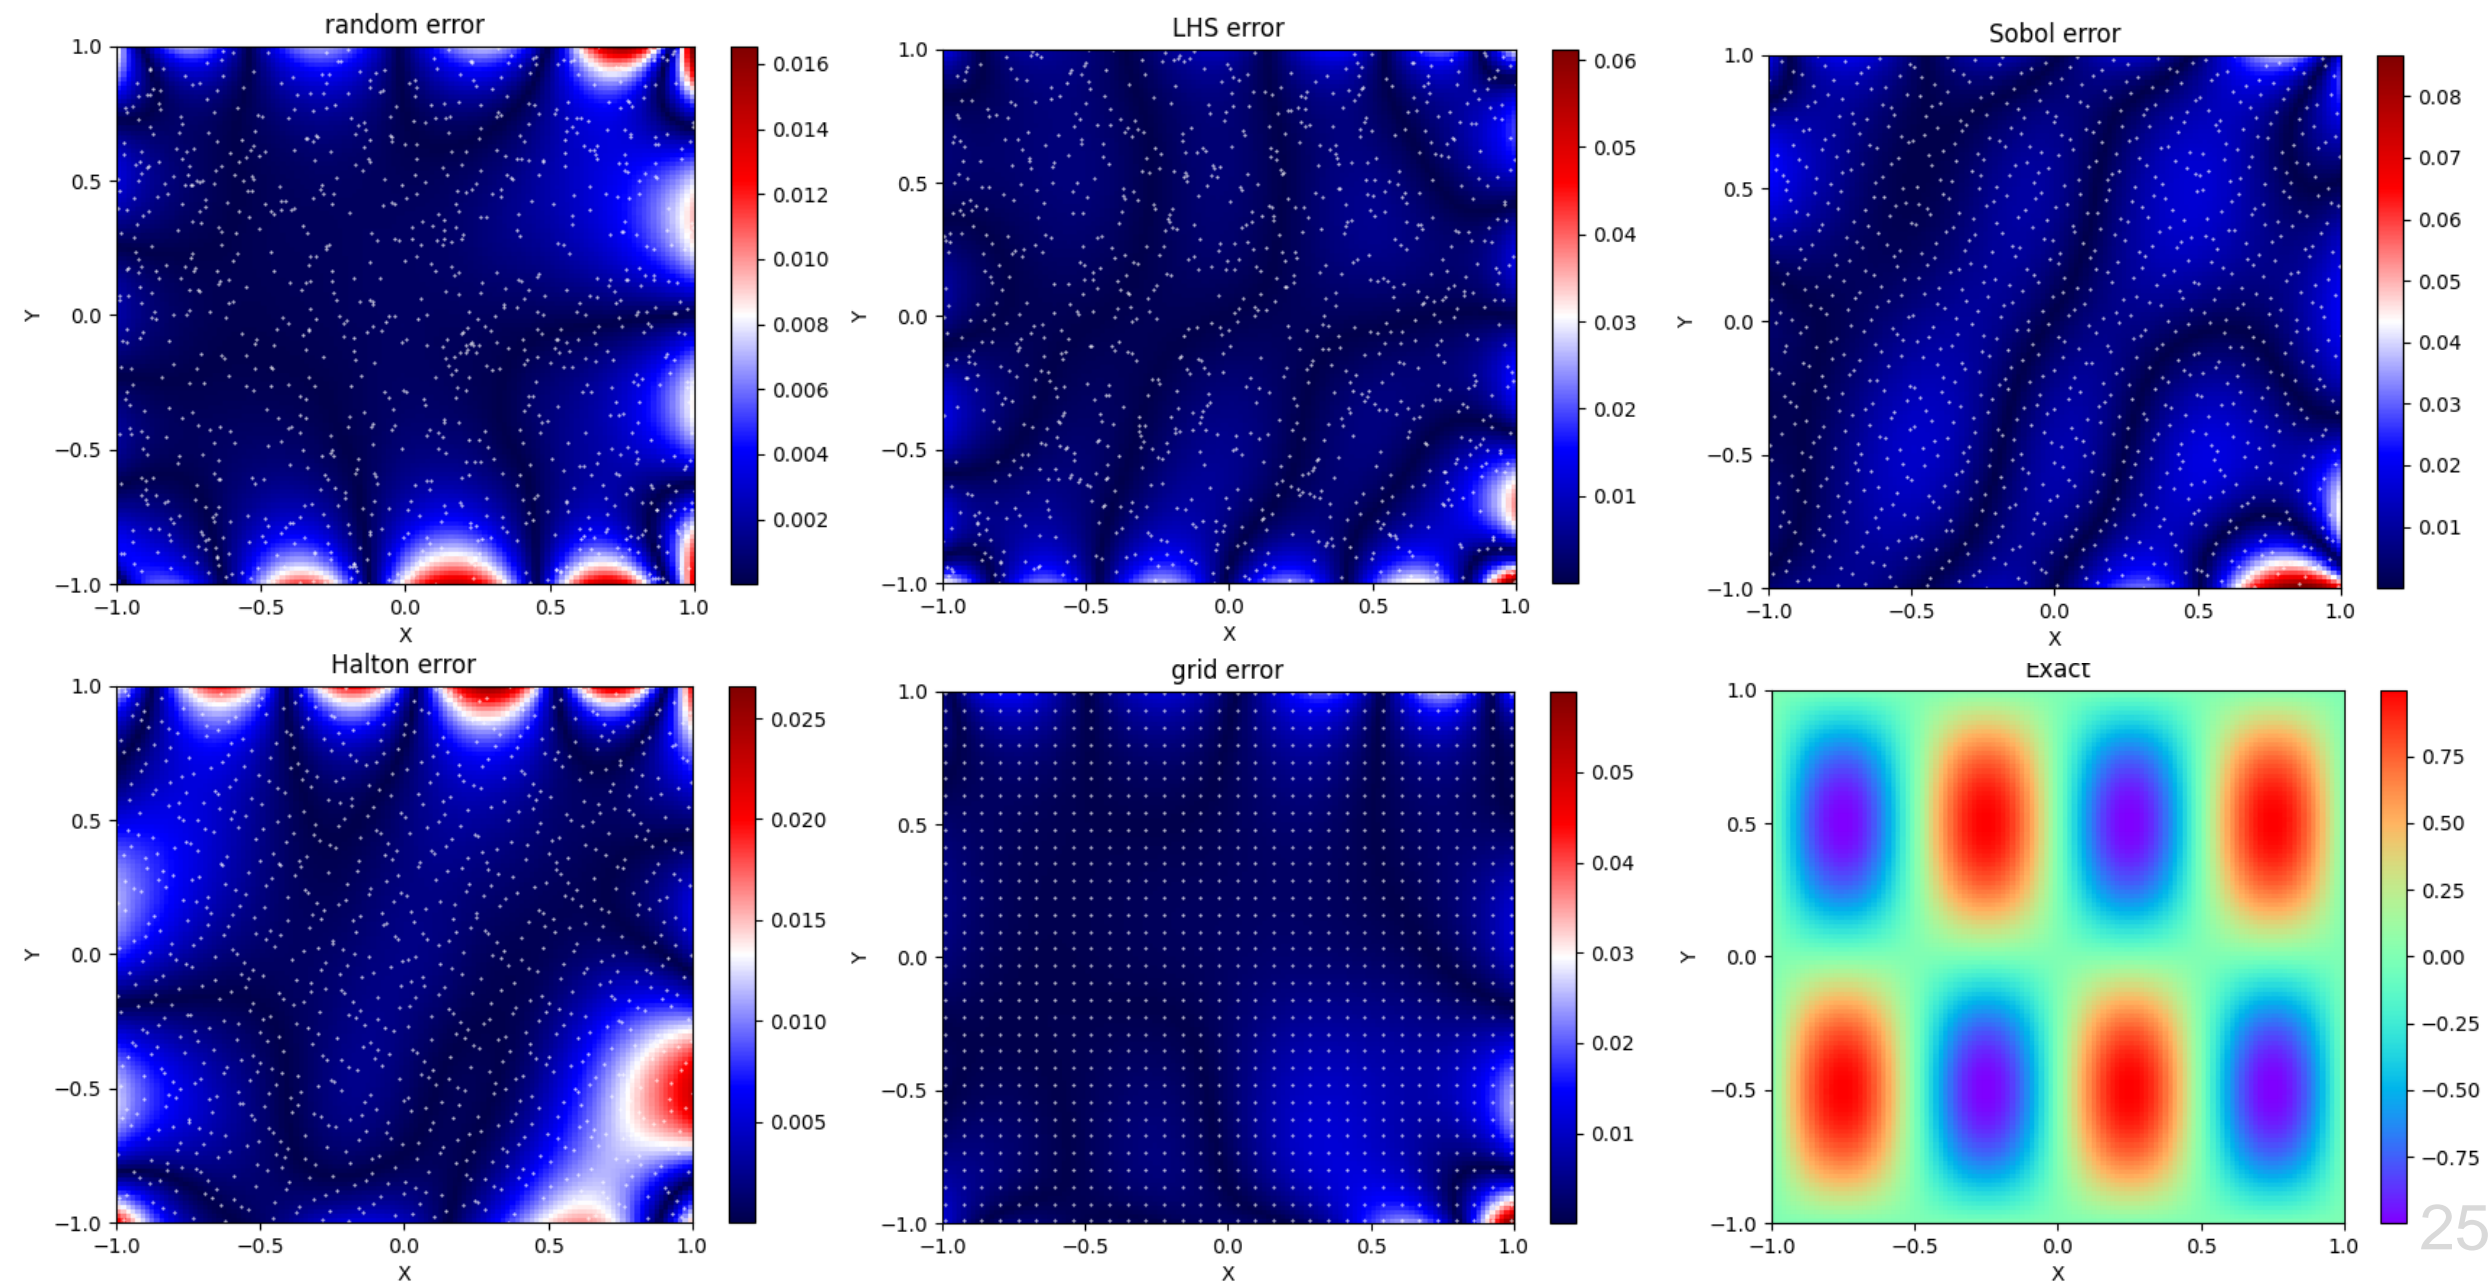
Basics

- ❑ batch [already seen, not covered here]
- ❑ number of points, how much can you offer

Sampling type



Non adaptive sampling



Residual based adaptive refinement with greed (RAR-G):

- ❑ Create sample of test points
- ❑ Determine PDE residual $\varepsilon(x) = |\mathcal{N}[\tilde{u}] - f|$
- ❑ Add points with largest residual to train sample

```
if self.iteration >= Start and (self.iteration - Start)%frequency == 0 and self.x.shape[0] < self.NCollocationMax:
    X_test, _ = self.sampler(Ntest) # Your hand crafted sampling method (grid, random uniform, ...)
    residu = self.residual_PDE(X_test[:,0:1], X_test[:,1:2])
    _, indices = torch.topk(residu**2, self.GreedNum, dim=0, sorted=False)
    self.x = torch.vstack ( (self.x, torch.take(X_test[:,0:1], indices).clone().detach().requires_grad_(True)) )
    self.y = torch.vstack ( (self.y, torch.take(X_test[:,1:2], indices).clone().detach().requires_grad_(True)) )
    self.f_target = torch.zeros(self.x.shape).to(device)
```

Residual based adaptive distribution (RAD):

- ❑ Create sample of test points
- ❑ Determine PDE residual $\varepsilon(x) = |\mathcal{N}[\tilde{u}] - f|$
- ❑ Determine probability distribution function according to PDE residual

$$p(x) = \frac{\varepsilon(x)}{A} \text{ with } A = \int \varepsilon(x) dx$$

- ❑ Take new sample of train points randomly according to PDF

```
if self.iteration >= Start and (self.iteration - Start)%frequency == 0:
    X_test, _ = self.sampler(Ntest) # Your hand crafted sampling method (grid, random uniform, ...)
    residu = self.residual_PDE(X_test[:,0:1], X_test[:,1:2])
    norm = torch.sum(torch.abs(residu))
    distribution = (torch.abs(residu) / norm).flatten()
    indices = torch.multinomial(distribution, self.x.shape[0])
    self.x = torch.take(X_test[:,0:1], indices[:,None])
    self.y = torch.take(X_test[:,1:2], indices[:,None])
    '# create new leaf variable to move grad_fn from UnsqueezeBackward0 to ToCopyBackward0'
    self.x = self.x.clone().detach().requires_grad_(True)
    self.y = self.y.clone().detach().requires_grad_(True)
```

Residual based adaptive distribution with greed (RAD-G):

- ❑ Create probability density function as in RAD
- ❑ Add points taken randomly according to PDF to train sample

```
if self.iteration >= Start and (self.iteration - Start)%frequency == 0 and self.x.shape[0] < self.NCollocationMax:
    X_test, _ = self.sampler(Ntest) # Your hand crafted sampling method (grid, random uniform, ...)
    residu = self.residual_PDE(X_test[:,0:1], X_test[:,1:2])
    norm = torch.sum(torch.abs(residu))
    distribution = (torch.abs(residu) / norm).flatten()
    indices = torch.multinomial(distribution, self.GreedNum)
    x_new = torch.take(X_test[:,0:1], indices[:,None])
    y_new = torch.take(X_test[:,1:2], indices[:,None])
    x_new = x_new.clone().detach().requires_grad_(True)
    y_new = y_new.clone().detach().requires_grad_(True)
    self.x = torch.vstack ( (self.x, x_new) )
    self.y = torch.vstack ( (self.y, y_new) )
    self.f_target = torch.zeros(self.x.shape).to(device)
```

Adaptive sampling

$$\frac{\partial u}{\partial t} - c \frac{\partial u}{\partial x} = 0, \quad x \in [0, 2\pi], \quad t \in [0, 1]$$


$$u(x, 0) = \sin(x), \quad x \in [0, 2\pi]$$

$$u(0, t) = u(2\pi, t), \quad t > 0$$

RAR-G




RAD

RAD-G




$c = 10$



Iteration = 500



Common troubles



$$c = 10$$






$$c = 50$$

Evolutionary sampling (EVO):




- ❑ Take absolute value of residu for each collocation points
- ❑ Keep collocation points for which the residu $>$ mean residual value
- ❑ Resample randomly the other collocation points



```
residu      = torch.abs(self.residual_PDE(self.x, self.t)).to(device)
threshold   = torch.mean(residu)
indices     = torch.stack(list((torch.nonzero((residu >= threshold)*1, as_tuple=True)[0])), dim=0)
x_evo       = torch.take(self.x, indices)[:,None]
t_evo       = torch.take(self.t, indices)[:,None]
x_rand      = (self.lb + (self.ub - self.lb) * torch.rand(self.x.shape[0]-x_evo.shape[0]).float()[:None]).to(device)
t_rand      = (torch.rand(self.t.shape[0] - t_evo.shape[0]).float()[:None]).to(device)
x_evo       = torch.vstack((x_evo, x_rand))
t_evo       = torch.vstack((t_evo, t_rand))
self.x      = x_evo.clone().detach().requires_grad_(True)
self.t      = t_evo.clone().detach().requires_grad_(True)
```

classic




EVO



$c = 10$  $c = 50$ 

Many others tricks

Sequence to sequence



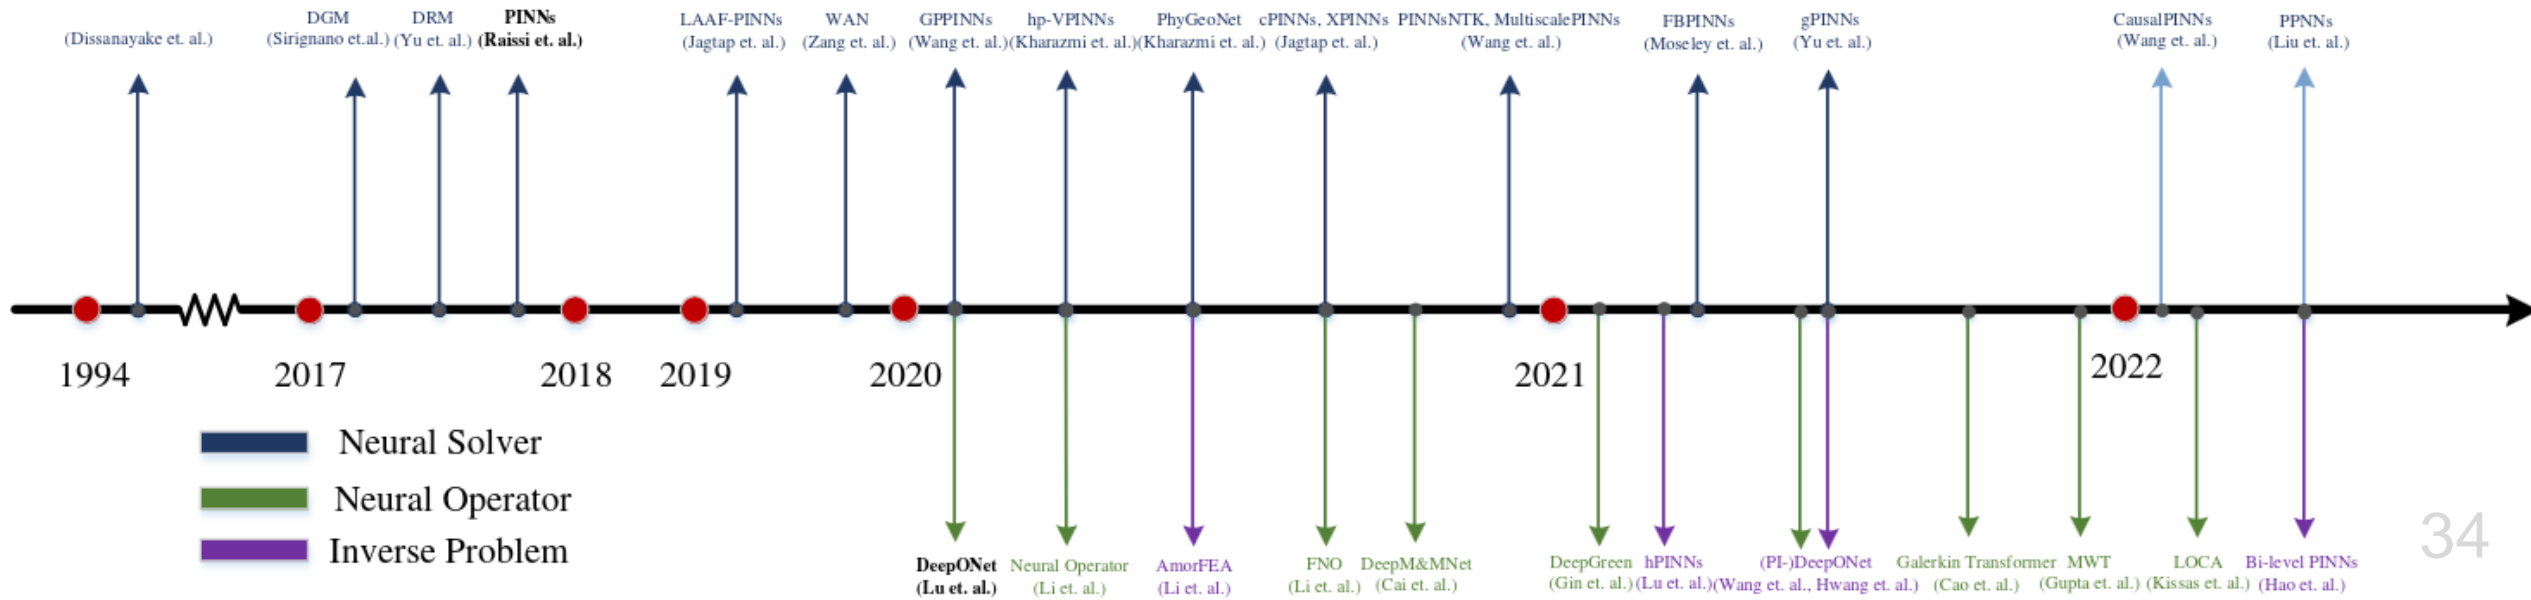
<https://arxiv.org/pdf/2109.01050.pdf>

Gradient enhanced

$$\mathcal{L}_{gPINN} = \mathcal{L}_{PINN} + \frac{1}{N} \sum_{i=1}^N \left| \frac{\partial r}{\partial x_i} \right|^2$$

<https://doi.org/10.1016/j.cma.2022.114823>

And so much more... (<https://doi.org/10.48550/arXiv.2211.08064>)



Architecture of NN


- ❑ MLP (scattered sensor measurements)
- ❑ RNN (discrete time series)
- ❑ CNN (gridded 2D domains, images)
- ❑ GNN, PointNet (Unstructured grids and irregular geometries)

Self adaptive activation functions (<https://doi.org/10.1016/j.jcp.2019.109136>)

Learning operators

- ❑ DeepONet (<https://doi.org/10.1038/s42256-021-00302-5>)
- ❑ FourierNeuralOperator (<https://doi.org/10.48550/arXiv.2010.08895>)


Questions Break



Next, on Fidle :

1^{re} Journée **Deep learning** pour la **Science**

ORSAY 12 MAI 2023




<http://www.idris.fr/annonces/jdls2023.html>

Rediffusion en live sur : <https://www.youtube.com/@CNRS-FIDLE>


Next, on Fidle :

Merci !



1^{re} Journée **Deep learning** pour la Science

ORSAY 12 MAI 2023



<http://www.idris.fr/annonces/jdls2023.html>

Rediffusion en live sur : <https://www.youtube.com/@CNRS-FIDLE>

To be continued...

contact@fidle.cnrs.fr

FIDLE

<https://fidle.cnrs.fr>



YouTube

<https://fidle.cnrs.fr/youtube>



FIDLE



Attribution-NonCommercial-NoDerivatives 4.0 International (CC BY-NC-ND 4.0)
<https://creativecommons.org/licenses/by-nc-nd/4.0/>